

Bias Adjusted Latent Factor Model of Collaborative Filtering for a Recommender System

JOSEPH O'CONNOR, M.S. CHEMISTRY

jpoconnor1961@gmail.com

September 25, 2023

Abstract

This report describes a bias adjusted latent factor model of collaborative filtering for a recommender system, and the key steps in development of the algorithms used in the model. One of the algorithms, TRLIM,¹ involves a classic ridge regression optimization approach that guarantees best-fit solutions to systems of linear equations that are not invertible. The approach is fully explained in the model and has many useful applications beyond recommender systems, as non-invertible systems of linear equations are frequently a problem in real world data. Furthermore, the TRLIM algorithm may be used to assign interaction biases between users and items, which has implications for increasing transparency in Artificial Intelligence.

The algorithm used to calculate the latent factors in the model involved matrix factorization of a very sparse high-dimensional user-movie dataset that was missing about 98.9% of the values due to unrated movies by the users. Sparse data conditions such as this are common in recommender systems, e.g., MovieLens, Netflix, YouTube, etc., which can cause performance problems, particularly in collaborative based recommender systems. A common approach to mitigate the sparse data problem is to impute the missing data values with estimates based on the available data and an understanding, or assumption, of why the data are missing. Due to the overwhelming proportion of data that imputation affects under sparse conditions, the method of imputation must be carefully chosen and evaluated because of its potential to profoundly impact the accuracy of the model. Three common methods of single imputation and one matrix completion method of imputation were evaluated for their relative improvements on the accuracy of the model: Zero Imputation, Column Mean Imputation, Average of Row & Column Means Imputation, and Low-Rank Matrix Completion (LRMC) Imputation. LRMC clearly outperformed the three common methods of imputation, with a significant 4.85% improvement in the model accuracy. Furthermore, key aspects of each imputation method are evaluated for their impact and are discussed in the report.

More specifically, the algorithm used to calculate the latent factors in the model involved matrix factorization, of a LRMC imputed high-dimensional user-movie dataset, with a Principal Components Regression (PCR) approach² that minimized prediction RMSE³ by optimizing truncation of the SVD-PCA⁴ principal component results. The SVD-PCA results were truncated after the first 106 (out of 10,677) principal components to denoise the latent factors that were uncovered from the aggregated dominant patterns of user and movie covariance proportions in the data.

Prior to matrix factorization, the model was adjusted by separation of the overall mean rating of the dataset, as well as separation of the regularized rating biases by various static factors such as movie influenced rating bias, user influenced rating bias, and genre influenced rating bias. The model also captured a regularized temporal factor in the movie influenced rating bias, as well as a regularized user-specific factor in the genre influenced rating bias, which were both included in the biases separated out during adjustments prior to matrix factorization. Furthermore, the calculation and regularization of the user-specific factor in the genre influenced rating bias was made possible by the optimization-centric TRLIM algorithm.

¹Tikhonov Regularized Left-Inverse Matrix (TRLIM)

²Gareth et al. (2021), *An Introduction to Statistical Learning (2nd Ed.)*, pp. 256-259

³Root Mean Squared Error (RMSE): $\sqrt{\frac{1}{n} \sum_1^n (\text{rating} - \text{prediction})^2}$ where n = total number of items rated & predicted.

⁴Singular Value Decomposition - Principal Component Analysis (SVD-PCA)

The TRLIM algorithm is a ridge regression solution that was derived analytically by minimizing the gradient of the Lagrangian form of a least squares cost function, containing Tikhonov L_2 regularization⁵ of the user’s genre-element vector of biases \mathbf{x} , in an over-determined system of linear equations. The derived solution to this constrained minimization problem uses a Lagrangian multiplier, in the form of a tuned Tikhonov regularization parameter, to further reduce prediction RMSE of the user-specific genre biases that were solved in an otherwise non-invertible system of linear equations. Non-invertible systems of linear equations arose from the over-determined non-square matrices of genre elements extracted from users’ movie ratings, and because the m rows or the n columns of the matrices often had lower ranks of independence. The TRLIM algorithm successfully solved such systems of ill-posed and ill-conditioned linear equations, and parsed user-specific biases for each of the genre-elements contained in each user’s data of rated movies. Furthermore, the algorithm also included non-linear scaling with a weighted hyperbolic-tangent function in order to constrain variability in the TRLIM generated user-specific genre biases. Overall, the algorithm made a 1.44% improvement in the model accuracy.

The data used to develop the model was the MovieLens 10M, which is a stable benchmark dataset with 10 million movie ratings that is maintained by GroupLens Research at the University of Minnesota. The accuracy of the final recommender model was measured at a prediction RMSE of 0.792085 on a *final-holdout-test* partition of the 10M dataset.

Keywords: recommender system, matrix factorization collaborative filtering, latent factor model, Low-Rank Matrix Completion (LRMC), imputation, truncated SVD-PCA, Tikhonov ridge regression, MovieLens

1 Introduction

The explosive growth of online information and its overwhelming choices are an all too common experience for the modern consumer. Thus, algorithms that generate recommendations based on data collected online have significant value for internet websites and e-commerce. As such, recommendation systems have played a vital role in improving and increasing the consumer/user experience through widespread adoption of recommenders by an ever-growing number of e-businesses.

The types of online data that recommender systems generally use are explicit and/or implicit user preferences, item features, and sequence information such as temporal or spatial data. Recommender systems are also categorized based on the use of these types of input data. For example, a simple content-based recommender system which only uses item feature data (e.g., keyword search engine), or a more complex collaborative filtering recommender system, which is based on the premise that users who agree on the ratings for some items also tend to agree on the ratings for other items. As such, collaborative filtering recommender systems may use both explicit and implicit user preferences as well as temporal changes in user preferences (e.g., personalized Netflix movie recommendations).

Recommender systems have also continued to evolve with the advances in neural networks [Zhang et al. 2019; Steck et al. 2021; Wu et al. 2022]. A review of neural network based recommender systems by Zhang et al. (2019) summarizes the significant advantages of the deep learning techniques that neural networks bring to recommender systems:

"The linear assumption, acting as the basis of many traditional recommenders, is oversimplified and will greatly limit their modeling expressiveness. It is well-established that neural networks are able to approximate any continuous function with an arbitrary precision by varying the activation choices and combinations . . . This property makes it possible to deal with complex interaction patterns and precisely reflect [a] user’s preference . . . [Thus] deep learning is able to effectively capture the non-linear and non-trivial user-item relationships, and enable the codification of more complex abstractions as data representations in the higher layers. Furthermore, it catches the intricate relationships within the data itself, from abundant accessible data sources such as contextual, textual and visual information [Zhang et al. 2019, pp. 2, 6]."

⁵In this paper, " L_2 regularization" refers to the square of the Euclidean (L^2) norm, i.e., $\|\mathbf{x}\|_2^2$, see Ng, Andrew Y. (2004). In this context, "Tikhonov" refers to a λ_T parameter in the form of a Lagrangian multiplier.

The review by Zhang et al. (2019) also summarized the important limitations of using deep learning neural networks in recommender systems. The myriad of weights and activations in deep neural networks “are generally non-interpretable, limiting explainability” and lead to comparisons with a black box [Zhang et al. 2019, p. 7]. Furthermore, deep learning requires huge amounts of data in order to adequately support its extensive parameters and extensive hyperparameter tuning. On the other hand, traditional methods of modeling recommender systems still have a useful place when data is limited and/or computing power is constrained (e.g., edge computing). Furthermore, current social issues with lack of transparency in Artificial Intelligence (AI) gives traditional recommender methods an ethical advantage (e.g., Explainable Machine Learning, and Mechanistic Interpretability) over their counterpart “black box” deep neural network recommenders.

The field of recommender systems has a significant research community, and studying such systems is a part of most data science curricula [Hahsler 2022]. A study on recommender systems research by Rendle, Zhang, and Koren (2019) showed that published results in the field were often questionable unless the baseline methods were properly run on standardized benchmark datasets, such as the MovieLens 10M, where the baseline results have been well-calibrated by the research community. “For those performing algorithmic research, much of the power of these datasets lies in the ability to compare results with prior research papers or with other algorithms [Harper and Konstan 2015, p. 15].” In addition, machine learning competitions, such as Kaggle⁶ or the annual KDDCup⁷, can also serve as standardized benchmarks with well-calibrated results [Rendle, Zhang, and Koren 2019, p. 12].

This report describes research on a bias adjusted latent factor model of collaborative filtering for a recommender system, and the key steps in development of the algorithms in the model. The data used to develop the model was the MovieLens 10M, which is a stable benchmark dataset with 10 million movie ratings that has been used by industry, education, and research in the field of recommendation systems for more than a decade [Harper and Konstan 2015]. The MovieLens 10M dataset is maintained by GroupLens Research⁸ at the University of Minnesota.

2 Datasets

The MovieLens 10M Dataset was downloaded from the GroupLens website at <https://grouplens.org/datasets/movielens/10m/>. The 10M data at the site are divided between three files: movies.dat, ratings.dat, and tags.dat; however, the tags.dat data was not used in the model for this report. The site states that MovieLens users were selected at random for inclusion in the 10M dataset and that their IDs have been anonymized. Furthermore, the users were selected separately for inclusion in the ratings.dat and tags.dat data sets. Therefore, some user IDs that were randomly selected for one set (ratings or tags) were not randomly selected for the other set (tags or ratings) purely by chance. The total number of unique users across both of the ratings.dat and tags.dat data sets are represented by 71,567 anonymized user IDs, however 1,689 of the user IDs that were in the tags.dat data, were not also in the rating.dat data and, therefore, were not represented in the model for this report. Hence, the recommender model of this report was trained on data from a total of 69,878 anonymized users.

The MovieLens 10M site states that all of the movie Ratings are contained in the ratings.dat file, where each line of the file represents one rating of one movie by one user with the following data format: UserID::MovieID::Rating::Timestamp. The lines within this file are ordered first by UserID, and then ordered (within a UserID) by the MovieID(s). The users’ Ratings were scored on a 5-star scale, with half-star increments, that include the following possible values: 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, and 5.0. The Timestamp for each movie rating represents the number of seconds since midnight of January 1, 1970, in Coordinated Universal Time (UTC). The total time period covered by all movie ratings in the ratings.dat file is from January 9, 1995 to January 5, 2009.

The MovieLens 10M site also states that the Movie information is contained in the movies.dat file, where each line of the file represents one movie with the following data format: MovieID::Title::Genres. Movie

⁶<https://www.kaggle.com>

⁷<https://kdd.org>

⁸<https://grouplens.org/datasets/movielens/10m/>

titles were entered identically to those found in the IMDb,⁹ including year of release, however they were entered manually and consequently errors or inconsistencies may exist. The genres listed, in the file line for each movie, are a pipe-separated list of one or more selections from the following 20 individual categories:

- Action
- Adventure
- Animation
- Children's
- Comedy
- Crime
- Documentary
- Drama
- Fantasy
- Film-Noir
- Horror
- IMAX
- Musical
- Mystery
- Romance
- Sci-Fi
- Thriller
- War
- Western
- (no genres listed)

Most of the information in the paragraphs above are from the README.txt file¹⁰ for the MovieLens 10M Dataset at the GroupLens site. However, three apparent discrepancies with the README information should be explained. First of all, the README.txt states that this data set contains 10,000,054 ratings and 95,580 tags applied to 10,681 movies by 71,567 users of the online movie recommender service MovieLens. As previously stated, this report does not use the tags.dat data and so 1,689 of the user IDs that were in the tags.dat data, that were not also in the rating.dat data, were not included in the model for this report. Second of all, only four of the movie IDs in the tags.dat data were not also in the rating.dat data, and thus were not included in the model for this report. Third of all, the README.txt states that all users had rated at least 20 movies. Indeed this is the case, however when the 10,000,054 rows of rating data gets *partitioned* between the training and the final-holdout-test datasets (see below), then the minimum number of ratings reported by each user in the *training* dataset drops to at least 10 movies.

The complete starting dataset in the model for this report contained 10,000,054 rows of data with a rating and timestamp in each row applied to one of 10,677 movies (i.e., movie ID, title and year of release, and genres) by one of 69,878 anonymized users (i.e., user ID).¹¹ A 10% *final-holdout-test* partition (a.k.a., holdout dataset) of the 10,000,054 rows of starting data contains a 999,999 row subset of the data that was randomly selected without replacement. The holdout partition also represents a 9,809 subset of the 10,677 movies and represents a 68,534 subset of the 69,878 users. A 10% *testing* partition (a.k.a., edx_test dataset) of the 9,000,055 remaining rows of starting data also contains a 999,999 row subset of the remaining data that was randomly selected without replacement. The testing partition also represents a 9,779 subset of the 10,677 movies and represents a 68,548 subset of the 69,878 users. The 80% *training* partition (a.k.a., edx_train dataset) contains the remaining 8,000,056 rows of starting data that were left after the randomized removal of the holdout and testing partitions. The training partition also represents all 10,677 movies and all 69,878 users. See Table 1 for a summary of this information.

Table 1: Proportions of MovieLens 10M Data used in the Training Partition (edx_train Dataset), Test Partition (edx_test Dataset), and Final-Holdout-Test Partition (Holdout Dataset).

MovieLens 10M Data (movies.dat, ratings.dat)		
Training Partition 80% Data	Test 10%	Holdout 10%
8,000,056 Ratings	999,999	999,999
69,878 Users	68,548	68,534
10,677 Movies	9,779	9,809
797 Movie Genres	765	773

⁹Internet Movie Database (IMDb) <https://www.imdb.com/>

¹⁰<https://files.grouplens.org/datasets/movielens/ml-10m-README.html>

¹¹The complete starting dataset in the model for this report is sometimes referred to as the *edx* dataset.

2.1 Genres in Datasets

For clarity and modeling purposes, this paper makes a distinction between movie genre/genreId and genre element/genre category, where movie genre is synonymous with genreId and genre element is synonymous with genre category, such that one or more genre elements/categories are combined to create each unique movie genre/genreId. Genre IDs were not part of the MovieLens 10M Dataset that was downloaded from the GroupLens website. Analysis of the dataset revealed a maximum total of 797 unique movie genre IDs that were identified and labeled, as applicable, in the edx_train, edx_test, and holdout datasets (see Table 1 for summary of Movie Genre results). The purpose was to facilitate leveraging of the predictive power of global biases for the various movie genres as well as individual user biases for the various movie genres.

Table 2: Complete List and Count of Genre-Elements/Individual Categories from All Movies in the MovieLens (edx) Data

	Genre-Element	Count
1	Drama	3910127
2	Comedy	3540930
3	Action	2560545
4	Thriller	2325899
5	Adventure	1908892
6	Romance	1712100
7	Sci-Fi	1341183
8	Crime	1327715
9	Fantasy	925637
10	Children	737994
11	Horror	691485
12	Mystery	568332
13	War	511147
14	Animation	467168
15	Musical	433080
16	Western	189394
17	Film-Noir	118541
18	Documentary	93066
19	IMAX	8181
20	(no genre listed)	7

Table 3: Partial List of the 797 Unique Combinations of Genre Elements and the Corresponding Movie Genre IDs in the MovieLens (edx) Data

genreId	genres	genreId	genres
1	Comedy Romance	778	Adventure Comedy Fantasy Romance
2	Action Crime Thriller	779	Adventure Drama Romance Western
3	Action Drama Sci-Fi Thriller	780	Adventure Comedy Drama Fantasy Mystery Sci-Fi
4	Action Adventure Sci-Fi	781	Adventure Horror Romance Sci-Fi
5	Action Adventure Drama Sci-Fi	782	Thriller Western
6	Children Comedy Fantasy	783	Action Comedy Romance Western
7	Comedy Drama Romance War	784	Adventure Fantasy Mystery
8	Adventure Children Romance	785	Crime Drama Film-Noir Romance
9	Adventure Animation Children Drama Musical	786	Action Romance Western
10	Action Comedy	787	Crime Documentary
11	Action Romance Thriller	788	Fantasy Mystery Western
12	Action Comedy Crime Thriller	789	Comedy Drama Fantasy Sci-Fi
13	Action Comedy War	790	Animation IMAX Sci-Fi
14	Comedy	791	Drama Musical Thriller
15	Comedy Drama Romance	792	Adventure Fantasy Film-Noir Mystery Sci-Fi
16	Adventure Animation Children Comedy Musical	793	Animation Documentary War
17	Action Sci-Fi	794	Adventure Animation Musical Sci-Fi
18	Animation Children Drama Fantasy Musical	795	Fantasy Mystery Sci-Fi War
19	Animation Children	796	Action War Western
20	Action Drama War	797	Adventure Mystery

There are 20 individual categories of genres in the MovieLens data, from which 797 different subsets of the 20 categories are combined as the genre elements for the 797 unique movie genres in the MovieLens (edx) data of this report (see Table 2 and Table 3). For example, movie genre $g = 1$ (genreId 1, see Table 3) in the MovieLens data is composed of genre categories “Comedy” and “Romance”, which are the two individual genre elements in the subset $k \in \{2, 6\}$ that came from the complete set of 20 genre elements in the MovieLens data (see Table 2 and also see Equations 4.0, 4.1, 4.2). In addition, this paper uses the

terms “genre” and “genres” when the meanings are clear from the context. However, in Table 3, the term “genres” has a double meaning as it is used to refer to both the *individual genres* and their combination as genre elements in the *movie genres* listed as genreIds.

2.2 Time Bins in Datasets

The *timestamp* data associated with each rating in the MovieLens 10 M dataset was used to determine the time bin that each rating of a movie fell into. Note, however, that the timestamp data was collected from the computer system at the instant the user entered a movie rating into the MovieLens database. As Harper and Konstan (2015) point out in *The MovieLens Datasets: History and Context* (p. 15), the timestamps do not necessarily represent the date of a movie’s consumption and that users often entered a large number of movie ratings in a single online session. Thus, ratings in MovieLens can happen years after watching a movie. Table 4 shows samples of the Training partition: the *edx_train* dataset as well as the *edx_train_date* version of the dataset with Timestamp, Date, UTC Time, and 2-Year Time Bin data included. Each 2-year time bin corresponds to 24 months of data, thus seven time bins span all the data timestamps in the MovieLens 10M dataset from start of 1995-01-09 11:46:49 UTC to finish of 2009-01-05 05:02:16 UTC – see Table 5 for a list of the Date & UTC Time Range of each of the seven 2-Year Time Bins in the MovieLens 10M dataset.

Table 4: Representative Sample of Training Partition *edx_train* Dataset, and the *edx_train_date* Version of the Dataset with Date & UTC Time Stamps and 2-Year Time Bin Data Included
(Table column headers were edited for clarity.)

edx_train					Title & Year of Release			Genres		GenreID
Row	UserID	MovieID	Rating	Timestamp						
1:	1	122	5.0	838985046		Boomerang (1992)		Comedy Romance	1	
2:	1	292	5.0	838983421		Outbreak (1995)	Action Drama Sci-Fi Thriller		3	
3:	1	316	5.0	838983392		Stargate (1994)		Action Adventure Sci-Fi	4	
4:	1	329	5.0	838983392		Star Trek: Generations (1994)	Action Adventure Drama Sci-Fi		5	
5:	1	355	5.0	838984474		Flintstones, The (1994)		Children Comedy Fantasy	6	

8000052:	59269	64903	3.5	1230162521		Nazis Strike, The (Why We Fight, 2) (1943)		Documentary War	513	
8000053:	59342	61768	0.5	1230070861		Accused (Anklaget) (2005)		Drama	34	
8000054:	63134	54318	2.5	1222631928	Cruel Story of Youth (Seishun zankoku monogatari) (1960)			Drama	34	
8000055:	65308	64611	3.5	1230097394		Forgotten One, The (1990)		Thriller	81	
8000056:	67385	7537	2.5	1188277406		Du côté de la côte (1958)		Documentary	128	

edx_train_date								
Row	UserID	GenreID	MovieID	Rating	Timestamp	Date	UTC Time	2Yr Timebin
1:	1	1	122	5.0	838985046	1996-08-02	11:24:06	1996-01-01
2:	1	3	292	5.0	838983421	1996-08-02	10:57:01	1996-01-01
3:	1	4	316	5.0	838983392	1996-08-02	10:56:32	1996-01-01
4:	1	5	329	5.0	838983392	1996-08-02	10:56:32	1996-01-01
5:	1	6	355	5.0	838984474	1996-08-02	11:14:34	1996-01-01

8000052:	59269	513	64903	3.5	1230162521	2008-12-24	23:48:41	2008-01-01
8000053:	59342	34	61768	0.5	1230070861	2008-12-23	22:21:01	2008-01-01
8000054:	63134	34	54318	2.5	1222631928	2008-09-28	19:58:48	2008-01-01
8000055:	65308	81	64611	3.5	1230097394	2008-12-24	05:43:14	2008-01-01
8000056:	67385	128	7537	2.5	1188277406	2007-08-28	05:03:26	2008-01-01

Table 5: Date & UTC Time Ranges of 2-Year Time Bins in the MovieLens 10M Dataset

2yr Timebin	Start Date and Time		End Date and Time	
1: 1996-01-01	1995-01-09	11:46:49 UTC	1997-01-09	11:46:48 UTC
2: 1998-01-01	1997-01-09	11:46:49 UTC	1999-01-09	11:46:48 UTC
3: 2000-01-01	1999-01-09	11:46:49 UTC	2001-01-09	11:46:48 UTC
4: 2002-01-01	2001-01-09	11:46:49 UTC	2003-01-09	11:46:48 UTC
5: 2004-01-01	2003-01-09	11:46:49 UTC	2005-01-09	11:46:48 UTC
6: 2006-01-01	2005-01-09	11:46:49 UTC	2007-01-09	11:46:48 UTC
7: 2008-01-01	2007-01-09	11:46:49 UTC	2009-01-09	11:46:48 UTC

3 Methods & Analysis

3.1 Modeling Approach – Overview

From a high level perspective the model is, overall, a linear algorithm that consists of an overall mean rating and a summed series of feature biases followed by a summed series of latent factors. The rationale for this overarching model design was best explained by Koren, Bell, and Volinsky (2009) in their paper on “Matrix Factorization Techniques for Recommender Systems.” They explain that much of the observed variation in rating values is due to biases that are independent of any latent factor interactions between users and items (movies). Thus it is best to separate out the portion of rating values due to biases and then subject the “true interaction portion of the data to factor modeling” [Koren, Bell, and Volinsky 2009, p. 45]. Finally, in contrast to the overall linear design of the model, non-linear scaling was used within the genre feature bias to constrain variability in the TRLIM output of *interaction bias* between users and the movie items.

More specifically, the model starts by building on a regularized baseline predictor algorithm that accounts for the static biases,¹² in movie ratings and user ratings, that are centered on the overall average rating of the dataset [Koren 2009; Irizarry 2022]. Next the model incorporates a temporal function in the movie bias algorithm that is similar to the time-bin approach proposed by Koren (2009). Addition of the temporal movie bias to the static movie bias renders the movie bias algorithm as “time-aware” [Koren 2009]. Details and the equations of the time-aware movie bias algorithm are discussed in Section 3.3.

As Figure 1 shows, there is strong evidence of a rating bias due to the genre category of a movie [Irizarry 2022]. Thus, the model also incorporates a genre bias algorithm that consists of two major parts. The first part is a regularized baseline predictor for the static biases in genre effects on the ratings, that is centered on the overall average rating of the dataset. The second part is a function of the user that incorporates an optimized predictor of the user-specific genre biases, which renders the genre bias algorithm as user-aware.




Figure 1: **Strong Evidence of Movie-Genre Bias based on Average & Standard Error of Movie Ratings within Genre ID Category (with 30 or more Ratings).** For clarity, Genre ID numbers not displayed.

¹²The *static biases* in this paper are determined by the training data and are treated as fixed properties throughout development of the model. The static biases are also globally applied in the sense that the bias of a *particular* movie (i.e., movieId), user (i.e., userId) or genre (i.e., genreId) is a unique and fixed property of that particular feature ID and affects all rating calculations it is a part of. Nonetheless, in a real world application, these so called “static biases” would be periodically updated as part of a model maintenance program when a threshold of new data had accumulated.

Adding the user function to the genre bias algorithm required a Tikhonov regularized ridge regression approach to reduce multicollinearity and condition number in a non-invertible system of linear equations. The non-invertible system of linear equations arose from the non-square matrix of genre elements that got extracted from a user’s movie ratings, in which the rows or the columns of the user’s matrix of genre elements were not necessarily linearly independent. Ridge regression gets around these non-invertible issues by using the pseudo-inverse of the user’s matrix of genre elements along with a Tikhonov regularization parameter that adds a “ridge” of positive elements on the diagonal of the pseudo-inverse matrix. This ridge is an example of Bias-Variance Tradeoff, where Tikhonov regularization by ridge regression decreases condition number and multicollinearity, while increasing numerical stability and positive definiteness, which makes for a well-posed problem with a well-conditioned solution in exchange for a tolerable amount of ridge bias. This Tikhonov regularized ridge regression approach was analytically derived (see Appendix C in Section 6.3) into a Tikhonov regularized left-inverse matrix equation (TRLIM Algorithm) that produces an optimized solution vector for the user-specific genre-element biases [Aggarwal 2020, pp. 79-85, 176-179; Deisenroth et al. 2023, pp. 85-88, 300-303, 313-314]. Furthermore, the user-specific genre-element biases were non-linearly transformed to real numbers \mathbb{R} within the interval $[-1, 1]$ by a hyperbolic-tangent “squashing” function, and then weighted by a tuned proportionality factor, which constrained variability and further reduced RMSE. Details and equations of the user-aware genre bias algorithm are discussed in Section 3.5.

Finally, the residuals of the rating data (i.e., remnants of the rating values that were left after separating out the rating biases of the previously described algorithms in the model) were collaboratively filtered by matrix factorization using the Singular Value Decomposition method of Principal Component Analysis (SVD-PCA). More specifically, the residuals of the rating data were SVD-PCA factorized by decomposing an imputed¹³ dense matrix of high-dimensional¹⁴ user rows and movie columns into a joint user-movie latent factor space as a form of collaborative filtering [Koren 2009; Irizarry 2022]. Latent factor filtering by SVD-PCA matrix factorization can be thought of as collaborative filtering in the sense that the user and movie “covariances” are hierarchically ordered by the strength of their loadings in the singular vectors (latent factor eigenvectors) of the SVD symmetric positive-semidefinite RR^T row-space (user “covariance”) matrix and of the SVD symmetric positive-semidefinite $R^T R$ column-space (movie “covariance”) matrix.¹⁵

A Principal Components Regression (PCR) approach was used to truncate the SVD-PCA principal component results and provide the best possible denoised low-rank approximation of the latent factors in terms of minimized squared error [Gareth et al. 2021, pp. 256-259; Aggarwal 2020, p. 318]. Out of the 10,677 principal component singular vectors in the full-rank *economy*¹⁶ SVD-PCA of the user-movie training dataset, truncation down to a level of only the first 106 principal component singular vectors, for use in calculation of the latent factors in the model, was found to denoise and minimize the RMSE of the test dataset predictions (see Figure 2).¹⁷ Details and equations of the SVD-PCA matrix factorization latent-factor collaborative filtering algorithm are discussed in Section 3.6.

The model was trained on a random partition of 80% of the MovieLens 10M dataset, with the remaining 20% randomly divided equally between the test partition and the final-holdout-test partition of the dataset.

¹³The 98.9% empty (sparse) matrix was imputed by Low-Rank Matrix Completion (LRMC). See Appendix A for details.

¹⁴SVD generalizes the factorization approach to any size matrix, rectangular or square, and is preferred for use with high dimensional data [Aggarwal 2020, pp. 300, 323].

¹⁵“Covariance” is in quotes because, technically, when the matrix R of residuals (with zero-mean centered data) is multiplied by its transpose R^T , the result is *proportional* to the actual User Covariance Matrix of R : $C_{u_R} = \frac{1}{n}RR^T$ where n is the number of columns (length of user row) in matrix R . Likewise, $R^T R$ is *proportional* to the actual Movie Covariance Matrix of R : $C_{i_R} = \frac{1}{m}R^T R$ where m is the number of rows (length of movie column) in matrix R . Nonetheless, the left and right singular vectors of the respective SVD row-space matrix RR^T and column-space matrix $R^T R$, accurately represent the proportional strengths within the respective user and movie covariances. Proportional covariances which therefore “collaborate” in hierarchically selecting factors based on the collective weight of their loadings for the variables represented in the respective principal component singular vectors (latent factor eigenvectors).

¹⁶The *economy* SVD has a rank $r = \min\{m, n\}$ of the m rows and n columns of the matrix R [Aggarwal 2020, p. 306].

¹⁷As more principal components are used in the regression, the bias decreases while the variance increases, which results in a typical U-shaped curve for the RMSE [Gareth et al. 2021, p. 258]. Due to the training time and complexity of SVD-PCA with LRMC Imputation, an acceptably optimized Truncation Level of the SVD-PCA was determined by PCR evaluation with RMSE from the (one) Test Dataset, as opposed to PCR Cross Validation with multiple training & test partitions of the MovieLens 10M dataset sans final-holdout-test partition.




Figure 2: **Tuning the Level of Principal Component Singular Vectors in Truncated SVD-PCA by Principal Components Regression (PCR) to Denoise the RMSE of Test Data Predictions.**

The accuracy of the model during development was evaluated by the RMSE measured between the rating predictions of the model and the actual ratings in the test partition (10%) of the MovieLens 10M dataset. The accuracy of the *final* model was measured to have an RMSE of 0.792085, as determined by the final model rating predictions evaluated against the actual ratings in the final-holdout-test partition (10%) of the MovieLens 10M dataset.¹⁸ Refer to Section 2 Datasets for more details on the partitions of the MovieLens 10M dataset.

3.2 Complete Model

The following two mathematical equations, Equation 1.0 and Equation 1.1, represent the complete model and algorithms for the predicted ratings by this recommendation system:

$$\hat{r}_{ui} = \mu + b_i(t) + b_u + b_g(u) + \left(\vec{p}_u^T \bullet \vec{q}_i \right)_f \quad (1.0)$$

where \hat{r}_{ui} denotes the predicted rating of user u for movie-item i . μ , denotes the overall mean rating of the dataset, $b_i(t)$ denotes the movie bias in its rating which includes a function of time t , b_u denotes the user bias in the movie rating, and $b_g(u)$ denotes the genre bias in the movie rating which includes a function of the user u . The dot product $\left(\vec{p}_u^T \bullet \vec{q}_i \right)_f$ denotes the total value of the imputed¹⁹ joint user-movie latent factors f from the inner product of the row vector \vec{p}_u^T of user u data in the SVD-PCA matrix $P\Sigma$ and the column vector \vec{q}_i of movie i data in the SVD-PCA matrix Q^T . The noise-error of the model was minimized when the lengths of the \vec{p}_u^T and \vec{q}_i vectors were truncated to the first $f = 106$ principal component latent factors from the respective left and right singular vectors of the SVD-PCA.

For further details and equations on each of the terms in Equation 1.0, see sections 3.3 Time-aware Movie Bias Algorithm, 3.4 User Bias Algorithm, 3.5 User-aware Genre Bias Algorithm, 3.6 SVD-PCA Latent Factor Collaborative Filtering Algorithm, and 6.1 Appendix A for Imputation by Low-Rank Matrix Completion.

¹⁸The source code of the model is publicly available at <https://github.com/jpoconnor1961/MovieLens10M.git>

¹⁹This involved a 98.9% empty (sparse) rating-residuals matrix that was imputed by Low-Rank Matrix Completion (LRMC) prior to factorization by SVD-PCA, because the SVD-PCA technique requires a complete matrix without any missing values, for accurate decomposition. See Section 6.1 Appendix A for details, equations and R code for Imputation by LRMC.

Approximately 0.4% of the results from Equation 1.0 fell outside of the 0.5 to 5.0 range of rating values in the MovieLens data set.²⁰ Predicted ratings above 5.0 or below 0.5 are expected due to the 0.7921 RMSE of the model, however these errors were reduced by min-max functions in the final step of the model:

$$\hat{r}'_{ui} = \max\{\min\{\hat{r}_{ui}, 5.0\}, 0.5\} \quad (1.1)$$

where \hat{r}_{ui} values above 5.0 were minimized (i.e., clipped) to \hat{r}'_{ui} values of 5.0, \hat{r}_{ui} values below 0.5 were maximized (i.e., clipped) to \hat{r}'_{ui} values of 0.5, and all other \hat{r}_{ui} values were equivalent to \hat{r}'_{ui} values.²¹ This final step of the model, Equation 1.1, improved (reduced) the RMSE of the Equation 1.0 results by about 0.04% RMSE (see Table 11 in Section 4 Results and Discussion for more details).

3.3 Time-aware Movie Bias Algorithm

The time-aware movie bias algorithm $b_i(t)$ includes a function of time t , such that the bias of movie i is split into a static part b_i and a time-based part $b_{i,t}$ that renders the movie bias algorithm as “time-aware” [Koren 2009, p. 450]:

$$b_i(t) = b_i + b_{i,t} \quad (2.0)$$

The static movie bias b_i of movie i is calculated as the regularized-average of all the deviations of the movie’s ratings r_i from the overall mean rating μ of all movies in the dataset:

$$b_i = \frac{\sum_1^{n_i} (r_i - \mu)}{n_i + \lambda} \quad (2.1)$$

where n_i denotes the number of ratings for movie i in the dataset and λ denotes the regularization parameter lambda, which is tuned to minimize RMSE and constrain variability in b_i .²²

The time-based part of the movie bias $b_{i,t}$, for movie i , is calculated as the difference between the regularized-mean rating of the movie, which occurred in a particular time bin t of the dataset, and the overall mean rating of the same movie μ_i in the entire dataset:

$$b_{i,t} = \frac{\sum_1^{n_{i,t}} r_{i,t}}{n_{i,t} + \lambda_t} - \mu_i \quad (2.2)$$

where $r_{i,t}$ denotes a rating of movie i in time bin t ,²³ $n_{i,t}$ denotes the number of ratings for movie i in time bin t , and λ_t denotes the regularization parameter lambda for the mean rating of the movie in time bin t , where λ_t is tuned to minimize test dataset RMSE and constrain variability in $b_{i,t}$. Thus, Equation 2.2 requires the sequential optimization of two parameters – first, the length of time spanned by a time bin t , and second, the strength of λ_t in regularizing movies with low rating frequencies in the optimized span of time bin t .

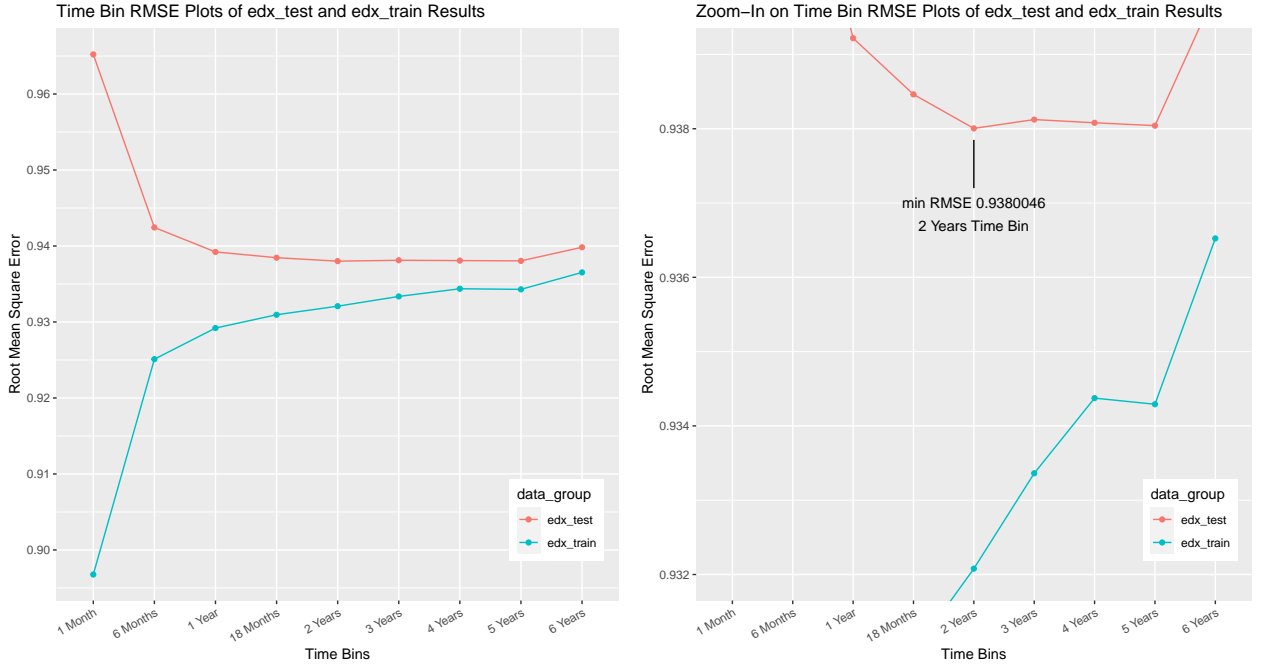
A range of time-bins from 1-month to 6-years were tested, and the 2-year time-bins produced the lowest RMSE based on the test partition of the dataset (see Figures 3a and 3b). Since each 2-year time bin corresponds to 24 months of data, there are 7 time bins spanning all the date & time stamps in the entire MovieLens 10M dataset (1995-01-09 11:46:49 UTC to 2009-01-05 05:02:16 UTC). See Section 2.2 Time Bins within Section 2 Datasets for details on date ranges of the seven 2-year time bins.

²⁰The users’ MovieLens Ratings were scored on a 5-star scale, with half-star increments, that included the following possible values: 0.5, 1.0, 1.5, 2.0, 2.5, 3.0, 3.5, 4.0, 4.5, and 5.0. For more information, see Section 2 Datasets.

²¹This min-max range step for error reduction is also known as “clipping” by other authors, see Paterek (2007) and Webb (2006).

²²See Footnote #26.

²³The *timestamp* data associated with each rating in the MovieLens 10 M dataset was used to determine the time bin that each rating of a movie fell into. Note, however, that the timestamp data was collected from the computer system at the instant the user entered a movie rating into the MovieLens database. As Harper and Konstan (2015) point out in *The MovieLens Datasets: History and Context* (p. 15), the timestamps do not necessarily represent the date of a movie’s consumption and that users often entered a large number of movie ratings in a single online session. Thus, ratings in MovieLens can happen years after watching a movie, which may partially explain why the optimum size time bin of this model turned out to be 2 years.



(a) Overfitting with Shorter Time Bins in edx_train. (b) Minimum RMSE with 2 Year Time Bin in edx_test.

Figure 3: (a) Time Bin RMSE Plots of edx_test and edx_train Results. RMSE Results Show Overfitting at Shorter Time Bins in edx_train. (b) Zoom-In on Time Bin RMSE Plots of edx_test and edx_train Results. RMSE Results Show Minimum at 2 Years Time Bin in edx_test.

The regularization parameter λ_t in Equation 2.2 was tuned to minimize the RMSE of the time-based part $b_{i,t}$ of the movie bias algorithm $b_i(t)$ with the time bin t set at the optimal 2-year size. However, to properly calculate RMSE, the movie bias prediction model $\hat{r}_i = \mu + b_i(t)$ was based on the full time-aware movie bias algorithm $b_i(t) = b_i + b_{i,t}$, where the regularized b_i component was treated as a constant across the range of λ_t tuning values tested in the $b_{i,t}$ component. Hence, a 30-fold boot-strap cross validation was performed to test a range of λ_t values for the minimum RMSE on the training data set (edx_train).²⁴ Oddly, the minimum RMSE result for each of the 30 cross validations was found at $\lambda_t = 0$. Nonetheless, the average of these minimum RMSE results from the 30 cross validations was an edx_train mean RMSE = 0.93602 with a standard deviation = 0.00004. Figure 4a shows RMSE and λ_t data from one of the 30 cross validations, with a minimum RMSE of 0.93606 at $\lambda_t = 0$. Further investigation was conducted to explain these results.

Further investigation found no issues with the logic, math, or coding that produced these consistent but curious $\lambda_t = 0$ results. Thus, the next step of logic was to conclude that the results were valid within the limits of their precision, which was the step size of 0.5 increments across the -10 to $+10$ range of λ_t values tested. This conclusion also suggested that the rounded value of the *true* λ_t was closer to zero than it was to either $+0.5$ or -0.5 . Thus, the next logical step was to search for λ_t in many tiny increments across a small range where zero was at the center. A 30-fold boot-strap cross validation, however, is a semi-empirical approach that is not very efficient for searching across a high number of very small increments. Thus, the approach for finding λ_t was changed from a computationally expensive semi-empirical method to a much more efficient analytical solution.

The analytical approach taken was to derive the *partial derivative* of the cost function with respect to λ_t for the model up to that point, i.e., $\hat{r}_i = \mu + b_i(t)$, and see which incremental value of λ_t close to zero produced no more than a ± 0.00001 slope in the cost function on the Training dataset. Figure 4b shows the Mean

²⁴The 30-fold boot-strap cross validation of λ_t for time-based part $b_{i,t}$ of the movie bias (Equation 2.2) was separate from another 30-fold boot-strap cross validation of the regularization parameter λ for the movie, user and genre static biases b_i , b_u and b_g respectively (see Footnote #26).

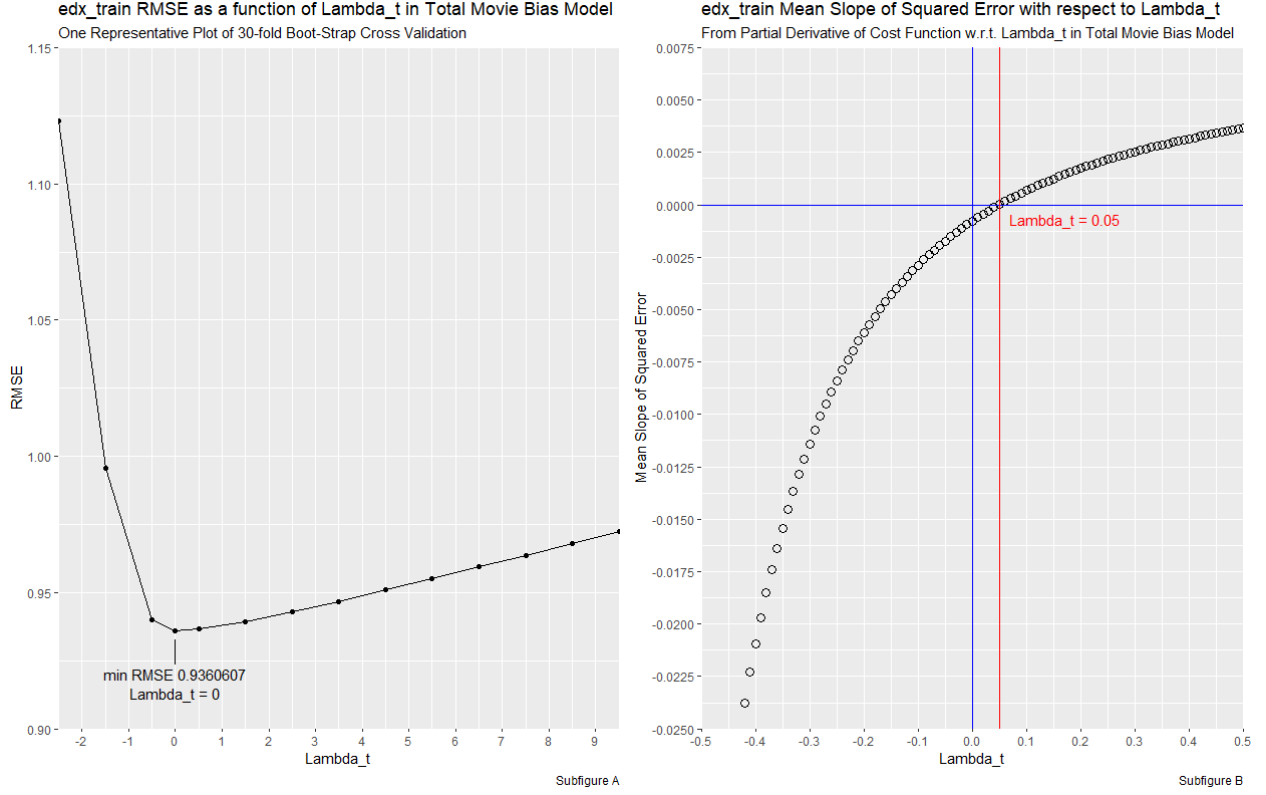


Figure 4: Finding λ_t by Inefficient Semi-Empirical Approach A vs Efficient Analytical Solution B
Subfigure A: Tuning Lambda Regularization Parameter λ_t of Movie Temporal Bias $b_{i,t}$ to Minimize RMSE of Training Data Predictions. (Zoom-In of the Minimum RMSE Region is Shown)
Subfigure B: Tuning Lambda Regularization Parameter λ_t of Movie Temporal Bias $b_{i,t}$ to Zero Mean Slope for Partial Derivative of Squared Error Function, with respect to λ_t , using Training dataset. (See Appendix B for details, equations, derivatives, and R code.)

Slope of Squared Error calculated for the edx_train dataset in 0.01 increments of λ_t , across a range of -0.50 to $+0.50$ λ_t , with the derivative of the cost function crossing the zero slope (within ± 0.00001) when closest to $\lambda_t = 0.05$, which is the final result for this parameter. Once setup, the computations took less than 5 minutes to complete, see Appendix B in Section 6.2 for more details, equations, derivatives, and R code.

3.4 User Bias Algorithm

The user bias b_u in movie ratings was calculated as the regularized-average of all the deviations of a user's ratings from the respective time-aware movie biases and the overall mean rating of all movies in the dataset:

$$b_u = \frac{\sum_1^{n_u} (r_{u,i} - \mu - b_i(t))}{n_u + \lambda} \quad (3.0)$$

where $r_{u,i}$ denotes the rating by user u of movie i , μ denotes the overall mean rating of all movies in the dataset, $b_i(t)$ denotes the time-aware bias of movie i , n_u denotes the number of movie ratings by user u in the dataset, and λ denotes the regularization parameter lambda, which is tuned to minimize RMSE and constrain variability in b_u .²⁵

²⁵See Footnote #26.

3.5 User-aware Genre Bias Algorithm

The user-aware genre bias algorithm $b_g(u)$ estimates the rating bias of movie-genre g that includes a function of the user u , such that $b_g(u)$ consists of the sum of a static movie-genre bias term, b_g , and a user-specific movie-genre bias term, $b_{g,u}$:

$$b_g(u) = b_g + b_{g,u} \quad (4.0)$$

The static movie-genre bias, b_g , was calculated as the regularized-average of all the deviations of a genre's (i.e., genreId's) movie ratings from the respective time-aware movie biases $b_i(t)$, the respective user biases b_u , and the overall mean rating μ of all movies in the dataset:

$$b_g = \frac{\sum_1^{n_g} (r_{u,i,g} - \mu - b_i(t) - b_u)}{n_g + \lambda} \quad (4.1)$$

where $r_{u,i,g}$ denotes the rating by user u of movie i which is in a particular movie-genre g (i.e., genreId). The variable n_g denotes the number of movie ratings for movie-genre g (genreId) in the dataset, and λ denotes the regularization parameter lambda, which is tuned to minimize RMSE and constrain variability in b_g .²⁶

The user-specific movie-genre bias $b_{g,u}$ involves an *optimization approach* to solve an *ill-conditioned* system of linear equations in order to parse user-specific biases for each of the genre-elements²⁷ contained in a user's data of rated movies in the training dataset (see Table 6 for an example of a user's data). In addition, an *indirect approach* was necessary to *infer estimates* of unknown user-specific movie-genre biases from the appropriate genre-element components, because many of the movie genres in ratings by specific users in the test dataset were not rated by the same users in the training dataset. For example, userId 1 has rated a movie with genreId 2 in the test dataset that is composed of genre-elements "Action", "Crime" and "Thriller". However, userId 1 has not rated a movie with genreId 2 in the training dataset (see Table 6). Thus, it is necessary to infer an estimate of userId 1 bias for genreId 2 from the genre-elements "Action", "Crime" and "Thriller" that are *also components in genre IDs for which the user has movie ratings in the training dataset*, i.e., genreId 3 (Action, Drama, Sci-Fi, Thriller), genreId 4 (Action, Adventure, Sci-Fi), genreId 5 (Action, Adventure, Drama, Sci-Fi), genreId 10 (Action, Comedy), genreId 11 (Action, Romance, Thriller), genreId 12 (Action, Comedy, Crime, Thriller), genreId 13 (Action, Comedy, War), and genreId 17 (Action, Sci-Fi); see Table 6.

Table 6: UserId 1 Data in the Training Dataset

row	userId	movieId	rating	timestamp	title	genres	genreId
1	1	122	5	838985046	Boomerang (1992)	Comedy Romance	1
2	1	292	5	838983421	Outbreak (1995)	Action Drama Sci-Fi Thriller	3
3	1	316	5	838983392	Stargate (1994)	Action Adventure Sci-Fi	4
4	1	329	5	838983392	Star Trek: Generations (1994)	Action Adventure Drama Sci-Fi	5
5	1	355	5	838984474	Flintstones, The (1994)	Children Comedy Fantasy	6
6	1	356	5	838983653	Forrest Gump (1994)	Comedy Drama Romance War	7
7	1	362	5	838984885	Jungle Book, The (1994)	Adventure Children Romance	8
8	1	364	5	838983707	Lion King, The (1994)	Adventure Animation Children Drama Musical	9
9	1	370	5	838984596	Naked Gun 33 1/3: The Final Insult (1994)	Action Comedy	10
10	1	377	5	838983834	Speed (1994)	Action Romance Thriller	11
11	1	420	5	838983834	Beverly Hills Cop III (1994)	Action Comedy Crime Thriller	12
12	1	466	5	838984679	Hot Shots! Part Deux (1993)	Action Comedy War	13
13	1	539	5	838984068	Sleepless in Seattle (1993)	Comedy Drama Romance	15
14	1	589	5	838983778	Terminator 2: Judgment Day (1991)	Action Sci-Fi	17
15	1	594	5	838984679	Snow White and the Seven Dwarfs (1937)	Animation Children Drama Fantasy Musical	18
16	1	616	5	838984941	Aristocats, The (1970)	Animation Children	19

²⁶Optimal value of the lambda λ regularization parameter was determined by a 30-fold boot-strap cross validation of the training dataset, edx_train, across a range of λ values in a simultaneous calculation of the regularization parameter λ for the movie, user and genre static biases b_i , b_u and b_g respectively. The minimum RMSE result from each of the 30 cross validations was found at an average lambda value of $\lambda = 3.973333$, which is the optimized lambda value for the movie b_i , user b_u and genre b_g static biases (see Equations 2.1, 3.0, and 4.1).

²⁷See section 2.1 Genres in Datasets for definitions of genre-elements, genres, genreId, movie-genre, etc.

Table 7: UserId 1 Matrix A of Genre Elements, Vector \vec{x} of User’s Genre-Bias Elements $x_{u,j}$ and Vector \vec{y} of User’s Residual Movie-Rating Elements $y_{u,i}$ in Linear Equations $A\vec{x} = \vec{y}$

row	Comedy	Romance	Action	Drama	Sci-Fi	Thriller	Adventure	Children	Fantasy	War	Animation	Musical	Crime	Bias	Xuj	Residual	Yui
[1,]	1	1	0	0	0	0	0	0	0	0	0	0	0	X11		0.50184622	
[2,]	0	0	1	1	1	1	0	0	0	0	0	0	0	X12		-0.10478826	
[3,]	0	0	1	0	1	0	1	0	0	0	0	0	0	X13		0.12583684	
[4,]	0	0	1	1	1	0	1	0	0	0	0	0	0	X14		0.03179144	
[5,]	1	0	0	0	0	0	0	1	1	0	0	0	0	X15		0.69491380	
[6,]	1	1	0	1	0	0	0	0	0	1	0	0	0	X16		-0.66032801	
[7,]	0	1	0	0	0	0	1	1	0	0	0	0	0	X17		-0.15184919	
[8,]	0	0	0	1	0	0	1	1	0	0	1	1	0	X18		-0.35159866	
[9,]	1	0	1	0	0	0	0	0	0	0	0	0	0	X19		0.29579077	
[10,]	0	1	1	0	0	1	0	0	0	0	0	0	0	X110		-0.33384881	
[11,]	1	0	1	0	0	1	0	0	0	0	0	0	1	X111		0.59643652	
[12,]	1	0	1	0	0	0	0	0	0	1	0	0	0	X112		0.42748775	
[13,]	1	1	0	1	0	0	0	0	0	0	0	0	0	X113		-0.24294248	
[14,]	0	0	1	0	1	0	0	0	0	0	0	0	0			-0.50016148	
[15,]	0	0	0	1	0	0	0	1	1	0	1	1	0			-0.28856733	
[16,]	0	0	0	0	0	0	0	1	0	0	1	0	0			-0.04001912	
Genre Element Column Totals	7	5	8	6	4	3	4	5	2	2	3	2	1				

The inferred estimates are done by solving for the best-fit solution of a system of linear equations, $A\vec{x} = \vec{y}$, to mathematically parse the user’s movie rating residuals²⁸ $y_{u,i}$ into user-specific biases $x_{u,j}$ for the genre-elements that are contained in the genre-element matrix (Table 7) of the userId 1 rated movies (Table 6). As can be seen in matrix A of Table 7, there were 8 instances of Action, 3 instances of Thriller, and 1 instance of Crime, included in a system of 16 linear equations, for which user-specific biases can be parsed in a best-fit solution from the associated residuals of movie-rating values, vector $\vec{y} = \langle y_{u,i_1}, y_{u,i_2}, y_{u,i_3}, \dots, y_{u,i_{n_u}} \rangle$, where $n_u = 16$ for userId 1. Subsequently, the inferred user-specific biases for the genre-elements “Action”, “Crime” and “Thriller” are summed together to estimate the user’s overall bias $b_{g,u}$ for genreId 2 (treated as unknown) in a movie that the user rated in the test dataset.

Calculation of the user-specific movie-genre bias, $b_{g,u}$, is described by three equations which are applied in a sequential process starting with Equation 4.4, followed by Equation 4.3b, and then Equation 4.2. From a high level perspective, the purpose of Equation 4.4 is to calculate a vector $\vec{y} = \langle y_{u,i_1}, y_{u,i_2}, y_{u,i_3}, \dots, y_{u,i_{n_u}} \rangle$ of residual rating values for a system of linear equations $A\vec{x} = \vec{y}$ where A is the matrix of binary genre counts (1 = presence, 0 = absence) from each of the movies rated by the same user. The solution vector \vec{x} in the user’s system of linear equations $A\vec{x} = \vec{y}$ consists of the inferred user-specific genre-element biases that are *actually determined* by the optimization-centric Tikhonov Regularized Left-Inverse Matrix (TRLIM) *form* of this system of linear equations (see Equation 4.3b). An optimization-centric form of linear equations is necessary because an exact solution by simply using $\vec{x} = A^{-1}\vec{y}$ is not possible for two reasons, either of which makes matrix A non-invertible: 1) matrix A is typically *rectangular* in row and column dimensions, and 2) matrix A typically contains *non-independent* linear relationships between its rows and/or columns. As alluded, the purpose of the TRLIM algorithm is to calculate all j of a specific user’s genre-element biases $x_{u,j}$ by getting around these non-invertible limitations through an optimization-centric approach (that is explained in more detail below) which finds the minimum RMSE best-fit solution vector $\vec{x} = \langle x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_e} \rangle$ ²⁹ for an otherwise unsolvable system of linear equations. Finally, a subset $\{x_{u,k_1}, x_{u,k_2}, \dots, x_{u,k_c}\}$ ³⁰ of a user’s complete set of genre-element biases $\{x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_e}\}$ are combined together in Equation 4.2 to estimate the user-specific bias for a particular movie-genre, $b_{g,u}$.

More specifically, $b_{g,u}$ is a variance-constrained optimized-predictor of the inferred user-specific movie-genre

²⁸In this paper, the term “rating residuals” is synonymous with the term “residual ratings”, and are used interchangeably.

²⁹Where $j \in \{1, 2, 3, \dots, e\}$ is the complete set of unique genre elements in a specific user’s movie ratings in the training dataset. For example, $e = 13$ for userId 1 because the algorithm finds only 13 unique genre elements contained in the 16 movies rated by userId 1 and, therefore, generates 13 columns for the unique genre elements in matrix A containing the genre-element binary counts (1 = presence, 0 = absence) from the 16 movies rated by userId 1 in the training dataset (see Tables 6 and 7).

³⁰Where $k \in \{1, 2, \dots, c\}$ is the *subset* of genre elements (from a specific user’s complete set of unique genre elements j) that make up the composition $\{1, 2, \dots, c\}$ of a particular movie’s genre ID categories, e.g., genreId 2 = {Action, Crime, Thriller}.

bias that is estimated from the w weighted summation of the genre-element bias components $x_{u,k}$ for a particular movie-genre g and user u . Prior to summation and weighting, the genre-element bias components $x_{u,k}$ are non-linearly transformed to real numbers \mathbb{R} within the interval $[-1, 1]$ by a hyperbolic-tangent squashing function, $\tanh x = \frac{e^{2x} - 1}{e^{2x} + 1}$, which also constrains the total variability of the user-specific movie-genre bias:

$$b_{g,u} = w \sum_{k=1}^{k=c} \tanh(x_{u,k}) \quad (4.2)$$

The weight $w = 0.638$ is a tuned proportionality factor that was determined by minimizing the RMSE of prediction using the test dataset (see Figure 6b). The genre-element bias component $x_{u,k}$ represents the user u bias x for genre-element k that belongs in the set of components $k \in \{1, 2, 3, \dots, c\}$ that make up the composition of movie-genre g .

The complete set of a user’s inferred user-specific biases of genre-elements $\{x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_e}\}$ were determined by the TRLIM algorithm in Equation 4.3b. The TRLIM algorithm is a ridge regression solution that was derived analytically by minimizing the gradient of the Lagrangian form of a least squares cost function, containing Tikhonov L_2 regularization of the user’s genre-element vector of biases \mathbf{x} , in an over-determined system of linear equations:³¹

$$\arg \min_{\mathbf{x}} \left\{ \|\mathbf{A}\mathbf{x} - \mathbf{y}\|_2^2 + \lambda_T \|\mathbf{x}\|_2^2 \right\} \quad (4.3a)$$

Where the loss function was minimized, with respect to a user’s genre-element biases, by finding the optimal coefficient/coordinate vector \mathbf{x} of the orthogonal projection $\pi_{\mathbf{A}}(\mathbf{y})$, which is the orthogonal projection of the vector of rating observations \mathbf{y} onto the *closest subspace vector* spanned by the genre dimensions (rows) of the *pseudo-inverse*³² of matrix \mathbf{A} [Aggarwal 2020, pp. 79-85, 176-179; Deisenroth et al. 2023, pp. 85-88, 300-303, 313-314].

The derived solution to this constrained minimization problem uses a Lagrangian multiplier, in the form of a Tikhonov regularization parameter λ_T , to *tune minimization* of the prediction RMSE of the coefficient/coordinate solution vector $\vec{x} = \langle x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_e} \rangle$, of user u specific genre-element biases $x_{u,j}$, for the set (excludes redundancies) of a user’s individual genre elements $j \in \{1, 2, 3, \dots, e\}$, contained in the columns of their matrix A of binary genre counts:

$$\vec{x}_{u,j} = (A^T A + \lambda_T I_d)^{-1} A^T \vec{y}_{u,i} \quad (4.3b)$$

The binary counts (1 = presence, 0 = absence) of genre-elements in the movie genre g , of each movie i that user u rated in the training dataset, were tallied in a row for each movie in the user’s matrix A , which usually created a non-square matrix (for example, see Tables 6 and 7). A non-square matrix does not have an inverse, thus in Equation 4.3b matrix A is multiplied by its transpose A^T to transform the information to a square matrix $A^T A$. Nonetheless, linear dependencies between rows or columns of the non-square matrix A will remain linearly dependent in the square matrix $A^T A$ as scaled transformations of the linearly dependent rows and/or columns. Linear dependencies in square matrix $A^T A$ are removed in Equation 4.3b by the addition of a square-matrix Tikhonov regularization term, $\lambda_T I_d$, where the Tikhonov lambda, λ_T , is a tuned scalar regularization parameter (see Figure 6a), and I_d is a square Identity matrix with the same dimensions d as the number of columns in the user’s matrix A .³³

The addition of even a tiny fraction of value along the main diagonal of a square non-invertible/singular matrix will force linear independence between all its rows and all its columns, making the matrix completely invertible at the cost of only a very minor increase along its main diagonal (i.e., an example of

³¹See Appendix C in Section 6.3 for the analytically derived TRLIM Algorithm from the User’s Genre-Element Bias Loss Function. “ L_2 regularization” refers to the square of the Euclidean (L^2) norm, i.e., $\|\mathbf{x}\|_2^2$, see Ng, Andrew Y. (2004).

³²The *pseudo-inverse* of matrix $A = (A^T A + \lambda_T I_d)^{-1} A^T$ and is also called: “Tikhonov regularized left-inverse of matrix A ”. For example, the pseudo-inverse of matrix A for userId 1 has 13 dimensions (rows) of genres which is the subspace dimension for the projection of the 16 dimensional vector of rating observations \mathbf{y} , by the 13 dimensional coefficient/coordinate vector \mathbf{x} .

³³In this paper $\lambda_T I_d$ is also referred to as a “lambda diagonal matrix” (i.e., a Tikhonov matrix).

Bias-Variance Tradeoff, where regularization by ridge regression decreases condition number and multi-collinearity while increasing numerical stability and positive definiteness, which makes a well-posed problem with a well-conditioned solution in exchange for a tolerable amount of bias). Furthermore, using a larger Tikhonov lambda λ_T value adds a large regularization value to each element along the main diagonal, which not only forces the matrix to become invertible, it also forces the *inverse* of the matrix to be full of scaled-down (smaller) values. As shown in Equation 4.3b, this inverse square-matrix of scaled-down values $(A^T A + \lambda_T I_d)^{-1}$ is returned to the transpose of its original non-square shape by multiplication with A^T , which produces the pseudo-inverse of matrix A with scaled down values from matrix A^T .³⁴ Finally the Tikhonov regularized/scaled-down left-inverse of matrix A (i.e., $(A^T A + \lambda_T I_d)^{-1} A^T$) is multiplied with vector \vec{y} to produce d elements of regularized/scaled-down $x_{u,j}$, where d is equivalent to the value e from the particular user's complete set of j -indexed genre-element biases $\{x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_e}\}$.

Vector $\vec{y} = \langle y_{u,i_1}, y_{u,i_2}, y_{u,i_3}, \dots, y_{u,i_{n_u}} \rangle$ consists of the *residual* rating values y remaining from each movie i that was rated by user u in the training dataset, and n_u is the total number of movies rated by the user. More specifically, each residual rating element $y_{u,i}$ is determined by subtracting the following from the respective movie rating $r_{u,i}$ – the overall mean rating μ , the respective time-aware movie bias $b_i(t)$, the user's rating bias b_u , and the respective static movie-genre bias b_g (see Equation 4.4).

$$y_{u,i} = r_{u,i} - \mu - b_i(t) - b_u - b_g \quad (4.4)$$

In practice, the best-fit values $x_{u,j}$ are simultaneously solved for all of the genre-elements contained in a user's genre-element matrix A of rated movies in the training dataset. Also in practice, the relatively low number of instances of user-specific genre-elements (for example see *Genre Element Column Totals* in Table 7) required regularization to penalize/reduce large estimates that can result from small sample sizes, due to relatively higher noise and uncertainty. Three methods were used in combination to reduce and constrain the total variability of the user-specific bias sizes, with no loss of data, in the genre bias algorithm. The methods used were built into the equations of this section (Equations 4.2 and 4.3b) and are summarized below to show the impact of each constraining method:

- 1) Lambda diagonal matrix $(\lambda_T I_d)$ was tuned as a Tikhonov regularization parameter within the left-inverse matrix term of Equation 4.3b to reduce/scale-down the absolute value estimates of the user-specific genre-element biases in the solution vector $\vec{x} = \langle x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_e} \rangle$ of the equation. The Tikhonov lambda value ($\lambda_T = 12.75$) was tuned by minimizing the RMSE of prediction for the first 4000 users in the `edx_test` dataset (see Figure 5) with the weight parameter w also tuned to a value of 0.575 based on the same 4000 users (see Equation 4.2).³⁵
- 2) Hyperbolic-tangent (tanh) squashing function, $\tanh x = \frac{e^{2x} - 1}{e^{2x} + 1}$, was used to non-linearly transform/further scale-down the user-specific genre-element biases into real numbers \mathbb{R} within the interval $[-1, 1]$ of bias (see Equation 4.2).
- 3) A tuned proportionality factor w used weighting to further scale-down the tanh transformed/scaled-down user-specific genre-element biases (see Equation 4.2). The *final value* of the weight parameter, $w = 0.638$, was determined by minimizing the RMSE of prediction for all 68,548 users in the `edx_test` dataset (see Figure 6) with the Tikhonov lambda parameter set to the pre-tuned value of $\lambda_T = 12.75$ (based on the tuning results for the first 4000 users in the `edx_test` dataset described in step 1 above).³⁶

³⁴ $(A^T A + \lambda_T I_d)^{-1}$ is a $d \times d$ square matrix multiplied by a $d \times c$ non-square matrix A^T , which produces a non-square $d \times c$ (transpose) pseudo-inverse of matrix A , where c is the number of rows and d is the number of columns of the original matrix A .

³⁵The Tikhonov lambda λ_T and weight w parameters were tuned by an alternating iterative process, where one parameter (e.g., weight) was held constant while the other parameter (e.g., λ_T) was iterated through a range of values to find which value minimized the RMSE of the first 4000 users in the `edx_test` dataset. Next, the process was repeated with the second parameter (e.g., λ_T) held at its value which minimized RMSE, while the first parameter (e.g., weight) was iterated through a range of values to update which w value minimizes the RMSE. This alternating iterative tuning process was repeated until the optimum value of each parameter stopped changing and had converged on the same minimum RMSE value. Tikhonov lambda tuning is very time consuming and computationally expensive, thus cross validation was not a viable option. Furthermore, the alternating iterative tuning process for the λ_T and weight parameters had to be limited to a sample of the first 4000 users from the test dataset. The λ_T and weight parameters converged at a RMSE of 0.8407484 with optimal values of 12.75 and 0.575, respectively, based on a final iterative set of λ_T calculations that took 36 hours to complete for the sample of 4000 users from the test dataset (see Figure 5).

³⁶After the Tikhonov lambda parameter was optimized at $\lambda_T = 12.75$ based on the first 4000 users in the `edx_test` dataset,




Figure 5: **Alternating Iterative Tuning of Tikhonov Lambda & Weight Parameters Converged on Same Minimum RMSE = 0.8407484 where $\lambda_T = 12.75$ and $w = 0.575$. Both parameters were optimized on the first 4000 users in the test dataset.**

Figures 7a, 7b, and 7c show the sequential effects on the distribution of the user-specific biases in the genre elements, produced by the three methods described above to *constrain the total variability* of the user-specific bias sizes, with *no loss of data*. Figure 7a shows that the distribution resulting from the $\lambda_T = 12.75$ diagonal matrix as a regularization term, in Equation 4.3b, constrained most results into a normal distribution between -1 and 1, however there was a small proportion of results (0.041% or $n = 447$) that still had large estimates of bias (greater than ± 1). Figure 7b shows the subsequent effect of the hyperbolic-tangent (tanh) squashing function on the distribution, which non-linearly transformed all bias results (from figure 7a) to be centered on zero (0) bias within the interval of $[-1, 1]$ bias. Figure 7c shows the final effect on the distribution of the user-specific biases for the genre elements, which further constrained the bias estimates from figure 7b through weighting by a tuned proportionality factor of $w = 0.638$ (see Equation 4.2).

then all of the Tikhonov regularized user-specific genre-element biases $x_{u,j}$ were calculated for each of the 69,878 users in the complete training dataset (edx_train) with λ_T set at 12.75 in Equation 4.3b. Following this, the weighted tanh training results for all 69,878 users in edx_train were quickly and easily calculated by Equation 4.2, over a range of weights w from 0.550 to 0.750, to produce all the user-specific movie-genre biases, $b_{g,u}$, at each of 10 different weights w across that range. Finally, RMSE *evaluation* of these $b_{g,u}$ training results over the 0.550 to 0.750 weight w range was performed with the full 68,548-user test dataset (edx_test), which resulted in a (full) edx_test *tuned optimization* of the weight proportionality factor at $w = 0.638$, with minimum RMSE = 0.8463952 and λ_T pre-tuned at 12.75 (see Figure 6).




Figure 6: (a) Tuning the Tikhonov lambda regularization parameter λ_T in the TRLIM solutions for user’s vector of genre-element biases $x_{u,j}$ in Equation 4.3b. Y-axis shows the RMSE of various λ_T in the model’s prediction of ratings for the first 4000 users in `edx_test`. Tikhonov lambda $\lambda_T = 12.75$ gives the minimum RMSE (0.8407484) when weight $w = 0.575$, for first 4000 users in the test dataset. (b) Tuning the weight w of the hyperbolic tangent transformation of TRLIM solutions for user’s genre-element biases $x_{u,k}$ in Equation 4.2. Y-axis shows RMSE of various weights in the model’s prediction of ratings for all 68,548 users in the test dataset. Weight $w = 0.638$ gives the minimum `edx_test` RMSE (0.8463952) when $\lambda_T = 12.75$.

3.6 SVD-PCA Latent Factor Collaborative Filtering Algorithm

The residuals of the rating data, that were left after separating out the rating biases of the previously described algorithms in the model, were parsed into *latent factors* by matrix factorization through the hierarchical effects of Principal Component Analysis (PCA) with Singular Value Decomposition (SVD).³⁷ More specifically, the residuals of the rating data were SVD-PCA factorized by decomposing the residuals *and imputed values*³⁸ from a dense matrix, of completely filled-in high-dimensional user rows and movie columns, into a joint user-movie latent factor space as a form of principal component collaborative filtering. Indeed, collaborative filtering that was based on the weights of importance (loadings) from the proportional covariances³⁹ between the variables represented in the orthogonal singular vectors (latent factor eigenvectors) of the SVD-PCA RR^T user row-space and $R^T R$ movie column-space.

3.6.1 Spectral Decomposition of the Residuals into Principal Component Latent Factors

In order to apply the joint user-movie latent factor space to the model, it is conceptually useful to dissect the individual latent factors from the economy-rank SVD-PCA matrix results ($P\Sigma Q^T$) by mathematical expansion out to a hierarchical ordered series of rank-1 *spectral decomposition* matrices, that represent the contributions of the principal components from the most to the least important. The spectral decomposition

³⁷The preferred method for *numerical accuracy* in PCA is by using SVD for the factorization (reference: “prcomp” function documentation in *stats* package of R software, version 4.3.0)

³⁸Imputation was required for valid SVD-PCA analysis. Thus prior to SVD-PCA decomposition, the 98.9% empty (sparse) matrix was completely filled-in by *imputation* using a Low-Rank Matrix Completion (LRMC) approach with the Collective Matrix Factorization (*CMF*) function and the *imputeX* function of the *cmfrec* package (ver.3.5.1-1) in R (ver.4.3.0). See Section 6.1 Appendix A for the mathematical process and equations, function details, and the R code used in this project for Imputation by LRMC.

³⁹See explanation of “proportional covariances” in Footnote #15.

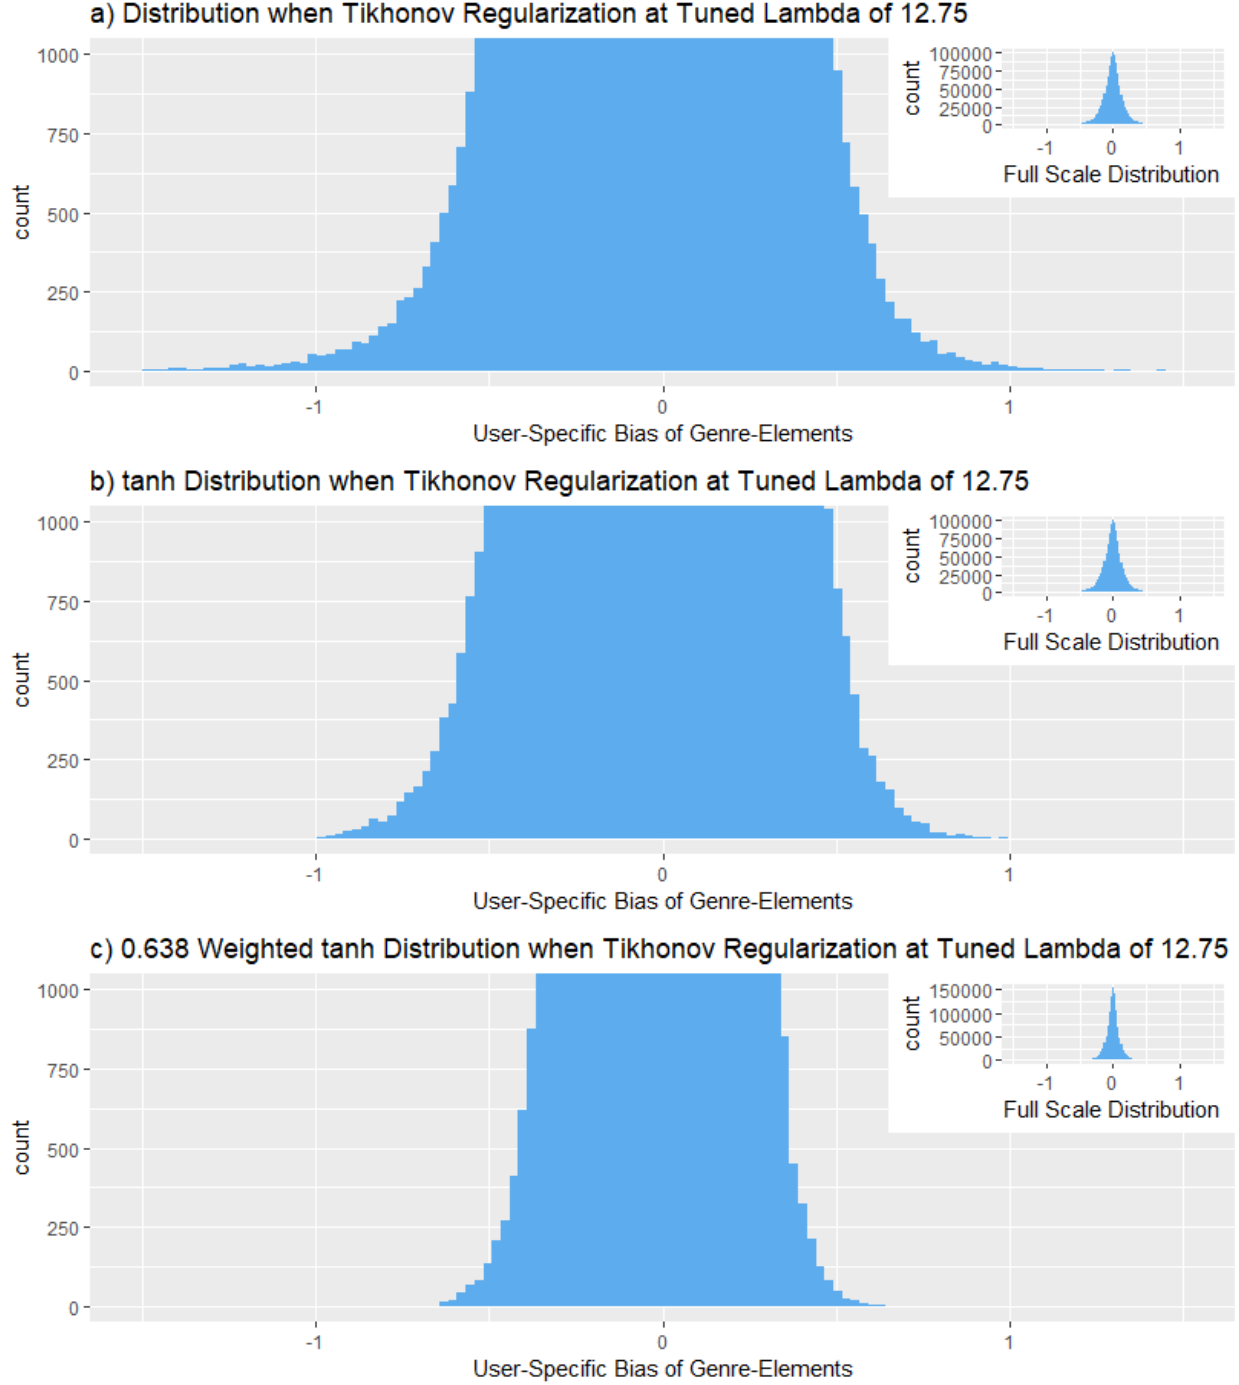


Figure 7: Effects on the distribution of the user-specific biases for the genre elements, produced by sequential application of the three methods (described in section 3.5) to constrain the total variability of the user-specific bias sizes, with no loss of data. [Figure 7a](#) shows that the distribution resulting from the $\lambda_T = 12.75$ diagonal matrix as a regularization term (see Equation 4.3b) constrained most bias estimates into a normal distribution between -1 and 1, however there was a small proportion (0.041% or $n = 447$) of the estimates that were still greater than ± 1 in magnitude. [Figure 7b](#) shows the subsequent effect of the hyperbolic-tangent (tanh) squashing function on the distribution, which non-linearly transformed all estimates (from [Figure 7a](#)) to be centered on zero bias within the interval of $[-1, 1]$ bias. [Figure 7c](#) shows the final effect on the distribution of the user-specific biases for the genre elements, which further constrained the bias estimates from [Figure 7b](#) through weighting by a tuned proportionality factor of $w = 0.638$ (see Equation 4.2).

can be written mathematically in terms of SVD notation by the following sum of rank-1 outer products:

$$R = P\Sigma Q^T = \sum_{k=1}^r \sigma_{kk} \vec{p}_k \vec{q}_k^T = \sigma_{1,1} \vec{p}_1 \vec{q}_1^T + \sigma_{2,2} \vec{p}_2 \vec{q}_2^T + \cdots + \sigma_{r,r} \vec{p}_r \vec{q}_r^T \quad (5.0)$$

Where R represents the residuals of the ratings and the imputed values in a dense matrix of $m \times n$ high-dimensional user rows and movie columns.⁴⁰ P and Q are the rank- r matrices of the economy SVD left and right singular vectors (eigenvectors) respectively, and Σ is the rank- r diagonal matrix of the singular values (eigenvalues).⁴¹ Singular value σ_{kk} is the (k, k) th entry of diagonal matrix Σ , left singular vector \vec{p}_k is the k th column of rank- r matrix P , and right singular vector \vec{q}_k^T is the k th row of rank- r matrix Q^T .

Each of the k terms in the right-hand side of Equation 5.0 is a rank-1 matrix of size $m \times n$, because it is obtained from the outer product of one independent m -dimensional column vector \vec{p}_k with one independent n -dimensional row vector \vec{q}_k^T . Furthermore, each of the $m \times n$ rank-1 k terms is referred to as a *latent component* of the original $m \times n$ matrix R , because each contains the joint user-movie latent factors that were hidden in the residuals data of matrix R [Aggarwal 2020, pp. 305-307]. In other words, each of the k latent components contains the joint user-movie latent factors, for each user row m and movie column n combination, in the $m \times n$ matrix of the respective k th principal component outer product $\sigma_{kk} \vec{p}_k \vec{q}_k^T$. Thus, to obtain the sum total of latent factors in the prediction model for a given user-movie combination, the factor at the intersection of the user row m and movie column n from each of the k terms of latent principal components are added together. Although easy to understand and visualize, this approach is an inefficient way to calculate a sum total of latent factors.

While the SVD *spectral decomposition outer product* form of the SVD-PCA provides the intuition for how the latent factor space is applied to the model, the PCA *inner (dot) product* form of the SVD-PCA provides the mathematically efficient way to calculate the same end results that are used in the model. The SVD outer product and equivalent PCA inner product forms can be written mathematically in terms of SVD and PCA notations by the following equation:⁴²

$$R = P\Sigma Q^T = \sum_{k=1}^r \sigma_{kk} \vec{p}_k \vec{q}_k^T = \begin{bmatrix} \vec{p}_{u_1}^T \bullet \vec{q}_{i_1} & \vec{p}_{u_1}^T \bullet \vec{q}_{i_2} & \cdots & \vec{p}_{u_1}^T \bullet \vec{q}_{i_n} \\ \vec{p}_{u_2}^T \bullet \vec{q}_{i_1} & \vec{p}_{u_2}^T \bullet \vec{q}_{i_2} & \cdots & \vec{p}_{u_2}^T \bullet \vec{q}_{i_n} \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ \vec{p}_{u_m}^T \bullet \vec{q}_{i_1} & \vec{p}_{u_m}^T \bullet \vec{q}_{i_2} & \cdots & \vec{p}_{u_m}^T \bullet \vec{q}_{i_n} \end{bmatrix}_{k=r} \quad (5.1)$$

where the first column (singular) vector $\vec{p}_{k=1}$ in matrix P is orthogonal to the first row (user) vector $\vec{p}_{u_1}^T$ in the matrix product $P\Sigma$, and where the first row (singular) vector $\vec{q}_{k=1}^T$ in matrix Q^T is orthogonal to the first column (movie) vector \vec{q}_{i_1} in matrix Q^T . In this case, the first row vector $\vec{p}_{u_1}^T$ contains all of the userId 1 principal component loadings from the sigma-scaled left singular vectors in matrix $P\Sigma$. Likewise, the first column vector \vec{q}_{i_1} contains all of the movieId 122 principal component loadings from the transposed right singular vectors in matrix Q^T . Thus in one efficient calculation, the dot product $\vec{p}_{u_1}^T \bullet \vec{q}_{i_1}$ of the respective

⁴⁰PCA requires row-wise pre-processing of $m \times n$ matrix R into *zero-mean centered data*. This is accomplished by first calculating the mean of each column \bar{x}_j of the matrix and storing these results into a row-wise vector of column means $\bar{\mathbf{x}}|_{j=1}^n = \langle \bar{x}_1, \bar{x}_2, \bar{x}_3, \dots, \bar{x}_n \rangle$. The vector of column means is then subtracted from each row of matrix R to achieve the zero-mean centered data. The pre-processing was performed by the default settings of the *prcomp* function (in *stats* package of R software, version 4.3.0) upon the 100% full matrix that was filled with all the known and all the imputed (see Footnote #38) rating-residual values in matrix R .

⁴¹The *economy* SVD is where the P , Q , and Σ matrices each has a rank $r = \min\{m, n\}$ of the m rows and n columns of the matrix R , and Σ is a diagonal matrix containing non-negative singular values [Aggarwal 2020, p. 306].

⁴²The singular value (eigenvalue) sigma term σ_{kk} in the SVD outer product notation is dropped in the PCA inner (dot) product notation because the sigma terms are understood to have already been multiplied with their respective left singular vectors (eigenvectors) \vec{p}_k in matrix P , i.e., the matrix product of $P\Sigma$. This PCA notation is consistent with the *prcomp* function used in the R coding of these equations (Reference: *prcomp* function documentation in *stats* package of R version 4.3.0).

principal component loadings produces the sum total of latent component factors for the userId 1 and movieId 122 combination.

The foregoing can be generalized to $\vec{p}_{u_\alpha}^T \bullet \vec{q}_{i_\beta}$ to calculate the sum total of latent factors for any userId u and movieId i combination, where $\alpha \in \{1, 2, 3, \dots, m\}$ rows and $\beta \in \{1, 2, 3, \dots, n\}$ columns in order to represent the complete $m \times n$ matrix array term. On the right-hand side of the matrix array term in Equation 5.1, the index k also becomes the length dimension of the $\vec{p}_{u_\alpha}^T$ and \vec{q}_{i_β} vectors in terms of the number of principal components (loadings) in each user vector and in each item (movie) vector, where $k = r$ is the full rank- r of principal component latent factors. The principal components dimension, index k , is also important for application of *truncated* SVD-PCA (see Section 3.6.2).

3.6.2 Truncated SVD-PCA User-Movie Latent Factor Collaborative Filtering Algorithm

Various truncation levels of the SVD-PCA latent factor algorithm were tested on the full factorization results of Equation 5.1 by PCR post-analysis, based on RMSEs from the model using Equation 5.2 (see Figure 2 and Table 9), which is the truncated version of Equation 5.1 (see for important details not repeated here), and is the SVD-PCA latent factor algorithm used in the final model:

$$R = P\Sigma Q^T \approx \sum_{k=1}^f \sigma_{kk} \vec{p}_k \vec{q}_k^T = \begin{bmatrix} \vec{p}_{u_1}^T \bullet \vec{q}_{i_1} & \vec{p}_{u_1}^T \bullet \vec{q}_{i_2} & \cdots & \vec{p}_{u_1}^T \bullet \vec{q}_{i_n} \\ \vec{p}_{u_2}^T \bullet \vec{q}_{i_1} & \vec{p}_{u_2}^T \bullet \vec{q}_{i_2} & \cdots & \vec{p}_{u_2}^T \bullet \vec{q}_{i_n} \\ \vdots & \vdots & \ddots & \vdots \\ \vec{p}_{u_m}^T \bullet \vec{q}_{i_1} & \vec{p}_{u_m}^T \bullet \vec{q}_{i_2} & \cdots & \vec{p}_{u_m}^T \bullet \vec{q}_{i_n} \end{bmatrix}_{k=f} = \tilde{P}\tilde{\Sigma}\tilde{Q}^T = \tilde{R} \quad (5.2)$$

where the matrix \tilde{R} contains the array of sums of only the first f principle-component user-movie latent factors for each user (u_α) and movie (i_β) combination in the array (see Table 8 for sample of matrix \tilde{R} when $f = 106$). Likewise, in the matrix array term in the center of Equation 5.2, f is the length (index k dimension) of the user $\vec{p}_{u_\alpha}^T$ and movie \vec{q}_{i_β} vectors in units of the number of principal components in each vector.⁴³ Furthermore, truncated SVD matrix \tilde{P} and truncated SVD matrix \tilde{Q}^T each contain only the first f principle component latent factors embodied as f left singular vectors and f right singular vectors, respectively, while the truncated SVD matrix $\tilde{\Sigma}$ contains only the first f singular values. Finally, the spectral decomposition form of the truncated SVD-PCA contains only the first f summation terms, where index $k \in \{1, 2, 3, \dots, f\}$, which represent the first f principle-component user-movie latent factors.

Table 8: **Sample* of Matrix \tilde{R} Array of Sums of the First $f = 106$ Principle-Components in the User-Movie Latent Factors for Each UserId (u) and MovieId (i) Combination**

userId \ movieId	122	292	316	329	355	356	362	364	370	377
1	0.48260789	0.23653283	0.35533192	0.3474575	0.58984433	-0.52324319	0.207408079	-0.07289591	0.53959384	-0.152524916
2	-0.04738856	-0.28591677	-0.37528951	-0.6026522	-0.19442097	-0.03402969	-0.113431827	-0.11578284	-0.04560144	-0.070438335
3	0.08292988	0.28315704	0.20187958	0.2428388	0.39327865	-0.06497949	0.093609016	0.15971138	-0.01862918	0.189619334
4	-0.03398417	-0.65592773	0.56302364	0.8647768	0.50490463	0.18232545	0.149386877	0.24833746	-0.10081180	-0.683909869
5	-0.37917389	-0.49475448	-0.77352238	-0.3877704	-0.34956085	-0.88243929	-0.001331489	-0.93242060	0.06451146	-0.711925660
6	-0.21185358	-0.07844738	-0.52287698	-0.2285524	-0.44075437	0.07266437	-0.018393208	-0.08481047	-0.32633912	0.094296358
7	-0.19096753	-0.21777020	-0.29862736	-0.4300856	-0.35839663	-0.22685366	-0.249630663	-0.40111337	0.08134932	-0.169477253
8	-0.06254857	0.03608508	0.14027859	-0.1567546	0.13909945	-0.34075098	-0.094190258	-0.07527481	-0.18828470	-0.009407292
9	-0.40367416	-0.44472518	-0.08402934	0.1590944	-0.27016028	-0.63953273	0.222083846	-0.52562000	0.02007647	-0.539286742
10	-0.08090274	-0.28263963	0.13260836	0.2197566	-0.04106331	-0.81239339	-0.086639934	0.01738328	0.08036595	-0.323648869

* Sample of first 10 User IDs and first 10 Movie IDs in the $m = 69,878$ rows (users) and $n = 10,677$ columns (movies) of the matrix.

⁴³The matrix array term $\left[\vec{p}_{u_\alpha}^T \bullet \vec{q}_{i_\beta} \right]_f$ in Equation 5.2 is equivalent to the dot product term $\left(\vec{p}_u^T \bullet \vec{q}_i \right)_f$ in Equation 1.0, where u and i indicate the userId and movieId, respectively. In the matrix array term, α and β respectively indicate the row number and column number of the array, in order to represent the complete $m \times n$ matrix array.

Table 9: RMSE of Truncation Levels Tested by PC Regression in the SVD-PCA Principal Components Post-Analysis of the User-Movie Latent Factors (Imputed by Low-Rank Matrix Completion)

PC Factors Truncation Level	PC Factors	edx_test	PC-106 Test	edx_train
Description	Number (f)	Clip-RMSE	Delta-RMSE	Clip-RMSE
1 28.8% CPV (1st Principal Component)	PC-1	0.8287763	0.0363446	0.7850463
2 41.7% CPV	PC-2	0.8258035	0.0333718	0.7804250
3 49.0% CPV	PC-3	0.8203130	0.0278813	0.7718524
4 61.0% CPV	PC-6	0.8098027	0.0173710	0.7534657
5 70.6% CPV	PC-13	0.8008788	0.0084471	0.7301808
6 80.0% CPV (See Figure 9a)	PC-37	0.7940755	0.0016438	0.6892763
7 83.8% CPV (PEV-Elbow, Figure 9c)	PC-58	0.7929332	0.0005015	0.6644925
8 88.8% CPV (Denoised Level, Figure 2)	PC-106	0.7924317*	0.0000000	0.6202779
9 90.0% CPV (See Figure 9a)	PC-124	0.7925751	0.0001434	0.6065232
10 93.3% CPV (SD-Elbow, Figure 9b)	PC-240	0.7941522	0.0017205	0.5352164
11 95.0% CPV (See Figure 9a)	PC-429	0.7947596	0.0023279	0.4573800
12 97.4% CPV	PC-1000	0.7945929	0.0021612	0.3184481
13 99.0% CPV	PC-2000	0.7945643	0.0021326	0.1901892
14 99.6% CPV	PC-3000	0.7946401	0.0022084	0.1239968
15 100 % CPV (Economy Rank-r)	PC-10677	0.7946946	0.0022629	0.0632548*

Key: CPV = Cumulative Proportion of Variance PC = Principal Component
 PEV = Percentage of Explained Variance SD = Standard Deviation
 RMSE = Root Mean Squared Error * = Minimum RMSE Value
 Clip-RMSE = Final Model RMSE of PC Level after Min & Max Adjustments (See Equation 1.1)
 Delta-RMSE = Increase in edx_test RMSE of PC Level from the (minimum) PC-106 Factors Level

4 Results & Discussion

4.1 Model Results and Performance Attributes

Improvement in the RMSE of each algorithm in the model was evaluated to determine the key performance attributes and their relative order of importance. Table 10, Table 11 and Figure 8 summarize the evolution and rank order of model improvements in terms of RMSE as various attributes were developed and added to the model. Model improvements were compared by the percent decrease in RMSE value relative to the starting RMSE of 1.06021, which was from the model starting point of the overall mean of all movie ratings by all users in the training dataset.

The first and second ranked model improvements were due to Movie Global Static Biases (10.982%) and User Global Static Biases (7.120%) respectively, which is consistent with observations by Koren, Bell, and Volinsky (2009) that “much of the observed variation in rating values is due to effects associated with either users or items, known as *biases* or *intercepts*, independent of any interactions” [Koren, Bell, and Volinsky 2009, p. 45]. Whereas, the third and fourth ranked model improvements were due to the *interactions between users and items*, with the third rank held by the Truncated SVD-PCA User-Movie Latent Factor Collaborative Filtering Algorithm (5.05%),⁴⁴ which summed the first 106 principle-component user-movie latent factors for each user and movie combination in the data – data that was imputed by Low-Rank Matrix Completion.⁴⁵ The fourth ranked model improvement was due to the User-Aware Movie-Genre Bias (1.467%),⁴⁶ which summed the movie-genre global static bias and an inferred user-specific movie-genre bias that was based on the Tikhonov Regularized Left-Inverse Matrix (TRLIM) algorithm.

⁴⁴RMSE Improvement 5.05% = 0.198% from PC-106 Truncation of the 10677 Principal Component Latent Factors + 4.852% from Imputation by Low-Rank Matrix Completion (see Table 11 and Figure 8).

⁴⁵Imputation was required for valid SVD-PCA, with LRMC clearly providing the best imputation in terms of % RMSE Improvement. Nonetheless, the Truncated SVD-PCA algorithm achieved the 3rd Rank in model improvement *regardless of which of the four Imputation Methods were used* (see Table 11 Model #11a through #14b).

⁴⁶RMSE Improvement 1.467% = 0.023% from the Movie-Genre Global Static Bias + 1.444% from the User’s Specific Movie-Genre Bias.

Table 10: Summary of Attributes in Model Improvements of Root Mean Squared Error (RMSE)

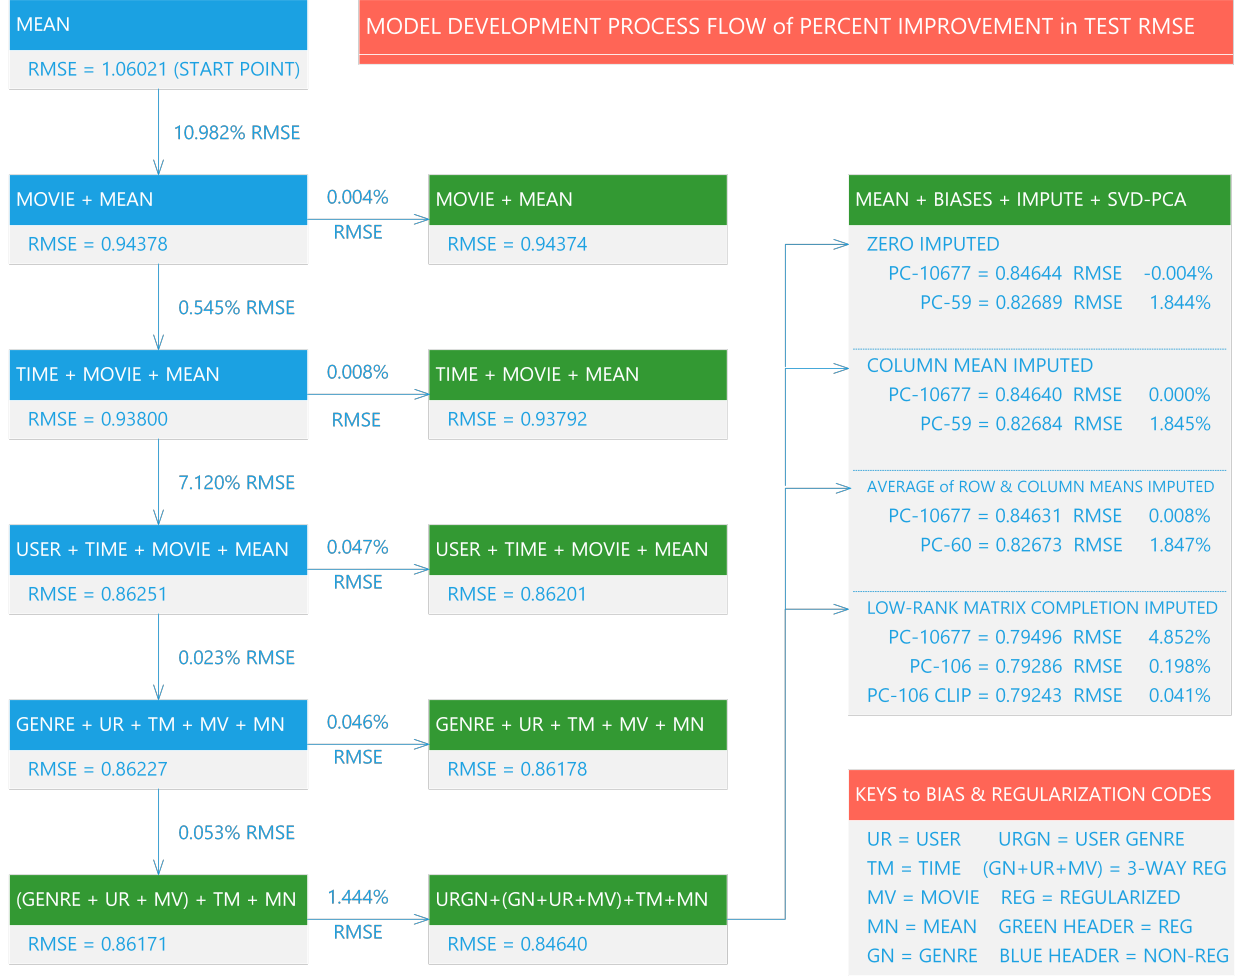
Model #	Model Name	RMSE	Type of Regularization	Static Global Mean	Static Global Movie Biases	Time Bin Biases	Static Global User Biases	Static Global Genre Biases	User Id Genre Biases	Truncated to PC# Latent Factors	Imputation Method for Missing Data
1	Mean	1.06021		Y
2	Movie + Mean	0.94378		Y	Y
3	Reg Movie + Mean	0.94374	1-Way	Y	Y
4	Time + Movie + Mean	0.93800		Y	Y	Y
5	Reg Model #4	0.93792	1-Way + 1-Way	Y	Y	Y
6	User + Model #4	0.86251		Y	Y	Y	Y
7	Genre + Model #6	0.86227		Y	Y	Y	Y	Y	.	.	.
8	Reg Model #6	0.86201	1-Way x 3	Y	Y	Y	Y
9a	Reg Model #7 (1-Way x 4)	0.86178	1-Way x 4	Y	Y	Y	Y	Y	.	.	.
9b	Reg Model #7 (3+1 Way)	0.86171	3-Way + 1-Way	Y	Y	Y	Y	Y	.	.	.
10	UserIdGenre + Model #9b	0.84640	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	.	.
11a	Zero PC10677 + Model #10	0.84644	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	.	Zero
11b	Zero PC59 + Model #11a	0.82689	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	Y	Zero
12a	CM PC10677 + Model #10	0.84640	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	.	Column Mean
12b	CM PC59 + Model #12a	0.82684	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	Y	Column Mean
13a	ARCM PC10677 + Model #10	0.84631	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	.	ARCM
13b	ARCM PC60 + Model #13a	0.82673	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	Y	ARCM
14a	LRMC PC10677 + Model #10	0.79496	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	.	LRMC
14b	LRMC PC106 + Model #14a	0.79286	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	Y	LRMC
15	MinMax Clip of Model #14b	0.79243	3+1 Way + TRLIM	Y	Y	Y	Y	Y	Y	Y	LRMC

Key: RMSE = Root Mean Squared Error 1-Way = Feature(s) Regularized Separately with Separate Lambda Values
Y = Presence of Attribute in the Model 3-Way = Movie, User and Genre Features Regularized Together with 1 Lambda Value
Reg = Regularized Version of Model TRLIM = Tikhonov Regularized Left-Inverse Matrix (See Equation 4.3)
UserIdGenre = User Specific Genre Bias PC# = SVD-PCA Results Truncated to the First # Principal Component Latent Factors
3+1 Way = 3-Way Movie, User and Genre Features Regularized Together with 1 Lambda Value + 1-Way Time Regularization
CM = Column Mean method of imputation in Rating Residuals Matrix (Part of Preprocessing for SVD-PCA)
ARCM = Average of Row and Column Means method of imputation in Rating Residuals Matrix (Part of Preprocessing for SVD-PCA)
LRMC = Low-Rank Matrix Completion method of imputation in Rating Residuals Matrix (Part of Preprocessing for SVD-PCA)
MinMax Clip = Minimum of 0.5 & Maximum of 5.0 Adjustments to Rating Predictions Outside of Scoring Range (See Equation 1.1)

Table 11: Percent Improvement in RMSE and Rank Order of Importance of Key Model Attributes

Model #	Model Name	RMSE	RMSE Improvement (Model A - B)	% RMSE Improvement	Rank of Importance (* In Final Model)	Key Model Attributes Measured by % RMSE
1	Mean	1.06021	NA	NA	NA	Starting Point with Overall Mean Rating
2	Movie + Mean	0.94378	0.11643 (1-2)	10.982 %	1 *	Movie Global Static Biases
3	Reg Movie + Mean (1-Way)	0.94374	0.00004 (2-3)	0.004 %	6e	1-Way Movie Regularization
4	Time + Movie + Mean	0.93800	0.00574 (2-4)	0.545 %	5 *	Temporal Dynamics of Movie Bias
5	Reg Model #4 (1-Way x 2)	0.93792	0.00008 (4-5)	0.008 %	6d	1-Way Movie Reg + 1-Way Time Reg
6	User + Model #4	0.86251	0.07549 (4-6)	7.120 %	2 *	User Global Static Biases
7	Genre + Model #6	0.86227	0.00024 (6-7)	0.023 %	4b *	Genre Global Static Biases
8	Reg Model #6 (1-Way x 3)	0.86201	0.00050 (6-8)	0.047 %	6b	1-Way Reg x 3 = User, Movie & Time Regs
9a	Reg Model #7 (1-Way x 4)	0.86178	0.00049 (7-9a)	0.046 %	6c	1-Way Reg x 4 = Genre, User, Movie & Time Regs
9b	Reg Model #7 (3Way+1Way)	0.86171	0.00056 (7-9b)	0.053 %	6a *	3-Way Genre, User, Movie Reg + 1-Way Time Reg
10	UserIdGenre + Model #9b	0.84640	0.01531 (9b-10)	1.444 %	4a *	User Specific Genre Bias with TRLIM
11a	Zero PC10677 + Model #10	0.84644	-0.00004 (10-11a)	-0.004 %	3h	PC10677 PCLF with Zero Imputation
11b	Zero PC59 + Model #11a	0.82689	0.01955 (11a-11b)	1.844 %	3g	PC59 Truncation of PCLF with Zero Imputation
12a	CM PC10677 + Model #10	0.84640	0.00000 (10-12a)	0.000 %	3f	PC10677 PCLF with Column Mean Imputation
12b	CM PC59 + Model #12a	0.82684	0.01956 (12a-12b)	1.845 %	3e	PC59 Truncation of PCLF with CM Imputation
13a	ARCM PC10677 + Model #10	0.84631	0.00009 (10-13a)	0.008 %	3d	PC10677 PCLF with Avg Row&Col Means Imputation
13b	ARCM PC60 + Model #13a	0.82673	0.01958 (13a-13b)	1.847 %	3c	PC60 Truncation of PCLF with ARCM Imputation
14a	LRMC PC10677 + Model #10	0.79496	0.05144 (10-14a)	4.852 %	3a *	PC10677 PCLF with LowRankMatrixComp Imputation
14b	LRMC PC106 + Model #14a	0.79286	0.00210 (14a-14b)	0.198 %	3b *	PC106 Truncation of PCLF with LRMC Imputation
15	MinMax Clip of Model #14b	0.79243	0.00043 (14b-15)	0.041 %	7 *	Min 0.5 & Max 5.0 Predicted Rating Adjustment

Key: RMSE = Root Mean Squared Error 1-Way = Feature(s) Regularized Separately with Separate Lambda Values
UserIdGenre = User Specific Genre Bias 3-Way = Movie, User and Genre Features Regularized Together with 1 Lambda Value
Reg(s) = Regularized Version of Model(s) TRLIM = Tikhonov Regularized Left-Inverse Matrix (See Equation 4.3)
PCLF = Principal Component Latent Factors PC# = SVD-PCA Results Truncated to the First # Principal Component Latent Factors
CM = Column Mean method of imputation in Rating Residuals Matrix (Part of Preprocessing for SVD-PCA)
ARCM = Average of Row and Column Means method of imputation in Rating Residuals Matrix (Part of Preprocessing for SVD-PCA)
LRMC = Low-Rank Matrix Completion method of imputation in Rating Residuals Matrix (Part of Preprocessing for SVD-PCA)
MinMax Clip = Minimum of 0.5 & Maximum of 5.0 Adjustments to Rating Predictions Outside of Scoring Range (See Equation 1.1)

Figure 8: Model Development Process Flow of Percent Improvement in `edx_test` RMSE

The fifth ranked model improvement was due to incorporation of the Temporal Dynamics of Movie Bias (0.545%), which is consistent with Yehuda Koren’s (2009) observation in *Collaborative Filtering with Temporal Dynamics* that “much of the temporal variability is included within the baseline predictors” [p. 450], and so “modeling temporal dynamics should be a key when designing recommender systems” [p.447]. In other words, there is a lot of untapped temporal bias, hidden within the movie biases and the user biases, that must be gleaned through temporal algorithms in order to maximize their potential for improving recommender model accuracy.

The sixth ranked model improvement was due to Regularization of Biases, and furthermore, *how the regularization parameters were calculated* made for interesting improvements in the overall accuracy of the model. In the case of regularization of the user, movie and genre global static biases, the best method of parameter calculation is an example of *Stein’s Phenomenon* (more on this below). In the case of regularization of the 2-year Time Bin Movie Bias, as previously discussed (Figure 4, and Section 6.2 Appendix B), the *analytical approach based on the partial derivative of the cost function* was the best method of regularization parameter calculation. Finally, the seventh ranked model improvement was due to Min & Max Clipping of the final model rating predictions that were outside of the minimum 0.5 and maximum 5.0 scoring range (See Equation 1.1). The forgoing summarizes the top seven key performance attributes of the model, and some of their interesting aspects that were previously discussed. Whereas, some additional interesting aspects and details of these attributes will be addressed in the remaining sections of the Results & Discussion below.

4.2 Comparison with Published Models and Results

Comparison of the final recommender model RMSE of the holdout dataset (0.7921) with published RMSE results in the literature on recommender models, that used the MovieLens 10M dataset, turned up a comparable RMSE (0.7907)⁴⁷ from the SVDFeature recommender model developed by Chen et al. (2012), which is described in their paper *SVDFeature: A Toolkit for Feature-based Collaborative Filtering*.⁴⁸ The SVD-Feature model was the *KDD Cup Track 1 Winner* in the 2011 and the 2012 machine learning competitions. The SVDFeature model is a synthesis, of many variants of matrix factorization models in the recommender literature, into a single generalized model abstraction. This overall approach lead to development of software tools for efficiently optimizing all model parameters of the linear regression feature biases and of the matrix factorization collaborative filtering. The SVDFeature model is designed to easily incorporate user and item implicit and explicit information as well as temporal, neighborhood, and hierarchical information.

Comparison of the SVDFeature model to the final recommender model of this report found that the bias adjusted latent factor model of collaborative filtering easily fits within the SVDFeature model paradigm, however with two differences in detail. The first difference is that the SVDFeature model uses Stochastic Gradient Descent to train all parameters and minimize the cost function, whereas the model of this report used Iterative Linear Regression for the feature bias parameters and also used Alternating Least Squares regression for imputation of the matrix factorization latent features. The second difference is that the MovieLens 10M tags.dat data was available for incorporation into the SVDFeature model, whereas the tags.dat data was not included in the model of this report. The tag.dat data was introduced by MovieLens in December of 2005 and consists of user comments, thumbs up/down, and/or “tag expressions” on each of the movies a user reviewed [Harper & Konstan 2015, p. 10]. This tag data would require some pre-processing by sentiment analysis before incorporation into the SVDFeature model as implicit user information on each rated item. This additional information can easily account for the 0.0014 improvement in RMSE of the SVDFeature model Benchmark for the MovieLens 10M dataset verses the final RMSE of the model for this report.

The bias adjusted latent factor model of collaborative filtering is essentially a type of Regularized SVD (RSVD) model when one carefully examines the similarities and differences.⁴⁹ Historically, RSVD is one of the most widely used matrix factorization methods for estimating the feature vectors for predicting user/item interaction factors. A Google Scholar search of the “Regularized SVD recommender model” literature produced 16,300 unique citations published between 2006 and 2023 (inclusive), 382 of which were published so far within 2023. Several of these citations are listed in the References of this report: Ackerman (2012); Koren, Bell & Volinsky (2009); Li et al. (2019); Paterek (2007); Pilászy (2010); Rendle, Zhang & Koren (2019); and Takács et al. (2007). Close examination of the RSVD models revealed a number of similarities with the bias adjusted latent factor model of collaborative filtering described by this report, such as pre-process subtraction of baseline predictions (i.e., biases), low-rank matrix approximations, matrix factorization regularization methods, early stopping⁵⁰, and clipping predictions to the $\{0.5, \dots, 5\}$ MovieLens range or the $\{1, \dots, 5\}$ Netflix range.

4.3 Weighted-tanh TRLIM Algorithm & User-Specific Item-Feature Bias

The weighted-tanh TRLIM process can be generalized to turn virtually any movie item feature into a user-aware version of the feature by assigning user-specific interaction biases to the movie item feature. This has implications for increasing Explainable Machine Learning and AI Transparency by pulling the influence

⁴⁷The 0.7907 RMSE value for the SVDFeature model was reported by Rendle et al. (2019) in a list of methods representing baseline Benchmarks for the MovieLens 10M dataset (see Table 1 in Rendle et al. 2019).

⁴⁸Also see for a more detailed description and background of their model: Chen et al. (2011) “Feature-Based Matrix Factorization.”

⁴⁹The SVDFeature model discussed above is a generalized abstraction of a Regularized SVD model.

⁵⁰Early stopping is a regularization technique used in machine learning to prevent overfitting. It involves monitoring the performance of a machine learning model on a test set during the training process and stopping the training process once the performance on the test set starts to degrade. Early stopping can improve the generalization performance of a model on new, unseen data, particularly in cases where the training dataset is small or noisy [Wikipedia (2022, December 24)].

of user-specific item-feature interaction-biases out of the subsequent matrix factorization part of the model, where any remaining rating values are factorized into latent user-item interactions. The exact nature of latent interactions are often difficult to define, which is not the case with assigned user-specific item-feature interactions.

For example, user-specific movie-*actor* biases can be inserted in the model by following a process very similar to the user-specific movie-genre biases. The weighted-tanh TRLIM calculations can be performed on the actor-matrix extracted from each user’s rated movies⁵¹ to optimize the bias (coefficients) that project the updated rating-residuals $y_{u,i}$ onto the pseudo-inverse of the actor-matrix (update includes deduction of the user-aware movie-genre bias, i.e., compare $y_{u,i} = r_{u,i} - \mu - b_i(t) - b_u - b_g(u)$ to Equation 4.4). This would generate the optimized solution vector of actor biases $\vec{x} = \langle x_{u,j_1}, x_{u,j_2}, x_{u,j_3}, \dots, x_{u,j_a} \rangle$, where $j \in \{1, 2, 3, \dots, a\}$ is the complete set of unique actors in a specific user’s movie ratings in the training dataset.

One significant drawback is the costly process of tuning Tikhonov regularization parameters λ_T for different types of user-specific item-features. However in the case of *similar types* of user-specific item-features, the tuned λ_T value for the first feature can be used as a starting point for the much easier task of fine-tuning the λ_T value of another similar feature. For example, user-specific movie-actor biases and user-specific movie-genre biases are similar biases that are both based on the user-movie interaction and are both measured in units of rating bias, and are therefore expected to have similar variance in interaction bias. Thus they are likely to have similar tuned λ_T values, because the Tikhonov parameter is simply a regulator of variance that scales down the relative proportions & differences in actor biases and the relative proportions & differences in genre biases, to a range within (approximately) -1 to +1 rating bias (see for example, Figure 7a).

It is also noteworthy that improvements in model accuracy from the user-specific movie-genre bias were partly due to inclusion of weighted non-linear scaling with a hyperbolic-tangent (tanh) function in order to further constrain variability in the TRLIM output of user-specific movie-genre bias. It is interesting that the tanh scaling combined with the learned weighting of the TRLIM output is similar to *Batch Normalization* in a neural network. Batch Normalization acts as a form of regularization by reducing variance and smoothing the optimization landscape in a neural network [Santurkar et al. 2018]. Thus, the weighted-tanh constraint of the TRLIM output may be viewed as a Batch Normalization of the interaction biases between users and items. This then leads to the insight that the improvements in model accuracy, from the weighted-tanh TRLIM algorithm for user-specific item-features, are due to λ_T scaled-down & weighted-tanh normalized variances in bias (projection coefficients), which likewise nudge the predictions in the right directions with normalized scaled-down relative biases for an overall improvement in model accuracy.

From an Information Science perspective, however, there is an ontological arbitrariness in the selection of the user-specific item-feature interactions for assigning biases from the data, as well as an ontological arbitrariness in the order in which the various user-specific item-feature biases are added to the model. Furthermore, the ontological arbitrariness is a caveat for potential benefits in Explainable Machine Learning and AI Transparency attributed to the assigned user-specific item-feature interactions. Nonetheless the common sense assumption is that, there is practical justification for a given user-specific item-feature bias when *accounting for such a bias leads to improvements in model accuracy*. This *a posteriori* epistemological justification also applies to the order in which various user-specific item-feature biases are added to the model. Likewise, the justification applies to weighted-tanh TRLIM *improvements in model accuracy* – improvements which are very likely when the weighted-tanh TRLIM algorithm is applied to highly-probable or common user-item interactions.

⁵¹This assumes pre-processing where lists of leading actor names have been added to each movie ID, for example, by matching the MovieLens movieId to the TMDB ID in The Movies Dataset available on Kaggle.com. This dataset is an ensemble of data collected from The Movie Database and from GroupLens that contains metadata for all 45,000 movies listed in the full MovieLens Dataset released in July 2017. The metadata includes cast, crew, plot keywords, budget, revenue, posters, release dates, languages, production companies, countries, TMDB vote counts and vote averages. More specifically, the MovieLens movieId can be matched to the corresponding TMDB ID via the links.csv data file in The Movies Dataset, and then to the corresponding cast of actors in the credits.csv data file in The Movies Dataset. Click link to The Movies Dataset for more information: <https://www.kaggle.com/datasets/rounakbanik/the-movies-dataset>

4.4 Benefits of Truncated SVD & Truncated SVD-PCA

Truncated SVD is a classical regularization method in matrix factorization, especially in regard to ill-conditioned matrices, that uses the best low-rank approximation for the model [Hansen 1987, Xiang & Zou 2013, Garg & Goyal 2014, Mesgarani & Azari 2019, Ito & Jin 2020]. The low-rank number “plays the role of a regularization parameter, which has to be carefully chosen. . . . Choosing its value amounts to finding a compromise between fidelity to the original model and numerical stability” [Pes & Rodriguez 2020, p. 468]. Thus truncated SVD can lead to a more robust model, that is also less prone to over-fitting and can generalize better to new data.

“Truncated singular value decomposition is, in fact, the primary way in which SVD is used in real applications” [Aggarwal 2020, p. 307]. Only a small proportion of the singular values/vectors are large in real-world matrices, thus truncated SVD can retain a large level of accuracy by retaining just the large singular value (low-rank) singular vectors [Aggarwal 2020, p. 308]. An important side effect of SVD dimensionality reduction is that it often reduces the amount of noise in the data [Aggarwal 2020, p. 324], i.e., a denoising effect. The assumption is that low-level noise is independent of the dominant patterns in the data. Therefore, this noise aggregates in the small singular value (high-rank) singular vectors of SVD, which are independent of the dominant data patterns that aggregate in the large singular value (low-rank) singular vectors [Aggarwal 2020, p. 324]. Thus, benefits of using truncated SVD are model improvements in accuracy and efficiency from much smaller low-rank matrix approximations that retain a high level of accuracy by dropping many components with smaller singular values containing mostly noise [Aggarwal 2020, pp. 307-313; Brunton & Kutz 2022, pp. 9-12, 34-39].

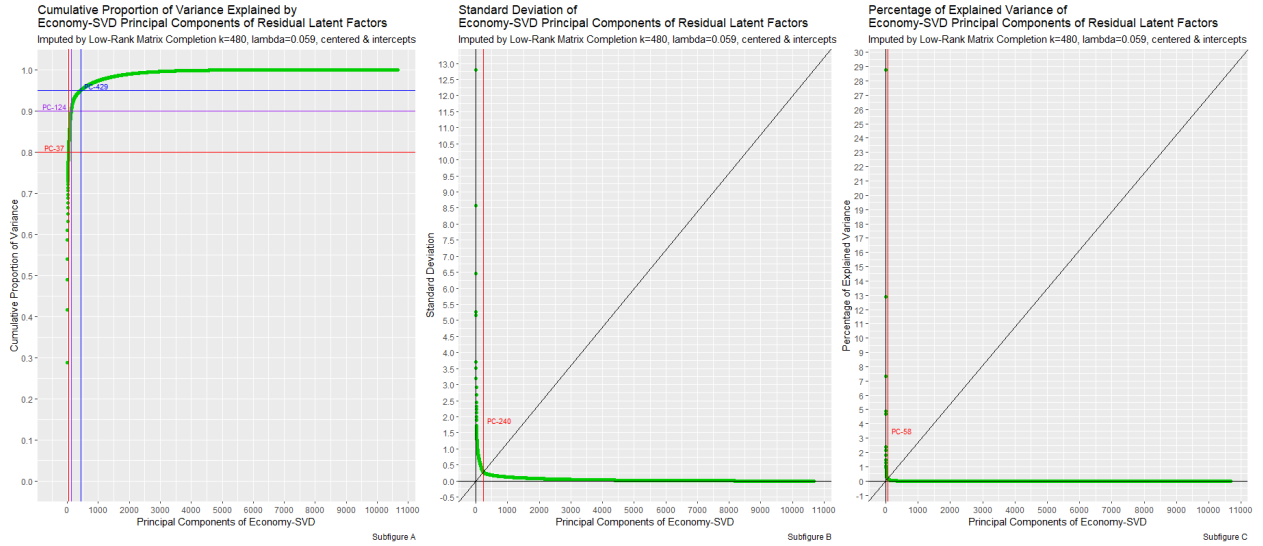


Figure 9: Comparison of PCR Optimized PC-106 to Graphical Selection Methods of SVD-PCA Truncation
 Subfigure A: 80%, 90% & 95% Cumulative Proportion of Variance Explained by PC-37, PC-124 & PC-429
 Subfigure B: Standard Deviation Scree Plot of Principal Components and PC-240 Elbow Bisection
 Subfigure C: Percentage of Explained Variance Scree Plot of Principal Components and PC-58 Elbow Bisection

Truncated SVD-PCA typically starts with the economy rank (full-size) matrix and then post-analysis drops many of the smaller and noisy components to create the truncated rank SVD-PCA. In contrast to post-analysis using arbitrary graphical-selection methods (see Figure 9 and Table 9), when using an optimization-centric post-analysis approach, truncated SVD-PCA provides the best possible low-rank approximation of a matrix in terms of minimizing squared error [Aggarwal 2020, p. 318]. All the benefits of truncated SVD-PCA were realized in the optimization-centric Principal Components Regression (PCR) approach in the post-analysis of the 10,677 principal components of the economy rank model, where only the first 106 principal component singular vectors were found to best minimize the RMSE of the test dataset predictions in the final model (see Section 3.6.2, Table 9, and Figure 2). Thus truncated SVD-PCA by PCR retained




Figure 10: Latent Movie Factors Uncovered by Mapping Top 10 & Bottom 10 Movie Titles of Principal Component Dimensions PC1 and PC2

maximum test accuracy with less than 1% of the original 10,677 principal component singular vectors. As a result, the optimal collaborative filtering of the final model was achieved by only the first 106 latent factors calculated for each user-movie combination. See Figure 10 for examples of movie latent factors uncovered in the final model from the movie title mappings of principal components PC-1 and PC-2.

4.5 Stein's Phenomenon

Models 9a & 9b in Table 11, as well as the lower left corner of Figure 8, show that the 0.053% RMSE improvement in the model achieved through a 3-way *simultaneous regularization* of the user, movie and genre global static biases was slightly better in RMSE than the 0.046% RMSE improvement achieved through three 1-way *individual regularizations* of the user, movie and genre global static biases. This small 0.007% of additional improvement in RMSE is *not an artifact* of noise or rounding errors. Indeed, this additional improvement in RMSE is based on *how* the regularizations were calculated – simultaneously versus individually – and is an example of Stein's phenomenon.

Stein's phenomenon (published in 1956) is a counter intuitive observation in decision theory and estimation theory, that revealed the Bayesian cracks in Frequentist statistics through the discovery of combined shrinkage

estimators (for the parameters of three or more variables) that are more accurate than any Frequentist method that estimates the parameters separately [Efron 2023]. For example, predicting the season batting average for each player on a baseball team based on their individual batting averages from their first 50 times at bat. The predictions are more accurate, in terms of total mean squared error, when calculated together using a shrinkage estimator based on the combined batting averages, because the individual batting averages are adjusted proportionally toward the overall grand mean of the batting averages [Efron & Morris 1977]. Stein’s phenomenon, also known as the James-Stein estimator (published in 1961), is a special case of Bias-Variance Tradeoff, where the variance of the parameter estimated across the samples is reduced by increasing the bias in the estimated parameter toward the overall grand mean. Furthermore, in the James-Stein case, both of these sources of error (bias & variance) *contribute* to the overall reduction in mean squared error of the parameters! For high dimensional data, however, Empirical Bayes methods provide better shrinkage estimators for parameters than the frequentist James-Stein method, whereas James-Stein is better for shrinkage estimation where there are from 3 to about 20 variables [Efron 2023]. For only one or two variables, either the frequentist individual average or the individual Maximum Likelihood Estimation are still the best estimates of the parameter(s) in terms of mean squared error.

4.6 Imputation of Missing Rating Values

According to Ghazanfar and Prugel (2013), “the most important task for a recommender system is to accurately predict the rating of the non-rated user/item combination and recommend items based on these predictions [p.65].” In recommender systems, such as MovieLens, Netflix, YouTube, etc., it is common for customers to provide ratings on less than 1% of the available content (items), which leads to the problem of making predictions under sparse data conditions. Collaborative filtering (CF) based recommender systems are disadvantaged under sparse data conditions, because the correlation coefficient between users (or items) is poorly approximated when there are very few or no similar users (or items) [Ghazanfar and Prugel 2013]. A common approach to mitigate the sparse data problem is to impute the missing data values with an estimate based on the available data and an understanding, or assumption, of why the data are missing.

4.6.1 Types of Missing Data

According to *Missing Data* by Wikipedia (2023, June 18), “understanding the reasons why data are missing is important for handling the remaining data correctly. If values are missing completely at random, the [remaining] data sample is likely still representative of the population. But if the values are missing systematically, analysis [of the remaining data] may be biased.” Missing data are categorized into three basic types: Missing Completely At Random (MCAR), Missing At Random (MAR), and Missing Not At Random (MNAR).

In the context of recommender systems, MCAR *assumes that a rating that is missing is completely independent of the value that the rating would have had if it were not missing*. An illustrative example, that violates the MCAR assumption, is the bias in missing values due to the tendency for people to avoid weak, poor or negative outcomes, which also biases the remaining observed (non-missing) values to have strong, good or positive outcomes. Marlin et al. (2012) in “Collaborative Filtering and the Missing at Random Assumption” showed that *the MovieLens rating distribution is skewed toward higher ratings*. Ghazanfar and Prugel (2013) are in agreement with this observation, and both the Marlin and the Ghazanfar research groups argued that this is because *MovieLens users are much more likely to watch and rate movies they think they will like*, rather than watch and rate movies they don’t think they will like. In other words, this rating bias is caused by a *naturally human selective rating behavior, which affects all content rating systems* – MovieLens, Netflix, YouTube, etc.

Although selective rating behavior explains the bias toward higher ratings, it does not explain *all* missing ratings in recommender systems, only a portion of the missing ratings. For most users in *large* recommender systems, MCAR is a safe assumption for explaining part of, if not the majority of, a user’s missing ratings simply because of the impractical reality of physically viewing or considering more than a modest fraction of the total long-form content that is available. For example, some current estimates of total long-form content

are: Netflix 17,000 titles, Amazon Prime 28,000 titles, and YouTube 800 Million videos.⁵² In other words, most missing ratings are of movies and videos that a user is simply unaware of and has nothing to do with avoiding a negative outcome.

Missing At Random (MAR) is a bit of a misnomer because MAR occurs when the missingness of the data is not fully at random, for the reason that the missing data can be accounted for by other variables or information that *are independent* of the values of the missing data, i.e., the missingness of the data still does not correlate with the value levels of those missing data. A good example that Wikipedia (2023, June 18) gives is that, “males are less likely to fill in a depression survey but this has nothing to do with their level of depression, after accounting for maleness.” In the context of MovieLens, the missing ratings due to selective rating behavior are not MAR because the reason for their missingness *is due to the values of the missing ratings* (i.e., the missing ratings correlate with lower rating values) and are not due to independent variables such as the gender or age of the MovieLens users. Nonetheless, other proportions of the missing ratings in recommender systems can still be categorized as MAR because of independent factors such as gender or age, for example, adults tend to watch less children oriented movies while children tend to watch less adult oriented movies. In addition to this natural age bias, film-content ratings (e.g., PG-13, R, and NC-17 ratings) are also independent MAR factors that account for missing ratings, especially for minors.

According to *Missing Data* by Wikipedia (2023, June 18), Missing Not At Random (MNAR) is data that is “neither MAR nor MCAR (i.e. the value of the variable that’s missing is related to the reason it’s missing). To extend the previous example, this would occur if men failed to fill in a depression survey *because* of their level of depression.” In the context of MovieLens and other content rating systems (e.g., Netflix, YouTube, etc.), the probability of observing a rating depends on the value of the rating, and thus the empty values may be considered MNAR from a probabilistic standpoint [Marlin et al. 2012]. However, as previously discussed, *all the reasons* for missing ratings in recommender systems are more nuanced, with various proportions covering all three categories of missingness – MCAR, MAR, and MNAR. Where the MNAR proportion, skews the reported ratings to higher values because of the selective nature of rating behavior in recommender systems.

4.6.2 Imputation Methods

The MovieLens 10M dataset used in the model for this report was missing about 98.9%⁵³ of values due to unrated movies by the users (see Figure 11). Comparative studies have shown that mean imputation performs better at higher percentages of missing values than median imputation [Jadhav et al. 2019]. Comparative studies have also shown that imputation with the average of the user-mean and item-mean results in better accuracy than imputing with either the item-mean or the user-mean alone [Ghazanfar and Prugel 2013]. Furthermore, comparative studies have shown that under conditions of large proportions of missing data, imputation by low-rank matrix completion outperforms imputation by means, nearest neighbors, and regression methods [Liu et al. 2020]. Three common methods of single imputation⁵⁴ and one low-rank matrix completion method of imputation were evaluated for their relative improvements on the RMSE of the test dataset in this report: Zero Imputation, Column Mean Imputation, Average of Row & Column Means Imputation, and Low-Rank Matrix Completion Imputation.

Zero Imputation was used as the baseline imputation method for comparison of the other three imputation methods. Furthermore, the impact of the method of imputation on the subsequent SVD-PCA analysis was also evaluated and is discussed below for each method of imputation. A brief background on this issue is that after an imputation of the dataset is completed, the subsequent SVD-PCA requires a pre-processing step where every element of data (observed data & imputed data) in the complete $m \times n$ rating-residuals matrix R is transformed into zero-mean centered data. In other words, the mean of each column (movie-

⁵²Estimates are as of 17 August 2023, from Firstpost Vantage news segment “The Rise & Rise of Streaming” at 43:23 minutes, in table/figure “‘Endless’ Content on Streaming”, at <https://www.youtube.com/watch?v=K36xPk7bNdW>.

⁵³The 98.9% value was determined by calculation of the percent of NAs in the training dataset Residuals Matrix R (for more background see Section 3.6 SVD-PCA Latent Factor Collaborative Filtering Algorithm, and see the first paragraph in Section 6.1.1 in Appendix A).

⁵⁴*Single imputation* refers to imputation methods that do not rely on a randomized initial set of imputation values, and are therefore not repeated to average out random error.

item) is subtracted from each value in its corresponding column.⁵⁵ This is a standard practice in Principal

⁵⁵See Footnote #40 for more technical details on the pre-processing into zero-mean centered data.




Figure 11: **Heat Map Sample of Missing Data in Rating Matrix of MovieLens Training Data**
Shows Typical 1% Sparsity of Rating Data (prior to imputation)

Component Analysis to reduce the influence of the absolute values in the data, which greatly improves numerical stability while keeping all the variance information of the data intact, as the variance is the key to sorting out the most influential latent factors by SVD-PCA. As expected, the method of imputation affects the means that are determined for each column and the subsequent values remaining, after subtraction of the mean from each data element in its column. This in turn affects the distribution of values in the zero-mean centered data, and is evaluated in the discussions below by Quantile-Quantile (Q-Q) Plot and Density Plot analysis. Q-Q plots are useful to show how the distributions of two different sized populations compare in terms of the spread of their probabilistic proportions (quantiles) of data, where perfect agreement is a Q-Q plot along the $x = y$ identity line of equality that bisects the graph at a $+45$ degree angle. However, Q-Q plots do not display any density information about the data distributions, thus density plots are also included to complete the analysis of the imputed distributions and their anticipated effects on the subsequent covariances and PCA analysis.




Figure 12: Mean-Centered ZERO-Imputation shifts Rating-Residuals Distribution to 93-fold Lower Variance
 Subfigure A: Distribution Statistics, Boxplot & Density of ed_train Rating Residuals (Before Imputation)
 Subfigure B: Distribution Statistics, Boxplot & Density of Zero Imputed edx_train Rating-Residuals
 Subfigure C: Distribution Statistics, Boxplot & Density of Mean-Centered Zero-Imputed Rating-Residuals

4.6.2.1 Zero Imputation is simply replacing the missing rating-residual values with numerical zero. However, the severity of the impact of Zero Imputation, followed by column-mean centering, on the distribution of the rating-residuals is shown in the Q-Q plot of Subfigure D in Figure 13, where the quantile distribution of the Zero Imputed & Mean Centered rating-residuals (y-axis) is plotted against the quantile distribution of the original (not imputed & not mean centered) rating-residuals (x-axis). Zero Imputation shows very poor agreement with the original distribution in a Q-Q plot that only partially aligns on the tail ends of the distributions. Wherever the Q-Q plot is flatter than the $x = y$ identity line, the distribution plotted on the x-axis is more dispersed than the distribution plotted on the y-axis. Indeed, the large flat central region of the Q-Q plot shows no variability at all within the precision of the y-axis dimension, which represents the Zero Imputed distribution after mean centering. A density plot of the Mean-Centered Zero-Imputed Rating-Residuals, in Subfigure C of Figure 12, shows a concentrated high density peak (variance = 0.007) with the central 99%⁵⁶ of the 746,087,406 rating-residual values in the distribution, located in a very narrow range with total width of 0.126 in terms of Movie Rating Residual units on the x-axis. This very narrow range explains the flat central region of the Q-Q plot, in Figure 13-D, from the perspective of a very narrow (0.126 units wide) distribution for the central 99% of Zero-Imputed & Mean-Centered edx_train

⁵⁶The *central 99%* is shorthand for the inter-percentile range between the 0.5th and 99.5th percentiles of the data, and is expressed in the units of the data, i.e., movie rating residual units.

Rating-Residuals on the y-axis. In contrast, a density plot of the original `edx_train` rating-residuals, in Subfigure A of Figure 12, shows a fairly Normal (variance = 0.651) spread-out low density peak with the central 99% of the 8,000,056 rating-residual values in the distribution covering a moderate range with a total

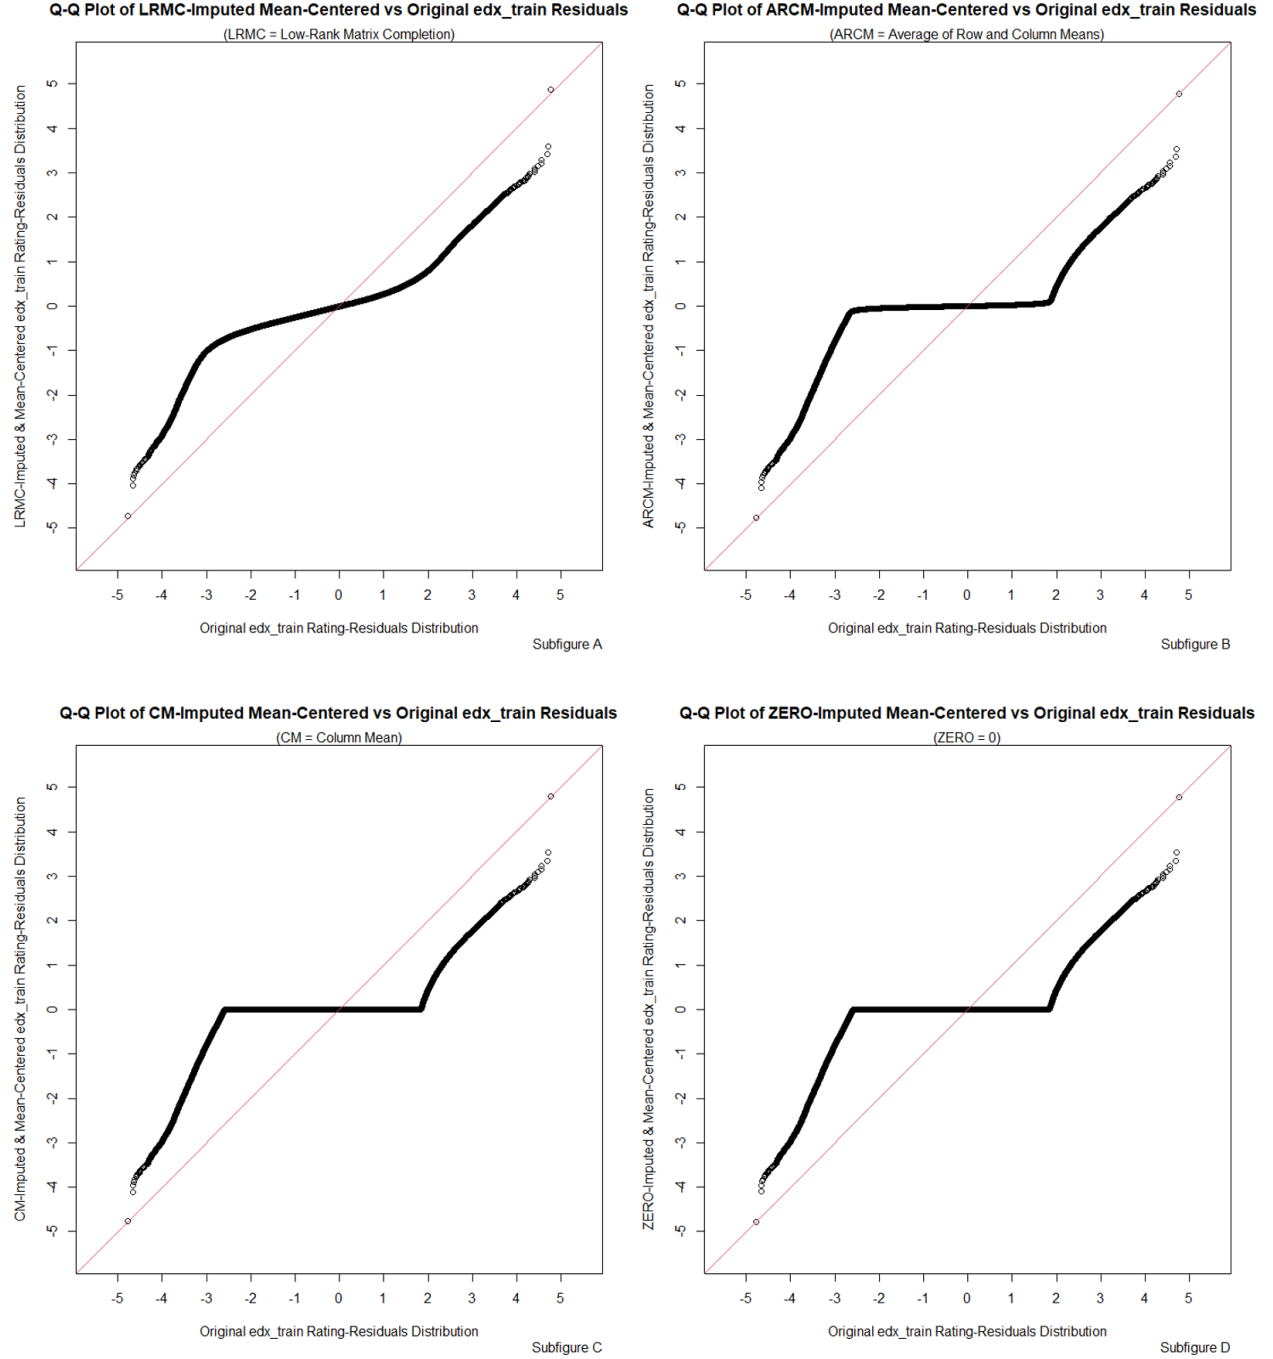


Figure 13: Q-Q Plots of Imputation Distributions **with** Mean Centering against the Original Distribution
 Subfigure A: LRMC Imputed & Mean Centered `edx_train` Rating Residuals vs Original Distribution
 Subfigure B: ARCM Imputed & Mean Centered `edx_train` Rating Residuals vs Original Distribution
 Subfigure C: Column Mean Imputed & Mean Centered `edx_train` Rating Residuals vs Original Distribution
 Subfigure D: ZERO Imputed & Mean Centered `edx_train` Rating Residuals vs Original Distribution

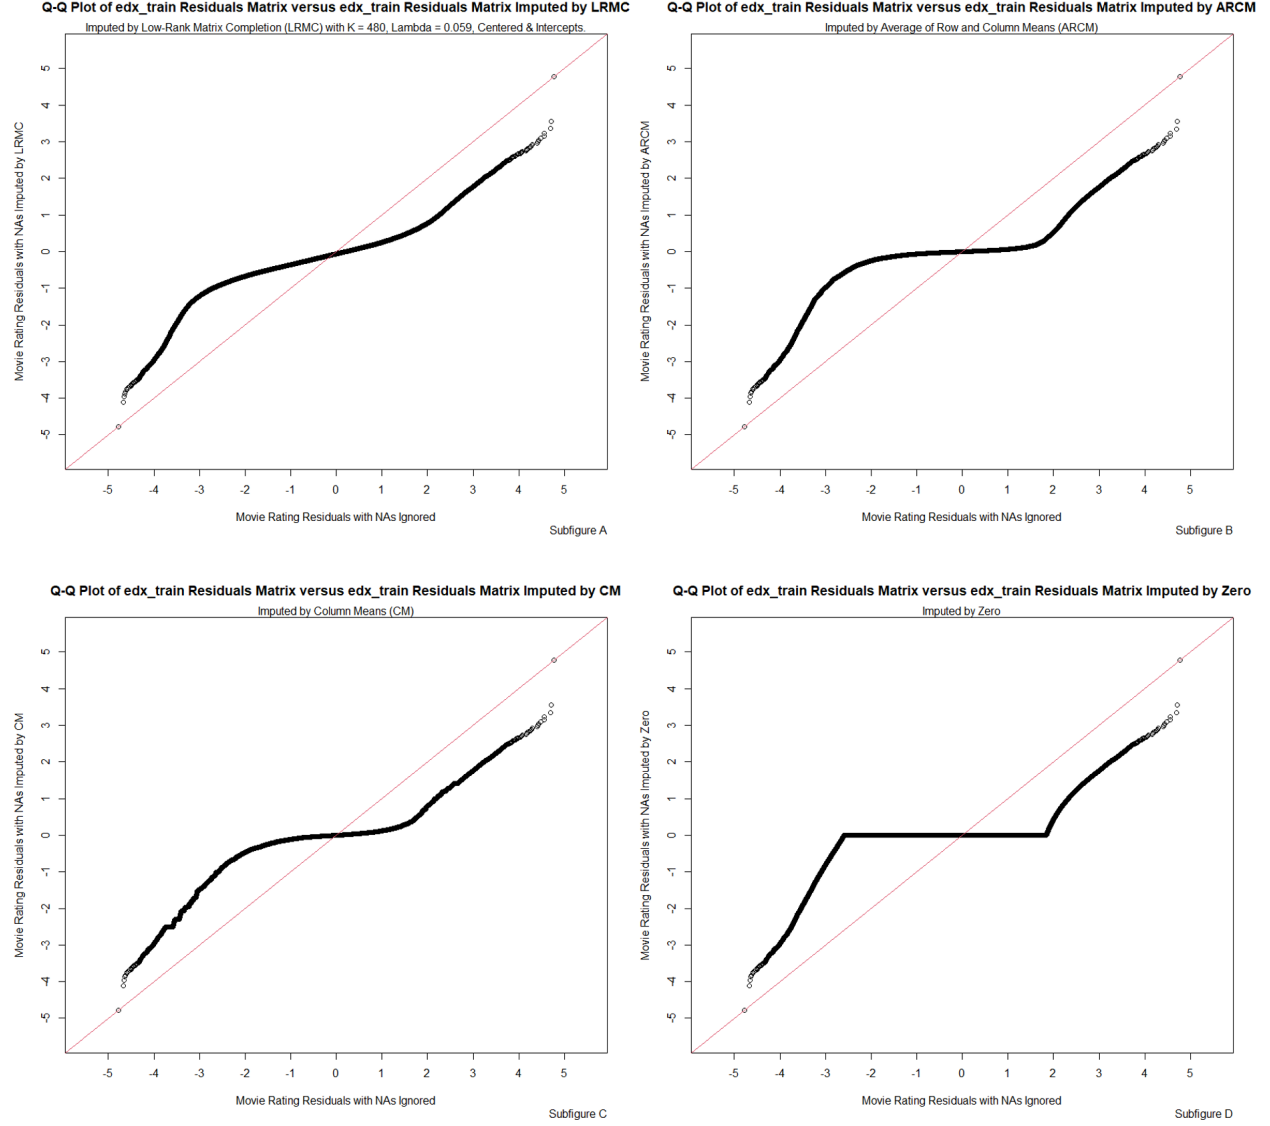


Figure 14: Q-Q Plots of Imputation Distributions **without** Mean Centering against the Original Distribution
 Subfigure A: LRMC Imputed edx_train Rating Residuals vs Original Distribution
 Subfigure B: ARCM Imputed edx_train Rating Residuals vs Original Distribution
 Subfigure C: Column Mean Imputed edx_train Rating Residuals vs Original Distribution
 Subfigure D: ZERO Imputed edx_train Rating Residuals vs Original Distribution

width of 4.5 in terms of Movie Rating Residual units on the x-axis. This moderate range also explains the flat central region of the Q-Q plot, in Figure 13-D, from the perspective of a moderately wide (4.5 units) distribution for the central 99% of Original edx_train Rating-Residuals on the x-axis.

Thus, Zero Imputation impacts the rating-residual data by *severely reducing the range and variance* of the central 99% of the distribution, *on the order of 35 and 90 times less range and variance, respectively* (see Density Plots in Subfigures A, B & C of Figure 12).⁵⁷ In addition, Zero Imputation severely impacts the subsequent SVD-PCA analysis which is focused on the variance and covariance of the data. In other words,

⁵⁷Subfigure B shows the effect of the Imputation Method *alone* and Subfigure C shows the *additional effect* of Mean Centering, which demonstrates that the Imputation Method itself is the dominant influence on the original distribution, and that the Mean Centering contributes to a lesser extent that is dependent on the method of imputation (compare Variance and Central 99 Percentile changes between Subfigures B & C in Figures 12, 15, 16, and 17).

the constancy of the zero values in the imputed positions of each column impacts the SVD-PCA because of the lack of variance in the data in all those positions, which accounts for the overwhelming majority of column positions in a 98.9% sparse matrix. Finally, the end result of Zero Imputation of the 98.9% Missing Values on the subsequent SVD-PCA analysis was a 0.004% higher (worse) RMSE of 0.84644, as shown on the right side of Figure 8 for PC-10677, when compared to the RMSE of 0.84640 from the previous stage in the model development.⁵⁸

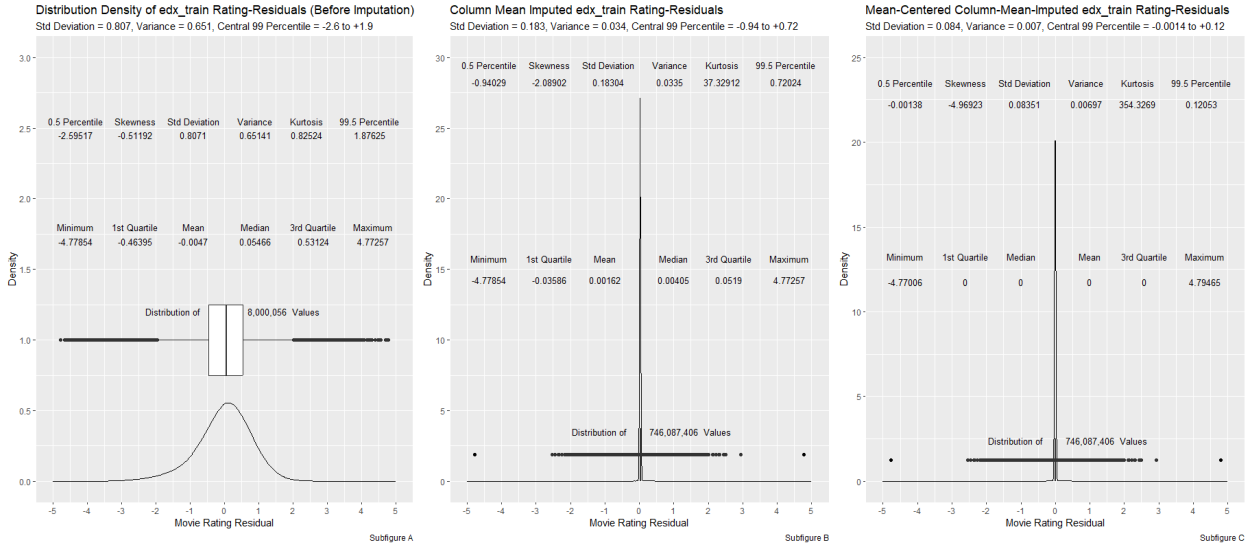


Figure 15: Mean-Centered Col-Mean-Imputation shifts Rating-Residuals Distribution 93-fold Lower Variance
 Subfigure A: Distribution Statistics, Boxplot & Density of edx_train Rating Residuals (Before Imputation)
 Subfigure B: Distribution Statistics, Boxplot & Density of Column Mean Imputed edx_train Rating-Residuals
 Subfigure C: Distribution Statistics, Boxplot & Density of Mean-Centered Col-Mean-Imputed Rating-Residuals

4.6.2.2 Column Mean Imputation is simply replacing the missing rating-residual values in each column (movie ID) with the mean of the non-missing values in that column. However, the severity of the impact of Column Mean Imputation, followed by column-mean centering, on the distribution of the rating-residuals is shown in the Q-Q plot of Subfigure C in Figure 13, where the quantile distribution of the Column Mean Imputed & Mean Centered rating-residuals (y-axis) is plotted against the quantile distribution of the original rating-residuals (x-axis). Similar to Zero Imputation, Column Mean Imputation also shows very poor agreement with the original distribution in a Q-Q plot that only partially aligns on the tail ends of the distributions. Indeed, Column Mean Imputation followed by Mean Centering is mathematically equivalent to Zero Imputation, and is demonstrated in the 1st through 3rd Quartile results listed in Subfigure C of Figure 15. Thus, the same type of analysis as the Zero Imputation (above) applies to Column Mean Imputation with Mean Centering in explaining the large flat central region of its Q-Q plot, in terms of the density plots in Subfigures A & C of Figure 15. Furthermore, a deeper analysis showed that the Q-Q plot of the original distribution plotted against its Column Mean Imputation *without* any subsequent mean centering, resulted in a slightly positive sloped large central region with rounded corners (compare Subfigure C in Figure 13 with Subfigure C in Figure 14). These Q-Q plots demonstrate a case where the Mean Centering makes an extensive contribution to reduction of variance in the distribution prior to SVD-PCA analysis. This is also evidenced by the relatively large reduction in Variance and relatively large reduction in Central 99 Percentile range in the density plot of Mean-Centered Column-Mean-Imputed edx_train Rating-Residuals (compare Subfigures B & C of Figure 15).

Thus, Mean Centering turns Column Mean Imputation into Zero Imputation during the SVD-PCA processing, which impacts the SVD-PCA analysis by severely reducing the range and variance of the central

⁵⁸The RMSEs of the Imputation Methods were compared, in isolation from their Truncated SVD-PCA results, based on the complete set of 10,677 Principal Components generated by their SVD-PCA of the Imputed Data.

99% of the distribution. As in Zero Imputation, the constancy of the zero values in the imputed positions of each column impacts the SVD-PCA because of the lack of variance in the data in all those positions, which accounts for the overwhelming majority of column positions in a 98.9% sparse matrix. Finally, the end result of Column Mean Imputation of the 98.9% Missing Values on the subsequent SVD-PCA analysis was a 0.000% net change with an RMSE of 0.84640, as shown on the right side of Figure 8 for PC-10677, which matches the RMSE of 0.84640 from the previous stage in the model development.⁵⁹

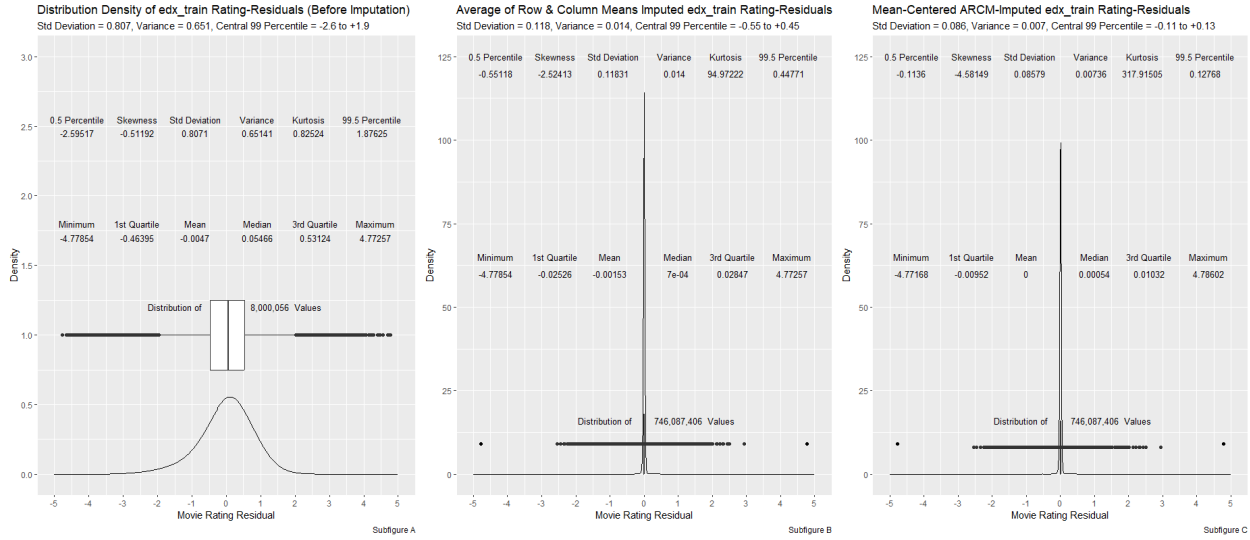


Figure 16: Mean-Centered ARCM-Imputation shifts Rating-Residuals Distribution to 93-fold Lower Variance
 Subfigure A: Distribution Statistics, Boxplot & Density of edx_train Rating Residuals (Before Imputation)
 Subfigure B: Distribution Statistics, Boxplot & Density of ARCM Imputed edx_train Rating-Residuals
 Subfigure C: Distribution Statistics, Boxplot & Density of Mean-Centered ARCM-Imputed Rating-Residuals

4.6.2.3 Average of Row & Column Means Imputation is simply replacing each missing data point (i.e., each NA value in the matrix) with the average of the row and column means in which it was located. However, the severity of the impact of Average of Row & Column Means (ARCM) Imputation, followed by column-mean centering, on the distribution of the rating-residuals is shown in the Q-Q plot of Subfigure B in Figure 13, where the quantile distribution of the ARCM Imputed & Mean Centered rating-residuals (y-axis) is plotted against the quantile distribution of the original rating-residuals (x-axis). Similar to Zero Imputation, ARCM Imputation also shows very poor agreement with the original distribution in a Q-Q plot that only partially aligns on the tail ends of the distributions. Indeed, ARCM Imputation followed by Mean Centering is mathematically equivalent to Zero Imputation in so far as the Column Mean part of ARCM Imputation is concerned. On the other hand, the Row Mean part of the ARCM Imputation produces a slightly positive increase in the slope of the straight and almost flat central region of the Q-Q plot, due to the influence of a different row mean on each imputed position of a column. This increase in variability is also manifested in the nearly doubling of the Central 99 Percentile range to 0.241 in the density plot of the Mean-Centered ARCM-Imputed edx_train Rating-Residuals (Subfigure C in Figure 16) relative to the Central 99 Percentile range of 0.122 in the density plot of Mean-Centered Column-Mean-Imputed edx_train Rating-Residuals (Subfigure C in Figure 15).

Nonetheless, ARCM Imputation with Mean Centering impacts the rating-residual data by severely reducing the range and variance of the central 99% of the distribution, on the order of 19 and 93 times less range and variance, respectively (compare Density Plots in Subfigures A & C of Figure 16), which explains the almost flat central region of the Q-Q plot (Subfigure B in Figure 13) and poor agreement with the original

⁵⁹Based on this analysis, it can also be inferred that the Mean Centering in the previous analysis of Zero Imputation was responsible for the 0.004% higher (worse) RMSE of 0.84644, as opposed to the 0.000% net change in RMSE in the SVD-PCA of Column Mean Imputation which is effectively a Zero Imputation *without* Mean Centering.

distribution. Furthermore, the very low variance of the Mean-Centered ARCM-Imputed rating-residuals impacts the subsequent SVD-PCA analysis which is focused on the variance and covariance of the data. On the other hand, the impact on SVD-PCA is mitigated by a small amount of *variability at the 98.9% scale of imputation* because, as previously stated, the imputed value is different in each position of a column due to the influence of a different row mean on each position of the column. This Row-Mean source of variability at the 98.9% scale allowed ARCM Imputation and the subsequent SVD-PCA analysis resulted in a 0.008% net improvement with an RMSE of 0.84631, as shown on the right side of Figure 8 for PC-10677, when compared to the RMSE of 0.84640 from Column Mean Imputation as well as the previous stage in the model development.

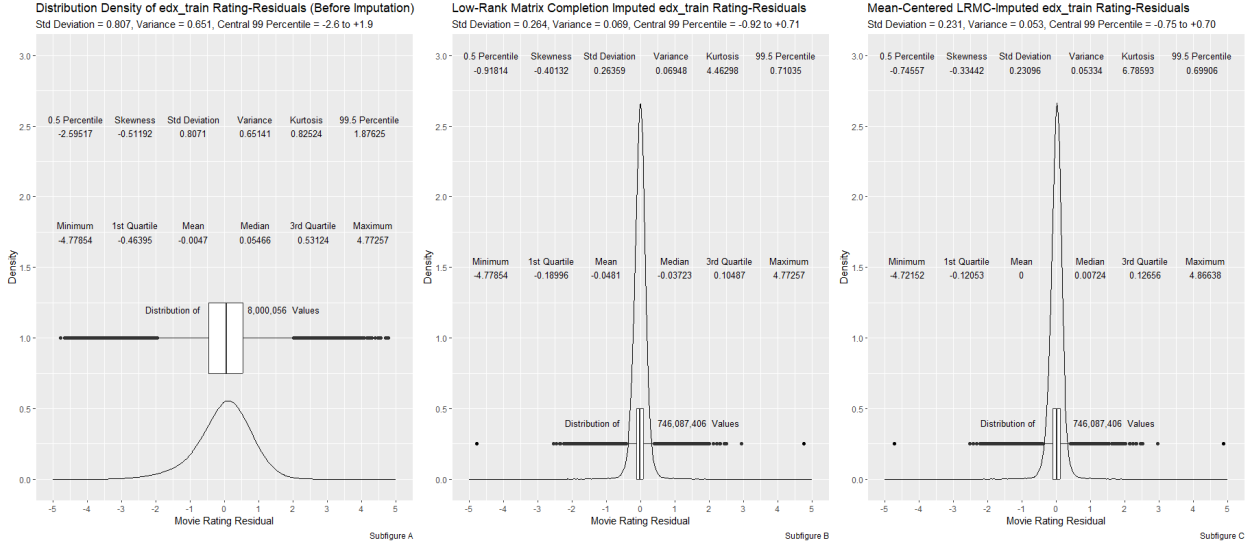


Figure 17: Mean-Centered LRM-Imputation shifts Rating-Residuals Distribution to 12-fold Lower Variance
 Subfigure A: Distribution Statistics, Boxplot & Density of edx_train Rating Residuals (Before Imputation)
 Subfigure B: Distribution Statistics, Boxplot & Density of LRM-Imputed edx_train Rating Residuals
 Subfigure C: Distribution Statistics, Boxplot & Density of Mean-Centered LRM-Imputed Rating-Residuals

4.6.2.4 Low-Rank Matrix Completion Imputation starts with factorizing an incomplete matrix of user-movie rating residuals R into two low-rank feature matrices: user-feature matrix A and movie-feature matrix B , respectively, where a low rank- k of latent features are determined for each user i and movie j by minimizing a L_2 regularized⁶⁰ least-squares loss function with respect to only the entries of R that are known.⁶¹ Minimizing the L_2 regularized least-squares loss function is a non-convex optimization problem for which local minima are found by an Alternating Least-Squares (ALS) algorithm. ALS takes advantage of the fact that when holding one of the low-rank matrices constant, the least-square values for the other matrix can be obtained through a closed-form solution of the system of linear equations. Cholesky Decomposition is the method used in the ALS algorithm for the closed-form numerical solutions of the least-squares fittings. The ALS algorithm alternates between solving one of the two low-rank matrices, A or B , while holding the other low-rank matrix constant, until convergence is reached. At the start of the ALS process, the two low-rank matrices A and B are initialized with random numbers, and then iteratively updated by ALS to decrease the least-squares error with respect to the known entries in R .⁶² After optimizing the low-rank matrices, A and B , their respective row vectors are used to impute the missing entries in R by taking the dot product $\vec{a}_i \cdot \vec{b}_j$ for user i and movie j of each of the missing entries. See Section 6.1 Appendix A for

⁶⁰“ L_2 regularized” refers to the square of the Euclidean (L^2) norm, i.e., $\|\mathbf{x}\|_2^2$, see Ng, Andrew Y. (2004).

⁶¹More specifically, regularization uses Weighted- λ -Regularization that scales (increases) the regularization parameter according to the number of non-missing entries in the data of matrix R for each user i and movie j . This allows for improved imputation by increased numbers of k features and by increased numbers of alternating least-squares iterations, *without overfitting*.

⁶²In addition, for better stability the prediction term in the loss function is centered on the overall mean of matrix R , and corrected in each iteration of ALS fitting for the biases/intercepts for each of the users and movies.

more details, equations and R code for Imputation by Low-Rank Matrix Completion (LRMC).

The impact of Imputation by LRMC, followed by column-mean centering, on the distribution of the rating-residuals is shown in the Q-Q plot of Subfigure A in Figure 13, where the quantile distribution of the LRMC-Imputed & Mean-Centered rating-residuals (y-axis) is plotted against the quantile distribution of the original rating-residuals (x-axis). Compared to the other three methods of imputation, LRMC shows far better agreement with the original distribution in the slope and smoother shape of the central region of the Q-Q plot, due to the influence of a randomly initialized LRMC process and a different optimized value fitted at each imputed position in the matrix. This resulted in a 757% increase in Variance (0.053) in the distribution of Mean-Centered LRMC-Imputed `edx_train` Rating-Residuals compared to the baseline Variance (0.007) in the distribution of Mean-Centered ZERO-Imputed `edx_train` Rating-Residuals (see Subfigure C density plots in Figures 17 and 12). Likewise, this resulted in a 1,150% increase in the Central 99 Percentile range (1.45) in the distribution of Mean-Centered LRMC-Imputed `edx_train` Rating-Residuals compared to the baseline Central 99 Percentile range (0.126) in the distribution of Mean-Centered ZERO-Imputed `edx_train` Rating-Residuals (see Subfigure C density plots in Figures 17 and 12).

LRMC Imputation with Mean Centering impacts the rating-residual data by moderately reducing the range and variance of the central 99% of the distribution, on the order of 3 and 12 times less range and variance, respectively (compare Density Plots in Subfigures A & C of Figure 17). The moderate level of variance in the LRMC Imputed residual-ratings is in accord with the SVD-PCA analysis, which is focused on the variance and covariance of the data. Furthermore, the 98.9% scale of imputation is favorable to the SVD-PCA analysis due to the influence of a randomly initialized LRMC process and a different optimized value fitted at each imputed position in the matrix. Thus, the end result of LRMC Imputation of the 98.9% Missing Values on the subsequent SVD-PCA analysis was a 4.852% net improvement with an RMSE of 0.79496, as shown on the right side of Figure 8 for PC-10677, when compared to the RMSE of 0.84640 from the previous stage in the model development. Indeed, SVD-PCA of the LRMC Imputed rating-residuals was 607 times better than SVD-PCA of the ARCM Imputed rating-residuals (i.e., 4.852% RMSE improvement versus 0.008% RMSE improvement, see Figure 8).

4.7 Limitations

Compared to neural networks, “the linear assumption, acting as the basis of many traditional recommenders, is oversimplified and will greatly limit their modeling expressiveness [Zhang et.al. 2018, p. 6].” Unfortunately this is a reality, and so commercial recommendation systems are at least a hybrid of back-box neural network and transparent traditional machine learning models. This tradeoff in limitations – limited expressiveness of transparent linear models versus limited understandability of black-box neural networks – will eventually give way to highly expressive and understandable neural networks. A relatively new area of neural network research, known as Mechanistic Interpretability, carefully studies and tests neural networks to understand the mathematical strategies the networks are simulating by reverse-engineering learned behaviors into their individual components [Nanda et al. 2023].

As previously discussed, the MovieLens rating distribution is biased toward higher ratings [Marlin et al. 2012; Ghazanfar & Prugel 2013], and is caused by a naturally human selective rating behavior, which affects all long-form content rating systems – MovieLens, Netflix, Amazon Prime, YouTube, etc. Due to the selective nature of rating behavior in recommender systems, use of test datasets drawn at random from the population of reported ratings will invariably have the same higher rating bias as the population, and lead to development of recommendation models that reinforce the selection biases of the users. This in turn leads to a cycle of more reinforced bias in future recommendations for the users. This entrenched bias in recommendations also leads to a lack of freshness and diversity, which are important aspects of real world recommendation systems that meet users’ changes in taste, mood, interests, and novelty.

The training data in typical recommendation systems do not accurately or completely represent the range of items that users don’t like, as opposed to the more accurate and well represented range of items that the users do like. Ideally training data should be a balanced dataset where users rated as many “unliked” items and “neutral” items as “liked” items. Unfortunately this ideal of *explicit* feedback is not realistically practical

on a large enough scale for recommender research, mainly because movies, books and videos are long-form content such that it takes more time to consume and produce ratings. Nonetheless, this is where *implicit* feedback may help to fill in the gaps, such as click tracking and time spent (or not) on items presented at random during browsing activities. Thus, incorporating implicit feedback approaches to help capture a user’s low-rating information should be a part of any well balanced recommendation system.

Lastly, in this project the data splits were done at random for the training, test, and holdout datasets, however in real-world applications, time-based data splits are more realistic scenarios for commercial recommendation systems [Meng 2020; Parashakis 2016]. Thus the performance of the model for this report may seem better than it would actually be in a real-world application.

5 Conclusions

This report thoroughly described the mathematical and practical development of the algorithms used in a latent factor model of collaborative filtering for a recommender system, where accuracy was improved by adjusting for biases that are independent of any latent factor interactions between users and items, as well as by adjusting for a particular interaction bias between users and items, namely the user-specific movie-genre biases. The latter used a Tikhonov-regularized left-inverse matrix (TRLIM) approach, to optimize the bias (coefficients) that projected the user’s movie rating-residuals onto the pseudo-inverse of the user’s genre-elements matrix. This TRLIM approach is based on classical ridge-regression optimization that guarantees best-fit solutions to systems of linear equations that are not invertible. Ridge regression got around these non-invertible issues by using the Tikhonov regularization parameter to add a “ridge” of positive elements on the diagonal of the pseudo-inverse matrix. This ridge is an example of Bias-Variance Tradeoff, where Tikhonov regularization by ridge regression decreased condition number and multicollinearity, increased numerical stability and positive definiteness, and produced a well-conditioned optimized solution in exchange for a tolerable amount of ridge bias.

It is also noteworthy that the improvement in model accuracy from the user-specific movie-genre bias was partly due to inclusion of weighted non-linear scaling with a hyperbolic-tangent (tanh) function in order to further constrain variability in the TRLIM output of user-specific movie-genre bias. It is interesting that the tanh scaling combined with the learned weighting of the TRLIM output is similar to *Batch Normalization* in a neural network. Batch Normalization acts as a form of regularization by reducing variance and smoothing the optimization landscape in a neural network [Santurkar et al. 2018]. Thus, the weighted-tanh constraint of the TRLIM output may be viewed as a Batch Normalization of the interaction biases between users and items. This then leads to the insight that the improvements in model accuracy, from the weighted-tanh TRLIM algorithm for user-specific item-features, are due to λ_T scaled-down & weighted-tanh normalized variances in bias (projection coefficients), which likewise nudge the predictions in the right directions with normalized scaled-down relative biases for an overall improvement in model accuracy.

The TRLIM algorithm is a ridge regression approach that is well established in the literature, and is an approach that may be used to produce an optimized solution for virtually any inconsistent system of linear equations. Furthermore, the weighted-tanh TRLIM algorithm used in the model may be generalized to assign bias to virtually any likely interaction between users and movie item-features, e.g., common interaction biases such as user-specific movie-actor biases, user-specific movie-director biases, user-specific movie-language biases, etc. This has implications for increasing Explainable Machine Learning and AI Transparency by pulling the influence of user-specific item-feature interaction-biases out of the subsequent matrix factorization part of the model, where any remaining rating values are factorized into latent user-item interactions. The exact nature of latent interactions are often difficult to define, which is not the case with assigned user-specific item-feature interactions.

However from an Information Science perspective, there is an ontological arbitrariness in the selection of the user-specific item-feature interactions for assigning biases from the data, as well as an ontological arbitrariness in the order in which the various user-specific item-feature biases are added to the model. Furthermore, the ontological arbitrariness is a caveat for potential benefits in Explainable Machine Learning and AI Transparency attributed to the assigned user-specific item-feature interactions. Nonetheless the common

sense assumption is that, there is practical justification for a given user-specific item-feature bias when *accounting for such a bias leads to improvements in model accuracy*. This *a posteriori* epistemological justification also applies to the order in which various user-specific item-feature biases are added to the model. Likewise, the justification applies to weighted-tanh TRLIM *improvements in model accuracy* – improvements which are very likely when the weighted-tanh TRLIM algorithm is applied to highly-probable or common user-item interactions.

One of the main lessons from this project is that recommendation modeling is really a problem of imputation, which is a very important area for improving the accuracy of the recommender model due to the overwhelming amount of missing ratings for the vast majority of users in these types of models. Thus, sparse data conditions are common in recommender systems, which can cause performance problems, particularly in collaborative based recommender models. Due to the overwhelming proportion of data that imputation affects under sparse conditions, the method of imputation must be carefully chosen and evaluated because of its potential to profoundly impact the accuracy of the model. No method is perfect, each has advantages and disadvantages, and imputation methods that more accurately model the known ratings will give better RMSE results when tested on data sets drawn from the same population of known ratings. This was the case with the four imputation methods compared in this report and evaluated by their Q-Q plots, range & variance of the central 99% of their distributions, as well as their RMSEs on the test dataset. With a much better match in Q-Q plots, much better match in range & variance in the central 99% distributions, and a significant 4.85% RMSE improvement in the model accuracy, Low-Rank Matrix Completion clearly outperformed the three common methods of imputation: zero, column mean, and average of row & column means.

This is in agreement with studies using simulations of missing data values, which have demonstrated usefulness in validating the accuracy of various imputation methods for particular types of data and levels of missingness [Jadhav et al. 2019; Kokla et al. 2019; Liu et al. 2020; Xu et al. 2020]. Results from these studies demonstrated that the choice of imputation depends on the specific characteristics of the data and the research question at hand. For example, if there are less than 5% missing values, those data points can be ignored, otherwise use mean or median imputation. If there are 5% to 20% missing values, then regression or kNN methods of imputation can work well depending on the type of missing data (e.g., use time series regression for time series data). However, the studies showed that if there are more than 20% missing values then Low-Rank Matrix Completion (LRMC) should be used for imputation [Liu et al. 2020; Xu et al. 2020].

Finally, it is important to understand that the column-mean centering of the data required by PCA amplifies the variance reduction problems of low variance imputation methods such as zero imputation, column mean imputation, and average of row & column means imputation. Furthermore, during model development it was interesting to observe that the more variance there was in the imputed values, the more variance that was explained by fewer Principal Components (PCs) in the SVD-PCA factorization of the rating residuals. Likewise, imputation by LRMC greatly contributed to increasing the variance explained in fewer PCs. Another interesting observation is that based on the mathematical process of LRMC imputation (in Discussion and Appendix A) LRMC results are expected to align with the distribution of the original (non-missing) data, including alignment with the higher rating bias of the reported (non-missing) values. Nonetheless, higher rating bias introduced by LRMC imputation is limited to the latent-factors component of the model.

5.1 Future Development Activities

It bears repeating that one of the main lessons from this project is that recommendation modeling is really a problem of imputation, and sophisticated methods such as Low-Rank Matrix Completion (LRMC) are key tools for handling this problem in sparse datasets. LRMC is a versatile tool that can also be used with content-based recommender models as well as with implicit-feedback recommender models. Furthermore, LRMC can *extend* explicit-feedback recommender models by incorporating additional matrices of user and item side information into the user-item rating factorization. The resulting latent features explain both the rating data as well as side information, such as user age or occupation and movie genre or year of release, thereby making the recommendation model generalize better to unseen data.⁶³

⁶³For more information see R documentation in *cmfrec* package vignette: “Matrix Factorization with Side Info.”

Conceptually, imputation is akin to extrapolation. LPMC imputation of unknown ratings using explicit-feedback can be thought of as an extrapolation that minimizes a cost function against the known ratings. In addition, LPMC imputation of unknown ratings using implicit-feedback or content information can also be thought of as extrapolations that minimize cost functions against those known values. Furthermore, results from multiple LPMC imputations using explicit-feedback, implicit-feedback, and content information can be joined together in a hybrid of weighted imputations that jointly minimizes the RMSE.

At the time that the user-aware movie-genre bias algorithm was developed, it was computationally prohibitive to optimize the TRLIM parameter λ_T (in the alternating iterative optimizations with the weight parameter w) at the 68,548 scale of users in the test dataset. Thus λ_T was optimized using just the first 4000 users in the test dataset. Since then, the computer system used to develop the algorithm was significantly upgraded to 8 CPU (16 Cores) with 128 GB RAM and 8 TB SSD, which can better handle several iterative cycles of 68,548 TRLIM calculations with efficient 8-threaded parallel processing, and optimize the λ_T parameter at the the full scale of users in the test dataset.

Finally, as concluded above, there are several likely interactions between users and movie item features that may be assigned biases by the weighted-tanh TRLIM algorithm, to increase the recommender model accuracy and its explainable machine learning. Furthermore, as previously discussed, there are also several untapped temporal biases, hidden within the movie biases and the user biases, that must be uncovered through time-based algorithms in order to maximize their potential for improving the recommender model accuracy; for example, algorithms to uncover time-aware user-bias, time-aware genre-bias, time-aware actor-bias, time-aware producer-bias, etc.

6 Appendices

6.1 Appendix A

The following was compiled, and adapted for use with this project, from the "cmfrec" package (ver.3.5.1-1) documentation in R (ver.4.3.0), from supporting literature citations referenced in the package, and from relevant Wikipedia articles. All references are cited in the text below.

6.1.1 Imputation by Low-Rank Matrix Completion with cmfrec Package in R

Imputation by Low-Rank Matrix Completion (LRMC) starts with factorizing the incomplete $n_u \times n_m$ matrix of user-movie rating residuals $\mathbf{R} = \{r_{ij}\}_{n_u \times n_m}$ into two low-rank $n_u \times n_k$ and $n_m \times n_k$ feature matrices: user-feature matrix $\mathbf{A} = \{a_{ik}\}_{n_u \times n_k}$ and movie-feature matrix $\mathbf{B} = \{b_{jk}\}_{n_m \times n_k}$ respectively, where each element r_{ij} represents the residual rating score of movie j that was rated by user i , with its value either being a real number (\mathbb{R}) or missing (NA):

$$\mathbf{R} \approx \mathbf{AB}^T \quad (\text{A.1})$$

The number of users is designated by n_u , the number of movies by n_m , and the number of features by n_k , where $n_k \ll \min\{n_u, n_m\}$. Thus, factorization by LRMC determines n_k latent features for each user i and movie j by minimizing a L_2 regularized⁶⁴ least-squares loss function with respect to only the entries of \mathbf{R} that are known. More specifically, regularization uses Weighted- λ -Regularization that scales (increases) the λ regularization parameter for each row vector \mathbf{a}_i and \mathbf{b}_j , in matrices \mathbf{A} and \mathbf{B} , according to the number of non-missing entries in the corresponding data of matrix \mathbf{R} for each user i and movie j . This scaled regularization greatly improves imputation, *without overfitting*, by allowing for increased numbers of features (k)⁶⁵ and for increased numbers of alternating least-squares iterations (niter), in the CMF function of the cmfrec package (see Figure 18 and Section 6.1.3 R-code for Imputation by LRMC).⁶⁶

$$(\mathbf{A}, \mathbf{B}) = \arg \min_{(\mathbf{A}, \mathbf{B})} \frac{1}{n} \sum_{(i,j) \in I_R} (r_{ij} - \mathbf{a}_i \mathbf{b}_j^T)^2 + \lambda \left[\sum_i n_{u_i} \|\mathbf{a}_i\|_2^2 + \sum_j n_{m_j} \|\mathbf{b}_j\|_2^2 \right] \quad (\text{A.2})$$

where I_R is the index set of the known ratings in \mathbf{R} , and n is the size of I_R .

Minimizing the L_2 regularized least-squares loss function in Equation A.2 is a non-convex optimization problem for which local minima are found by an Alternating Least-Squares (ALS) algorithm. ALS takes advantage of the fact that when holding one of the low-rank matrices constant, the least-square values for the other matrix can be obtained through a closed-form solution of the system of linear equations. Cholesky Decomposition was the method used, in the ALS algorithm of the CMF function, for the closed-form numerical solutions of the least-squares fits, because of its better numerical stability and precision than the faster conjugate gradient method (in the CMF function).⁶⁷

The ALS algorithm alternates between solving one of the two low-rank matrices, \mathbf{A} or \mathbf{B} , while holding the other low-rank matrix constant, until convergence is reached. At the start of the ALS process, the two low-rank matrices \mathbf{A} and \mathbf{B} are initialized with random numbers, and then iteratively updated by ALS to decrease the least-squares error with respect to the known entries in \mathbf{R} . In addition, for better stability the prediction term $\mathbf{a}_i \mathbf{b}_j^T$ is centered by subtracting the overall mean μ of the matrix \mathbf{R} , and by correcting for (subtracting) the biases/intercepts in \mathbf{R} for each of the users as parameter c_i ($i = 1, \dots, n_u$) and each

⁶⁴“ L_2 regularized” refers to the square of the Euclidean (L^2) norm, i.e., $\|\mathbf{x}\|_2^2$, see Ng, Andrew Y. (2004).

⁶⁵For clarity in this paper, the term “feature” is used with LRMC factorization and the term “factor” is used with SVD-PCA factorization, otherwise the two terms are completely synonymous.

⁶⁶For more information refer to CMF function documentation in cmfrec package (ver.3.5.1-1) of R (ver.4.3.0). For more information on Weighted- λ -Regularization refer to Zhou et al. (2008). For more information on “Matrix Completion” refer to Wikipedia Contributors (2023, March 20).

⁶⁷For more information refer to CMF function documentation in cmfrec package (ver.3.5.1-1) of R (ver.4.3.0), and refer to “Cholesky Decomposition” by Wikipedia Contributors (2023, February 13).

of the movies as parameter d_j ($j = 1, \dots, n_m$), which are updated as vectorized bias/intercept parameters, $\text{user_bias} = \langle c_1, \dots, c_{n_u} \rangle$ and $\text{item_bias} = \langle d_1, \dots, d_{n_m} \rangle$, at each iteration of ALS fitting in the CMF function:⁶⁸

$$(\mathbf{A}, \mathbf{B}) = \arg \min_{(\mathbf{A}, \mathbf{B})} \frac{1}{n} \sum_{(i,j) \in I_R} [r_{ij} - (\mathbf{a}_i \mathbf{b}_j^T + \mu + c_i + d_j)]^2 + \lambda \left[\sum_i n_{u_i} \|\mathbf{a}_i\|_2^2 + \sum_j n_{m_j} \|\mathbf{b}_j\|_2^2 \right] \quad (\text{A.3})$$

After optimizing the low-rank feature matrices, \mathbf{A} and \mathbf{B} , it is then possible to use their respective row vectors to impute the missing entries of \mathbf{R} by taking the dot product $\mathbf{a}_i \bullet \mathbf{b}_j$ for user i and movie j of each of the missing entries. This can be best accomplished by supplying the output of the CMF function (\mathbf{A}, \mathbf{B}) and the original matrix with missing values \mathbf{R} as arguments in the *imputeX* function of the *cmfrec* package in R.⁶⁹

6.1.2 Test RMSE Tuning of CMF Parameters, Lambda & k Features, in LRMC of Matrix R

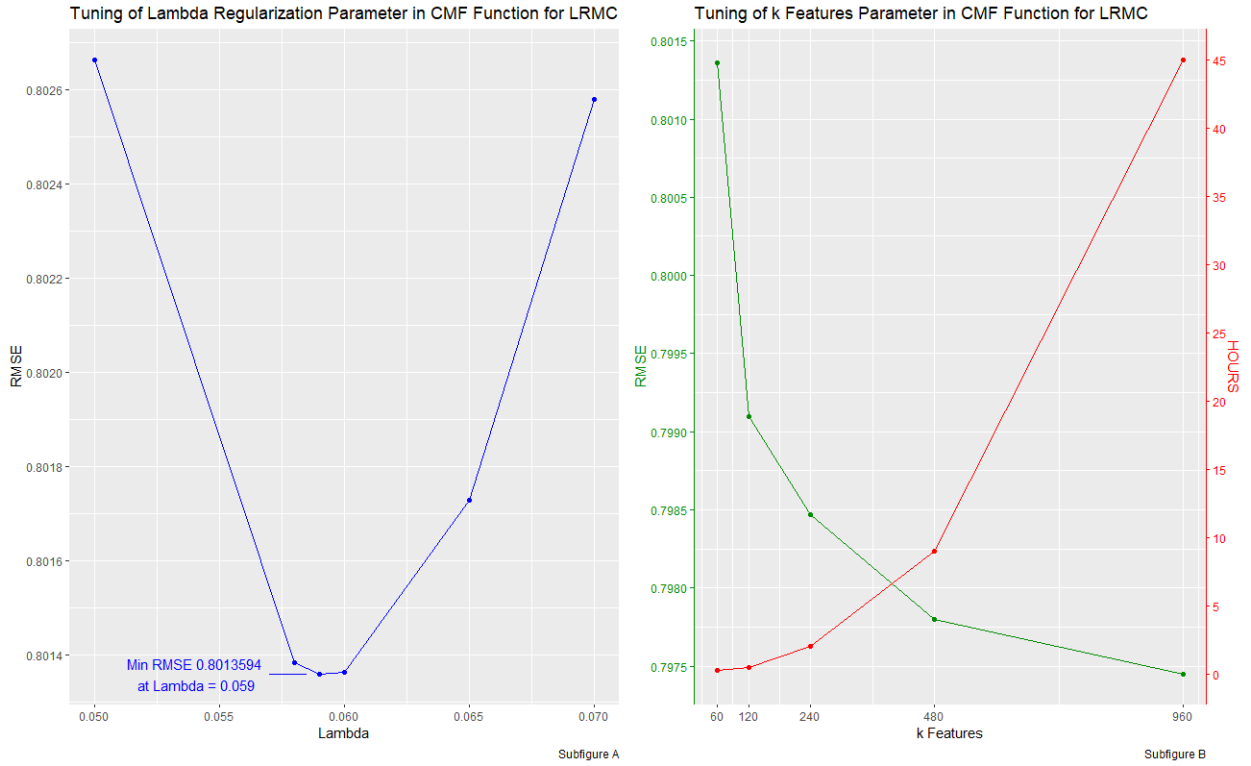


Figure 18: **Optimum CMF Parameters Lambda = 0.059 & k = 480 Features (9-Hour Processing Time)**
Subfigure A: Test RMSE Tuning of Lambda Regularization Parameter in CMF Function for LRMC
Subfigure B: Test RMSE Tuning of k Parameter in CMF Function for LRMC & HOURS for Processing

⁶⁸For more information refer to *cmfrec_vignette* “Matrix Factorization with Side Info” in the *cmfrec* package (ver.3.5.1-1) of *R* (ver.4.3.0), and refer to Cortes (2018) for deeper theoretical discussions.

⁶⁹For more information refer to Section 6.1.3 and to the *imputeX* function documentation in the *cmfrec* package (ver.3.5.1-1) of *R* (ver.4.3.0).

6.1.3 R-code for Imputation by Low-Rank Matrix Completion (LRMC)

```
# Load R package "cmfrec" (Collective Matrix Factorization for Recommender Systems)
library(cmfrec)

residuals_matrix_na <- residuals_matrix # Preserve original residuals_matrix (with NAs) by making copy as residuals_matrix_na
residuals_matrix_na[1:10, 1:10] # 10 x 10 sample of the Residuals Matrix
      122      292      316      329      355      356      362      364      370      377
1  0.7068175  0.1312068  0.3464287  0.2705741  0.8496777 -0.4190664  0.1155877 -0.06079595  0.4322812 -0.1062309
2      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
3      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
4      NA -1.1969667  1.0173922  0.8217935      NA      NA      NA  0.25059994      NA -1.1518907
5      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
6      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
7      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
8      NA  0.0262706      NA      NA      NA      NA      NA -0.54480292 -1.6222095 -0.4571868
9      NA      NA      NA      NA      NA      NA      NA      NA      NA      NA
10     NA      NA      NA      NA      NA      NA -1.1483094      NA      NA      NA

# Imputation of NAs by Low-Rank Matrix Completion (LRMC) uses the Collective Matrix Factorization for Recommender Systems package
# cmfrec, which contains the CMF() function to factorize the original sparse data matrix X (i.e., R) into a LRMCmodel that consists
# of two low-rank matrices of user features (matrix A) and movie-item features (matrix B). The CMF() function assumes that X is a
# sparse matrix in which users represent rows, items represent columns, and the non-missing values denote explicit feedback movie
# ratings from users on items. LRMC imputation of the sparse data matrix X is completed by the imputeX() function in the cmfrec
# package, which takes the LRMCmodel and the sparse matrix X as arguments. Optimized CMF() parameters, for the LRMCmodel to minimize
# RMSE on the edx_test dataset, are set at: k = 480, lambda = 0.059, niter = 30, center = TRUE, user_bias = TRUE, and item_bias = TRUE.

      # LRMCmodel = (A, B) output of the CMF function
LRMCmodel <- CMF( # CMF function takes about 9 hours to process a 69878 rows x 10677 cols matrix when k = 480 and niter = 30
  X = residuals_matrix_na, # residuals_matrix_na is an identical working copy of the original residuals_matrix (with NAs)
  k = 480, # feature dimension of the low-rank factorization for user matrix A and item matrix B (where k_max = 10677)
  lambda = 0.059, # L2-norm regularization parameter optimized at 0.059 for minimum edx_test RMSE (tuned for scale_lam = TRUE)
  method = "als", # alternating least-squares optimization to fit LRMCmodel (faster and more memory efficient than gradient)
  use_cg = FALSE, # defaults to Cholesky algorithm to solve closed-form least squares, instead of less exact/stable gradient
  user_bias = TRUE, # adds user (row) intercepts to the LRMCmodel, which improves stability and RMSE of edx_test predictions
  item_bias = TRUE, # adds item (column) intercepts to the LRMCmodel, which improves stability and RMSE of edx_test predictions
  center = TRUE, # centers X data by subtracting mean values in the LRMCmodel, which improves stability and RMSE of edx_test
  scale_lam = TRUE, # increases lambda for each row in A & B according to number of non-missing entries in the X data for that row
  niter = 30, # number of ALS iterations, better accuracy & lower RMSE with more iterations, balanced against time required
  finalize_chol = TRUE, # performs final iteration with Cholesky solver, improves accuracy and is consistent with use_cg = FALSE
  NA_as_zero = FALSE, # default setting = FALSE = do not take missing entries in X as zeros (to not interfere with imputation)
  nonneg = FALSE, # default = FALSE = do not constrain LRMCmodel to be non-negative (residuals_matrix_na contains negative values)
  precompute_for_predictions = TRUE, # default = TRUE (to prepare LRMCmodel for use in imputeX function)
  verbose = TRUE, # monitor LRMCmodel progress through each of the 30 iterations of fitting low-rank matrices A and B
  handle_interrupt = TRUE, # default = TRUE (to stop model with usable fitted model object if interrupt is necessary)
  seed = 1, # seed for random number generation in the initial matrices of A & B, for reproducibility of imputed missing values
  nthreads = parallel::detectCores()/2 # number of parallel threads for parallel processing in the imputeX() function step below
) # ALS method in the CMF() function does NOT use parallel processing, so nthreads ignored

# Impute the NAs in the Residuals Matrix with the LRMCmodel object using parallel processing (should only take about 15 minuts)
residuals_matrix_LRMCimputed <- imputeX(LRMCmodel, residuals_matrix_na, nthreads = LRMCmodel$info$nthreads)

residuals_matrix_LRMCimputed[1:10, 1:10] # 10 x 10 sample of the Imputed Residuals Matrix (compare with Residuals Matrix above)
      122      292      316      329      355      356      362      364      370      377
1  0.70681746  0.1312068  0.34642868  0.27057414  0.8496770 -0.41906637  0.115586992 -0.06079595  0.43228125 -0.10623087
2 -0.11275461 -0.2859850 -0.39728850 -0.50004171 -0.1865108 -0.03442356 -0.117905348 -0.03537676 -0.15135524 -0.11329069
3  0.08737603  0.1800540  0.16998775  0.21741491  0.3397911  0.05116116  0.080488352  0.11376302  0.01929575  0.20284162
4 -0.02538663 -1.1969667  1.01739215  0.82179351  0.4597019  0.12572960  0.089626581  0.25059994 -0.11332941 -1.15189069
5 -0.30853230 -0.4348464 -0.60586481 -0.43799718 -0.3627595 -0.77151677 -0.100173464 -0.75742333  0.12018838 -0.61945370
6 -0.25407688 -0.1099476 -0.40221712 -0.30397864 -0.4470596  0.08707153 -0.032594767 -0.06973103 -0.32251466 -0.07546702
7 -0.17820901 -0.3433118 -0.42807820 -0.36384573 -0.3206117 -0.31544016 -0.191604615 -0.31578820  0.01758948 -0.11240441
8 -0.06695511  0.0262706  0.06276693 -0.07132330  0.1203327 -0.20883589 -0.056004616 -0.54480292 -1.62220947 -0.45718676
9 -0.50288924 -0.3414170 -0.07655912  0.02918981 -0.2683232 -0.48168915  0.060723537 -0.46408739 -0.04962225 -0.47611516
10 -0.05832001 -0.1053899  0.03464984  0.12593818 -0.0281441 -1.14830938 -0.001208181 -0.05544321  0.05254412 -0.10951058

dim(residuals_matrix_LRMCimputed) # Imputed Residuals Matrix is full-size with correct dimensions
[1] 69878 10677

sum(is.na(residuals_matrix_LRMCimputed)) # confirmed that no NAs remain in the Imputed Residuals Matrix
[1] 0
```


6.2 Appendix B

6.2.1 Derivation of Partial Derivative of Squared Error Cost Function with respect to λ_t for Prediction Model $\hat{r}_i = \mu + b_i(t)$ from Total Movie Bias $b_i(t) = b_i + b_{i,t}$

DESCRIPTION	MODEL FUNCTION	PARTIAL DERIVATIVE
Root Mean Squared Error	$\text{RMSE} = \sqrt{\frac{1}{n} \sum_1^n (\hat{r}_i - r_i)^2} = \sqrt{\text{MSE}}$	Not Required*
Mean Squared Error	$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (\hat{r}_i - r_i)^2 = \frac{1}{n} \sum_{i=1}^n \text{SE}_i$	Not Required*
Squared Error	$\text{SE}_i = (\hat{r}_i - r_i)^2$	$\frac{\partial \text{SE}_i}{\partial \hat{r}_i} = 2(\hat{r}_i - r_i)$
Prediction Model	$\hat{r}_i = \mu + b_i(t)$	$\frac{\partial \hat{r}_i}{\partial b_i(t)} = 1$
Total Movie Bias	$b_i(t) = b_i + b_{i,t}$	$\frac{\partial b_i(t)}{\partial b_{i,t}} = 1$
Static Movie Bias	$b_i = \frac{\sum_1^{n_i} (r_i - \mu)}{n_i + \lambda}$	Not Required**
Temporal Movie Bias	$b_{i,t} = \frac{\sum_1^{n_{i,t}} r_{i,t}}{n_{i,t} + \lambda_t} - \mu_i$	$\frac{\partial b_{i,t}}{\partial \lambda_t} = -\frac{\sum_1^{n_{i,t}} r_{i,t}}{(n_{i,t} + \lambda_t)^2}$

Total Chain of

**Partial Derivatives for
Squared Error w.r.t. λ_t**

$$\frac{\partial \text{SE}_i}{\partial \lambda_t} = \frac{\partial \text{SE}_i}{\partial \hat{r}_i} \cdot \frac{\partial \hat{r}_i}{\partial b_i(t)} \cdot \frac{\partial b_i(t)}{\partial b_{i,t}} \cdot \frac{\partial b_{i,t}}{\partial \lambda_t} = -2(\hat{r}_i - r_i) \frac{\sum_1^{n_{i,t}} r_{i,t}}{(n_{i,t} + \lambda_t)^2} \quad (\text{B.1})$$

* The calculus of deriving partial derivatives involving the cost function is greatly simplified by considering just the squared error term, which is after all, what makes the cost function, *a cost function*.

** The static movie bias does not contain any terms that involve λ_t and therefore the partial derivative of the static movie bias is not required.

Overview of Algorithm: Equation B.1 is the partial derivative of the Squared Error (SE) of the total movie bias model, with respect to λ_t , and was used to determine the *slope* of the SE_i in the prediction of each i movie rating \hat{r}_i across all movies in the training dataset. The slopes of the squared errors of all the movies in the training dataset were averaged together to report the *mean slope of the squared error* for the λ_t value that was used in all the slope calculations. Calculation of the Mean Slope of the Squared Error (MSSE) was repeated across a range of closely spaced intervals of λ_t , in order to find the λ_t value with a MSSE that was closest to zero slope, as well as No More Than ± 0.00001 from zero slope. The algorithm used to find this tuned λ_t value is presented in Sections 6.2.2 and 6.2.3: *R-code for Calculation of Mean Slope of Squared Error Cost Function from Partial Derivative with respect to λ_t for Prediction Model from Total Movie Bias*.

6.2.2 R-code for Calculation of Mean Slope of Squared Error Cost Function from Partial Derivative with respect to λ_t for Prediction Model from Total Movie Bias

```
# Calculation of Mean Slope of Squared Error Cost Function from Partial Derivative with respect to Lambda_t
# for Prediction Model from Total Movie Bias.

library(tidyverse)
library(data.table)

# Determine the average rating of each movie in edx_train:
edx_train_movie_rating_avg <- edx_train %>% group_by(movieId) %>% summarise(movie_rating_avg = mean(rating))

edx_train_movie_rating_avg
# A tibble: 10,677 x 2
  movieId movie_rating_avg
  <dbl>         <dbl>
1     1             3.93
2     2             3.21
3     3             3.14
4     4             2.86
5     5             3.06
6     6             3.82
7     7             3.36
8     8             3.14
9     9             3.01
10    10            3.43
# ... with 10,667 more rows

# Partial Derivative of Squared Error of r_hat_i Prediction from Total Movie Bias Model with respect to Lambda_t:
#      dSEi/dLt = -2*(r_hat_i - r_i)*sum(r_it)/(n_it + L_t)^2          (From Equation B.1 in Appendix B)

# 3-Way Lambda_FMR_GenreUserMovieBias_edx_train = 3.973333
# Cross Validated/Full Model Regularization Parameter (CVR/FMR)
Lambda_m <- Lambda_FMR_GenreUserMovieBias_edx_train # Lambda_m applied for regularization of Static Movie Bias.
Lambdas_t <- seq(-0.50, 0.50, 0.01) # Range and precision of Lambda_t tested to pinpoint value where mean slope is
# no more than ±0.00001 from zero slope result of partial derivative of
# squared error of cost function, for prediction model from total movie bias.
# mu_hat_edx_train <- 3.51242785050505 # Global Mean Rating of All Movies in Training Dataset

# Algorithm for calculation of Mean Slope of Squared Error of Total Movie Bias Prediction Model w.r.t. Lambda_t:
dSEiLt_train <- sapply(Lambdas_t, function(L_t){
  reg_movie_bias <- edx_train_date %>% # Lambda_t parameter is tuned on entire Training dataset
  group_by(movieId) %>% # that also includes bin_2_year date data.
  summarize(b_m_reg = sum(rating - mu_hat_edx_train)/(n() + Lambda_m)) # regularized static movie bias
  reg_2yr_avg_rating <- edx_train_date %>%
  group_by(movieId, bin_2_year) %>%
  summarise(reg_2yr_avg_rating = sum(rating)/(n() + L_t)) # regularized part of temporal movie bias
  r_hat_i <- edx_train_date %>%
  left_join(reg_movie_bias, by = 'movieId') %>%
  left_join(reg_2yr_avg_rating, by = c('movieId', 'bin_2_year')) %>%
  left_join(edx_train_movie_rating_avg, by = 'movieId') %>%
  mutate(r_hat_i = mu_hat_edx_train + b_m_reg + reg_2yr_avg_rating - movie_rating_avg) %>%
  select(userId, movieId, bin_2_year, r_hat_i) # total movie bias model prediction r_hat_i
  regSqrd_2yr_avg_rating <- edx_train_date %>%
  group_by(movieId, bin_2_year) %>%
  summarise(regSqrd_2yr_avg_rating = sum(rating)/(n() + L_t)^2) # sum(r_it)/(n_it + L_t)^2 term of dSEi/dLt
  SEiLt_df <- edx_train_date %>%
  left_join(r_hat_i, by = c('userId', 'movieId', 'bin_2_year')) %>%
  left_join(regSqrd_2yr_avg_rating, by = c('movieId', 'bin_2_year')) %>%
  mutate(SEiLt = -2 * (r_hat_i - rating) * regSqrd_2yr_avg_rating) # Individual Slope of Squared Error
  SEiLt_avg <- SEiLt_df %>% mutate(SEiLt_avg = sum(SEiLt)/n()) %>% pull(SEiLt_avg) %>% first()
  return(SEiLt_avg) # Mean Slope of Squared Error
})
```

6.2.3 R-code (continued) for Calculation of Mean Slope of Squared Error Cost Function from Partial Derivative with respect to λ_t for Prediction Model from Total Movie Bias

```
# Data frame of Mean Slope of Squared Error results from range of Lambda_t values tested:
dSEidLt_train_df <- data.frame(Lambda_t = seq(-0.50, 0.50, 0.01), MeanSlope_SquareError = dSEidLt_train)

# Minimum of Absolute Value of MeanSlope_SquareError = Optimized Value of Lambda_t Regularization Parameter:
Lt <- dSEidLt_train_df$Lambda_t[which.min(abs(dSEidLt_train_df$MeanSlope_SquareError))]
Lt
[1] 0.05

# DISPLAY OF DATA FROM dSEidLt_train_df DATAFRAME, FOR PLOT FIGURE 4B OF REPORT:
dSEidLt_train_df
```

	Lambda_t	MeanSlope_SquareError		Lambda_t	MeanSlope_SquareError
1	-0.50	-4.024594e-02	51	0.00	-7.746359e-04
2	-0.49	-3.756881e-02	52	0.01	-6.056021e-04
3	-0.48	-3.510317e-02	53	0.02	-4.425634e-04
4	-0.47	-3.282858e-02	54	0.03	-2.852265e-04
5	-0.46	-3.072688e-02	55	0.04	-1.333156e-04
6	-0.45	-2.878194e-02	56	0.05	1.342904e-05
7	-0.44	-2.697938e-02	57	0.06	1.552521e-04
8	-0.43	-2.530637e-02	58	0.07	2.923840e-04
9	-0.42	-2.375141e-02	59	0.08	4.250420e-04
10	-0.41	-2.230422e-02	60	0.09	5.534312e-04
11	-0.40	-2.095554e-02	61	0.10	6.777452e-04
12	-0.39	-1.969707e-02	62	0.11	7.981666e-04
13	-0.38	-1.852131e-02	63	0.12	9.148684e-04
14	-0.37	-1.742148e-02	64	0.13	1.028014e-03
15	-0.36	-1.639148e-02	65	0.14	1.137757e-03
16	-0.35	-1.542577e-02	66	0.15	1.244245e-03
17	-0.34	-1.451932e-02	67	0.16	1.347616e-03
18	-0.33	-1.366758e-02	68	0.17	1.448001e-03
19	-0.32	-1.286638e-02	69	0.18	1.545525e-03
20	-0.31	-1.211196e-02	70	0.19	1.640306e-03
21	-0.30	-1.140086e-02	71	0.20	1.732455e-03
22	-0.29	-1.072994e-02	72	0.21	1.822079e-03
23	-0.28	-1.009632e-02	73	0.22	1.909279e-03
24	-0.27	-9.497366e-03	74	0.23	1.994151e-03
25	-0.26	-8.930656e-03	75	0.24	2.076786e-03
26	-0.25	-8.393978e-03	76	0.25	2.157271e-03
27	-0.24	-7.885295e-03	77	0.26	2.235689e-03
28	-0.23	-7.402731e-03	78	0.27	2.312117e-03
29	-0.22	-6.944564e-03	79	0.28	2.386632e-03
30	-0.21	-6.509202e-03	80	0.29	2.459304e-03
31	-0.20	-6.095178e-03	81	0.30	2.530202e-03
32	-0.19	-5.701137e-03	82	0.31	2.599389e-03
33	-0.18	-5.325825e-03	83	0.32	2.666929e-03
34	-0.17	-4.968083e-03	84	0.33	2.732880e-03
35	-0.16	-4.626837e-03	85	0.34	2.797298e-03
36	-0.15	-4.301089e-03	86	0.35	2.860238e-03
37	-0.14	-3.989917e-03	87	0.36	2.921750e-03
38	-0.13	-3.692460e-03	88	0.37	2.981885e-03
39	-0.12	-3.407921e-03	89	0.38	3.040688e-03
40	-0.11	-3.135556e-03	90	0.39	3.098206e-03
41	-0.10	-2.874674e-03	91	0.40	3.154480e-03
42	-0.09	-2.624631e-03	92	0.41	3.209553e-03
43	-0.08	-2.384824e-03	93	0.42	3.263463e-03
44	-0.07	-2.154693e-03	94	0.43	3.316249e-03
45	-0.06	-1.933713e-03	95	0.44	3.367946e-03
46	-0.05	-1.721393e-03	96	0.45	3.418590e-03
47	-0.04	-1.517276e-03	97	0.46	3.468214e-03
48	-0.03	-1.320932e-03	98	0.47	3.516850e-03
49	-0.02	-1.131958e-03	99	0.48	3.564528e-03
50	-0.01	-9.499771e-04	100	0.49	3.611278e-03
			101	0.50	3.657128e-03

<- Minimum Slope rounds to
No More Than 0.00001
from Zero Slope

6.3 Appendix C

Analytically derived Tikhonov Regularized Left-Inverse Matrix (TRLIM) Algorithm from the User's Genre-Element Bias Loss Function (for context, refer to Section 3.5 and Equations 4.3a and 4.3b):

$$\text{Loss Function} \quad \arg \min_{\mathbf{x}} \left\{ \|\mathbf{Ax} - \mathbf{y}\|_2^2 + \lambda_T \|\mathbf{x}\|_2^2 \right\} \quad (\text{C.1})$$

$$\arg \min_{\mathbf{x}} \left\{ (\mathbf{Ax} - \mathbf{y})^\top (\mathbf{Ax} - \mathbf{y}) + \lambda_T (\mathbf{x}^\top \mathbf{x}) \right\} \quad (\text{C.2})$$

$$\nabla_{\mathbf{x}} \left\{ (\mathbf{Ax} - \mathbf{y})^\top (\mathbf{Ax} - \mathbf{y}) + \lambda_T (\mathbf{x}^\top \mathbf{x}) \right\} = \mathbf{0} \quad (\text{C.3})$$

$$\nabla_{\mathbf{x}} \left\{ ((\mathbf{Ax})^\top - \mathbf{y}^\top) (\mathbf{Ax} - \mathbf{y}) + \lambda_T (\mathbf{x}^\top \mathbf{x}) \right\} = \mathbf{0} \quad (\text{C.4})$$

$$\nabla_{\mathbf{x}} \left\{ (\mathbf{x}^\top \mathbf{A}^\top - \mathbf{y}^\top) (\mathbf{Ax} - \mathbf{y}) + \lambda_T (\mathbf{x}^\top \mathbf{x}) \right\} = \mathbf{0} \quad (\text{C.5})$$

$$\nabla_{\mathbf{x}} \left\{ \mathbf{x}^\top \mathbf{A}^\top \mathbf{Ax} - \mathbf{y}^\top \mathbf{Ax} - \mathbf{x}^\top \mathbf{A}^\top \mathbf{y} + \mathbf{y}^\top \mathbf{y} + \lambda_T (\mathbf{x}^\top \mathbf{x}) \right\} = \mathbf{0} \quad (\text{C.6})$$

$$\mathbf{A}^\top \mathbf{Ax} + (\mathbf{A}^\top \mathbf{A})^\top \mathbf{x} - (\mathbf{y}^\top \mathbf{A})^\top - \mathbf{A}^\top \mathbf{y} + 2\lambda_T \mathbf{x} = \mathbf{0} \quad (\text{C.7})$$

$$\mathbf{A}^\top \mathbf{Ax} + \mathbf{A}^\top (\mathbf{A}^\top)^\top \mathbf{x} - \mathbf{A}^\top (\mathbf{y}^\top)^\top - \mathbf{A}^\top \mathbf{y} + 2\lambda_T \mathbf{x} = \mathbf{0} \quad (\text{C.8})$$

$$\mathbf{A}^\top \mathbf{Ax} + \mathbf{A}^\top \mathbf{Ax} - \mathbf{A}^\top \mathbf{y} - \mathbf{A}^\top \mathbf{y} + 2\lambda_T \mathbf{x} = \mathbf{0} \quad (\text{C.9})$$

$$2\mathbf{A}^\top \mathbf{Ax} - 2\mathbf{A}^\top \mathbf{y} + 2\lambda_T \mathbf{x} = \mathbf{0} \quad (\text{C.10})$$

$$2\mathbf{A}^\top \mathbf{Ax} + 2\lambda_T \mathbf{x} = 2\mathbf{A}^\top \mathbf{y} \quad (\text{C.11})$$

$$\mathbf{A}^\top \mathbf{Ax} + \lambda_T \mathbf{x} = \mathbf{A}^\top \mathbf{y} \quad (\text{C.12})$$

$$(\mathbf{A}^\top \mathbf{A} + \lambda_T \mathbf{I}) \mathbf{x} = \mathbf{A}^\top \mathbf{y} \quad (\text{C.13})$$

$$(\mathbf{A}^\top \mathbf{A} + \lambda_T \mathbf{I})^{-1} (\mathbf{A}^\top \mathbf{A} + \lambda_T \mathbf{I}) \mathbf{x} = (\mathbf{A}^\top \mathbf{A} + \lambda_T \mathbf{I})^{-1} \mathbf{A}^\top \mathbf{y} \quad (\text{C.14})$$

$$\text{TRLIM Algorithm} \quad \mathbf{x} = (\mathbf{A}^\top \mathbf{A} + \lambda_T \mathbf{I})^{-1} \mathbf{A}^\top \mathbf{y} \quad (\text{C.15})$$

Table 12: Zwillinger (2018), *CRC Standard Mathematical Tables and Formulas, 33rd Edition*, p. 287

Differentiation with respect to a vector \mathbf{x} . (Here A is constant.) $\frac{\partial A\mathbf{x}}{\partial \mathbf{x}^\top} = A$ and

y (a scalar or a vector)	$\frac{\partial y}{\partial \mathbf{x}}$ (where \mathbf{x} is a vector)
\mathbf{x}^\top	I
$A\mathbf{x}$	A^\top
$\mathbf{x}^\top A$	A
$\mathbf{x}^\top \mathbf{x}$	$2\mathbf{x}$
$\mathbf{x}^\top A\mathbf{x}$	$A\mathbf{x} + A^\top \mathbf{x}$

Table 13: Deisenroth et al. (2023), *Mathematics for Machine Learning*, p. 25

The following are important properties of inverses and transposes:

$$\mathbf{A}\mathbf{A}^{-1} = \mathbf{I} = \mathbf{A}^{-1}\mathbf{A} \quad (\text{2.26})$$

$$(\mathbf{AB})^{-1} = \mathbf{B}^{-1}\mathbf{A}^{-1} \quad (\text{2.27})$$

$$(\mathbf{A} + \mathbf{B})^{-1} \neq \mathbf{A}^{-1} + \mathbf{B}^{-1} \quad (\text{2.28})$$

$$(\mathbf{A}^\top)^\top = \mathbf{A} \quad (\text{2.29})$$

$$(\mathbf{A} + \mathbf{B})^\top = \mathbf{A}^\top + \mathbf{B}^\top \quad (\text{2.30})$$

$$(\mathbf{AB})^\top = \mathbf{B}^\top \mathbf{A}^\top \quad (\text{2.31})$$

7 References

- Abdi, Hervé, and Lynne J. Williams. (2010). “Principal Component Analysis.” *Wiley interdisciplinary reviews: computational statistics*, Vol. 2, No. 4, pp.433-459.
<https://doi.org/10.1002/wics.101>
<https://personal.utdallas.edu/~herve/abdi-awPCA2010.pdf>
- Ackerman, Brian. (2012). *Context-Aware Rank-Oriented Recommender Systems*. MS Thesis, June 2012. Arizona State University.
https://keep.lib.asu.edu/_flysystem/fedora/c7/66292/tmp/package-7wF0rK/Ackerman_asu_0010N_12443.pdf
- Aggarwal, Charu C. (2020). “An Optimization-Centric View of Linear Systems.” In *Linear Algebra and Optimization for Machine Learning*, pp. 79-82. Springer Nature Switzerland AG, ©2020. ISBN: 978-3-030-40344-0
<https://doi.org/10.1007/978-3-030-40344-7>
- Bates, Douglas, Martin Mächler, Ben Bolker, and Steve Walker. (2015). “Fitting Linear Mixed-Effects Models Using lme4”. *Journal of Statistical Software*, Vol. 67, No. 1, pp. 1-48.
<https://doi.org/10.18637/jss.v067.i01>
<https://www.jstatsoft.org/index.php/jss/article/view/v067i01/946>
- Bell, Robert M., and Yehuda Koren. (2007). “Scalable Collaborative Filtering with Jointly Derived Neighborhood Interpolation Weights.” In *Seventh IEEE International Conference on Data Mining (ICDM 2007)*, pp. 43-52. IEEE, 2007.
<https://citeseerx.ist.psu.edu/document?repid=rep1&type=pdf&doi=0c953e630cccd64d2ad5fbb09f08425e7f82b7a3>
- Brunton, Steven L., and J. Nathan Kutz. (2022). *Data-Driven Science and Engineering: Machine Learning, Dynamical Systems, and Control (2nd Edition)*. Cambridge: Cambridge University Press, ©2022. ISBN: 978-1-009-09848-9
<https://doi.org/10.1017/9781009089517>
- Chen, Tianqi, Zhao Zheng, Qiuxia Lu, Weinan Zhang, and Yong Yu. (2011). “Feature-Based Matrix Factorization.” *arXiv preprint*, arXiv:1109.2271v3 [cs.AI] 29 Dec 2011.
<https://arxiv.org/pdf/1109.2271.pdf>
- Chen, Tianqi, Weinan Zhang, Qiuxia Lu, Kailong Chen, Zhao Zheng and Yong Yu. (2012). “SVDFeature: A Toolkit for Feature-based Collaborative Filtering”. *Journal of Machine Learning Research*, Vol. 13, No. 116, pp. 3619-3622.
<https://jmlr.org/papers/volume13/chen12a/chen12a.pdf>
- Cortes, David. (2018). “Cold-Start Recommendations in Collective Matrix Factorization.” *arXiv preprint*, arXiv:1809.00366.
<https://arxiv.org/pdf/1809.00366.pdf>
- Deisenroth, Marc Peter, A. Aldo Faisal, and Cheng Soon Ong. (2023). *Mathematics for Machine Learning*. Cambridge: Cambridge University Press, ©2020. ISBN: 9781108455145
 Free Online 2023-02-15 Draft: <https://mml-book.github.io/book/mml-book.pdf>
- Efron, Bradley. (2023). “Machine Learning and the James–Stein Estimator.” *Japanese Journal of Statistics and Data Science*, pp. 1-10. Springer Online: 30 June 2023.
<https://link.springer.com/article/10.1007/s42081-023-00209-y>
<https://link.springer.com/content/pdf/10.1007/s42081-023-00209-y.pdf>
- Efron, Bradley, and Carl Morris. (1977). “Stein’s Paradox in Statistics.” *Scientific American*, Vol. 236, No. 5, pp. 119-127, May 1977.
<https://efron.ckirby.su.domains//other/Article1977.pdf>

- Feuerverger, Andrey, Yu He, and Shashi Khatri. (2012). “Statistical Significance of the Netflix Challenge.” *Statistical Science*, Vol. 27, No. 2, pp. 202-231.
<https://projecteuclid.org/journals/statistical-science/volume-27/issue-2/Statistical-Significance-of-the-Netflix-Challenge/10.1214/11-STS368.pdf>
<https://doi.org/10.1214/11-STS368>
- Garcia, Cyril, and Luca Rona. (2015?). “The Netflix Challenge.”
https://www.researchgate.net/profile/Cyril-Garcia/publication/326694752_The_Netflix_Challenge/links/5b5fb50c458515c4b2543cd1/The-Netflix-Challenge.pdf
- Gareth, James, Daniela Witten, Trevor Hastie, and Robert Tibshirani. (2021). *An Introduction to Statistical Learning: with Applications in R (2nd Edition)*. New York: Springer Nature, ©2021.
 ISBN: 978-1-0716-1418-1
<https://doi.org/10.1007/978-1-0716-1418-1>
- Garg, Amita, and Komal Goyal. (2014). “Regularization Methods to Solve Various Inverse Problems.” *International Journal of Science and Research*, Vol. 3, No. 8, pp. 1840-1845.
https://www.researchgate.net/profile/Komal-Goyal/publication/317101803_Regularization_Methods_to_Solve_Various_Inverse_Problems/links/592678d5a6fdcc444342271c/Regularization-Methods-to-Solve-Various-Inverse-Problems.pdf
- Gavish, M. and D. L. Donoho. (2014). “The optimal hard threshold for singular values is $4/\sqrt{3}$.” *IEEE Transactions on Information Theory*, Vol. 60, No. 8, pp. 5040–5053, August 2014.
<https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=6846297>
- Ghazanfar, Mustansar Ali, and Adam Prugel. (2013). “The Advantage of Careful Imputation Sources in Sparse Data-Environment of Recommender Systems: Generating Improved SVD-Based Recommendations.” *Informatica*, Vol. 37, No. 1, pp. 61-92.
<https://www.informatica.si/index.php/informatica/article/download/435/439>
- Goldberg, David, David Nichols, Brian M. Oki, and Douglas Terry. (1992). “Using Collaborative Filtering to Weave an Information Tapestry.” *Communications of the ACM*, Vol. 35, No. 12, pp. 61-70.
<https://dl.acm.org/doi/pdf/10.1145/138859.138867>
- Hahsler, Michael. (2022). “recommenderlab: An R framework for Developing and Testing Recommendation Algorithms.” *arXiv e-prints*, arXiv:2205.12371v1 [cs.IR] 24 May 2022. DOI: 10.48550/arXiv.2205.12371
<https://arxiv.org/pdf/2205.12371.pdf>
- Hansen, Per Christian. (1987). “The Truncated SVD as a Method for Regularization.” *BIT Numerical Mathematics*, Vol. 27, pp. 534-553.
<http://i.stanford.edu/pub/cstr/reports/na/m/86/36/NA-M-86-36.pdf>
- Harper, F. Maxwell and Joseph A. Konstan. (2015). “The MovieLens Datasets: History and Context.” *ACM Transactions on Interactive Intelligent Systems (TiiS)* Vol. 5, No. 4, Article 19 (22 December 2015), pp. 1-19. <https://dx.doi.org/10.1145/2827872>
<https://files.grouplens.org/papers/harper-tiis2015.pdf>
- Hastie, Trevor, Rahul Mazumder, Jason D. Lee, and Reza Zadeh. (2015). “Matrix Completion and Low-Rank SVD via Fast Alternating Least Squares.” *The Journal of Machine Learning Research*, Vol. 16, No. 1, pp. 3367-3402.
<https://www.jmlr.org/papers/volume16/hastie15a/hastie15a.pdf>
- Hu, Yifan, Yehuda Koren, and Chris Volinsky. (2008). “Collaborative Filtering for Implicit Feedback Datasets.” In *2008 Eighth IEEE International Conference on Data Mining*, pp. 263-272. IEEE, 2008.
https://www.researchgate.net/profile/Yifan-Hu-25/publication/220765111_Collaborative_Filtering_for_Implicit_Feedback_Datasets/links/0912f509c579ddd954000000/Collaborative-Filtering-for-Implicit-Feedback-Datasets.pdf
- Irizarry, Rafael A. (2022). “Introduction to Data Science - Data Analysis and Prediction Algorithms with R.” 2022-05-17. <https://rafalab.github.io/dsbook/>

- Ito, Kazufumi, and Bangti Jin. (2020). “Regularized Linear Inversion with Randomized Singular Value Decomposition.” In *Mathematical and Numerical Approaches for Multi-Wave Inverse Problems: CIRM, Marseille, France, April 1–5, 2019*, pp. 45-72. Springer International Publishing, 2020.
<https://arxiv.org/pdf/1909.01947.pdf>
- Jadhav, Anil, Dhanya Pramod, and Krishnan Ramanathan. (2019). “Comparison of Performance of Data Imputation Methods for Numeric Dataset.” *Applied Artificial Intelligence*, Vol. 33, No. 10, pp. 913-933.
<https://www.tandfonline.com/doi/pdf/10.1080/08839514.2019.1637138?needAccess=true&role=button>
- Jolani, Shahab. (2018). “Hierarchical Imputation of Systematically and Sporadically Missing Data: An Approximate Bayesian Approach Using Chained Equations.” *Biometrical Journal*, Vol. 60, No. 2, pp. 333-351.
<https://doi.org/10.1002/bimj.201600220>
- Kokla, Marietta, Jyrki Virtanen, Marjukka Kolehmainen, Jussi Paananen, and Kati Hanhineva. (2019). “Random Forest-Based Imputation Outperforms Other Methods for Imputing LC-MS Metabolomics Data: a Comparative Study.” *BMC Bioinformatics*, Vol. 20, No. 1, pp. 1-11.
<https://bmcbioinformatics.biomedcentral.com/counter/pdf/10.1186/s12859-019-3110-0.pdf>
- Koren, Yehuda. (2009). “Collaborative Filtering with Temporal Dynamics.” *KDD '09: Proceedings of the 15th ACM SIGKDD international conference on Knowledge discovery and data mining*. (June 2009), pp. 447–456. <https://doi.org/10.1145/1557019.1557072>
<https://cseweb.ucsd.edu/classes/fa17/cse291-b/reading/p447-koren.pdf>
- Koren, Yehuda, Robert Bell, and Chris Volinsky. (2009). “Matrix Factorization Techniques for Recommender Systems.” *Computer*, Vol. 42, No. 8, pp. 30-37.
<https://www.inf.unibz.it/~ricci/ISR/papers/ieeecomputer.pdf>
- Li, Dongsheng, Chao Chen, Qin Lv, Junchi Yan, Li Shang, and Stephen Chu. (2016). “Low-rank Matrix Approximation with Stability.” In *International Conference on Machine Learning*, pp. 295-303. PMLR.
<https://proceedings.mlr.press/v48/lib16.pdf>
- Liu, Xiaofeng, Xue Wang, Lang Zou, Jing Xia, and Wei Pang. (2020). “Spatial Imputation for Air Pollutants Data Sets via Low Rank Matrix Completion Algorithm.” *Environment International*, Vol. 139, p. 105713.
<https://www.sciencedirect.com/science/article/pii/S0160412019341704>
- Marlin, Benjamin, Richard S. Zemel, Sam Roweis, and Malcolm Slaney. (2012). “Collaborative Filtering and the Missing at Random Assumption.” *arXiv preprint*, arXiv:1206.5267, 2012 - arxiv.org
<https://arxiv.org/pdf/1206.5267>
- Meng, Zaiqiao, Richard McCreadie, Craig Macdonald, and Iadh Ounis. (2020). “Exploring Data Splitting Strategies for the Evaluation of Recommendation Models.” In *Proceedings of the 14th ACM Conference on Recommender Systems*, pp. 681-686.
<https://arxiv.org/pdf/2007.13237.pdf>
- Mesgarani, Hamid, and Yaqub Azari. (2019). “Numerical investigation of Fredholm integral equation of the first kind with noisy data.” *Mathematical Sciences*, Vol. 13, pp. 267-278.
<https://link.springer.com/article/10.1007/s40096-019-00296-7>
- Nanda, Neel, Lawrence Chan, Tom Liberum, Jess Smith, and Jacob Steinhardt. (2023). “Progress Measures for Grokking via Mechanistic Interpretability.” *arXiv preprint arXiv:2301.05217v2[cs.LG]* 13 Jan 2023.
<https://arxiv.org/pdf/2301.05217.pdf>
- Ng, Andrew Y. (2004). “Feature selection, L_1 vs. L_2 regularization, and rotational invariance.” In *Proceedings of the 21st International Conference on Machine Learning-ICML*, Vol. 4, pp. 78-85. Banff, Canada, July 2004. <https://icml.cc/Conferences/2004/proceedings/papers/354.pdf>
- Nguyen, Luong Trung, Junhan Kim, and Byonghyo Shim. (2019). “Low-Rank Matrix Completion: A Contemporary Survey.” *IEEE Access*, Vol. 7, pp. 94215-94237.
<https://ieeexplore.ieee.org/stamp/stamp.jsp?arnumber=8759045>

- Paraschakis, Dimitris. (2016). “Recommender Systems from an Industrial and Ethical Perspective.” In *Proceedings of the 10th ACM Conference on Recommender Systems*, pp. 463-466.
<https://www.diva-portal.org/smash/get/diva2:1409571/FULLTEXT01.pdf>
- Paterek, Arkadiusz. (2007). “Improving regularized singular value decomposition for collaborative filtering.” In *Proceedings of KDD Cup and Workshop*, vol. 2007, pp. 5-8.
<https://www.cs.uic.edu/~liub/KDD-cup-2007/proceedings/Regular-Paterek.pdf>
- Pes, Federica, and Giuseppe Rodriguez. (2020). “The Minimal-Norm Gauss-Newton Method and Some of its Regularized Variants.” *Electronic Transactions on Numerical Analysis*, Vol. 53, pp. 459-480.
<https://iris.unica.it/bitstream/11584/293576/1/nonlinreg20.pdf>
- Pilászy, István. (2010). “Factorization-based large scale recommendation algorithms.” PhD Thesis, January 2009. Department of Measurement and Information Systems, Budapest University of Technology and Economics, Hungary.
<https://repozitorium.omikk.bme.hu/bitstream/handle/10890/977/ertekezes.pdf?sequence=1>
- Rendle, Steffen, Li Zhang, and Yehuda Koren. (2019). “On the Difficulty of Evaluating Baselines - A Study on Recommender Systems.” *arXiv e-prints*, arXiv:1905.01395v1 [cs.IR] 4 May 2019.
<https://arxiv.org/pdf/1905.01395.pdf>
- Resche-Rigon, Matthieu, and Ian R. White. (2018). “Multiple Imputation by Chained Equations for Systematically and Sporadically Missing Multilevel Data.” *Statistical Methods in Medical Research*, Vol. 27, No. 6, pp. 1634-1649.
<https://www.ncbi.nlm.nih.gov/pmc/articles/PMC5496677/pdf/emss-71930.pdf>
- Santurkar, Shibani, Dimitris Tsipras, Andrew Ilyas, and Aleksander Madry. (2018). “How Does Batch Normalization Help Optimization?” *Advances in Neural Information Processing Systems*, Vol. 31.
<https://proceedings.neurips.cc/paper/2018/file/905056c1ac1dad141560467e0a99e1cf-Paper.pdf>
<https://arxiv.org/pdf/1805.11604.pdf> 2019 ArXiv version includes Appendices of Experimental Setup, Omitted Figures, and Mathematical Proofs.
- Shlens, Jonathon. (2014). “A Tutorial on Principal Component Analysis.” *arXiv preprint*, arXiv:1404.1100.
<https://arxiv.org/pdf/1404.1100.pdf>
- Steck, Harald, Linas Baltrunas, Ehtsham Elahi, Dawen Liang, Yves Raimond, and Justin Basilico. (2021). “Deep Learning for Recommender Systems: A Netflix Case Study.” *AI Magazine*, Vol. 42, No. 3, pp. 7-18.
<https://ojs.aaai.org/index.php/aimagazine/article/view/18140/18876>
- Takács, Gábor, István Pilászy, Botyán Németh, and Domonkos Tikk. (2007). “Major Components of the Gravity Recommendation System.” *Acm Sigkdd Explorations Newsletter*, 9(2), pp.80-83.
https://web.archive.org/web/20190924051502id_/http://www.sze.hu/~gtakacs/download/sigkdd_explorations_2007.pdf
- Webb, Brandyn (a.k.a., Simon Funk). (2006). “Netflix Update: Try This at Home.”
<https://sifter.org/~simon/journal/20061211.html>
- Wikipedia Contributors. (2022, December 24). “Early Stopping.” In *Wikipedia, The Free Encyclopedia*. Retrieved 23:10 UTC, February 19, 2023, from https://en.wikipedia.org/w/index.php?title=Early_stopping&oldid=1129334456
- Wikipedia Contributors. (2023, February 13). “Cholesky Decomposition.” In *Wikipedia, The Free Encyclopedia*. Retrieved 15:59 UTC, March 30, 2023, from https://en.wikipedia.org/w/index.php?title=Cholesky_decomposition&oldid=1139072844
- Wikipedia Contributors. (2023, June 18). “Missing Data.” In *Wikipedia, The Free Encyclopedia*. Retrieved 09:21 UTC, July 13, 2023, from https://en.wikipedia.org/w/index.php?title=Missing_data&oldid=1160761517
- Wikipedia Contributors. (2023, March 20). “Matrix Completion.” In *Wikipedia, The Free Encyclopedia*. Retrieved 17:09 UTC, March 30, 2023, from https://en.wikipedia.org/w/index.php?title=Matrix_completion&oldid=1145758143

Wu, Le, Xiangnan He, Xiang Wang, Kun Zhang, and Meng Wang. (2022). “A Survey on Accuracy-Oriented Neural Recommendation: From Collaborative Filtering to Information-Rich Recommendation.” *IEEE Transactions on Knowledge and Data Engineering*.

<https://arxiv.org/pdf/2104.13030.pdf>

Xiang, Hua, and Jun Zou. (2013). “Regularization with Randomized SVD for Large-Scale Discrete Inverse Problems.” *Inverse Problems*, Vol. 29, No. 8: 085008.

https://www.researchgate.net/profile/Hua-Xiang-11/publication/258228821_Regularization_with_randomized_SVD_for_large-scale_discrete_inverse_problems/links/00b49527f4ce979d29000000/Regularization-with-randomized-SVD-for-large-scale-discrete-inverse-problems.pdf

Xu, Junlin, Lijun Cai, Bo Liao, Wen Zhu, and JiaLiang Yang. (2020). “CMF-Impute: an Accurate Imputation Tool for Single-Cell RNA-seq Data.” *Bioinformatics*, Vol. 36, No. 10, pp. 3139-3147.

<https://academic.oup.com/bioinformatics/article/36/10/3139/5740569>

Zhang, Shuai, Lina Yao, Aixin Sun, and Yi Tay. (2018). “Deep Learning Based Recommender System: A Survey and New Perspectives.” *ACM Computing Surveys*, Vol. 1, No. 1, Article 1 (July 2018), 35 pages. *arXiv e-prints*, arXiv:1707.07435v7 [cs.IR] 10 Jul 2019.

<https://arxiv.org/pdf/1707.07435.pdf>

Zhou, Yunhong, Dennis Wilkinson, Robert Schreiber, and Rong Pan. (2008). “Large-Scale Parallel Collaborative Filtering for the Netflix Prize.” In *Algorithmic Aspects in Information and Management: 4th International Conference, AAIM 2008, Shanghai, China, June 23-25, 2008*. Proceedings 4, pp. 337-348. Springer-Verlag Berlin Heidelberg, ©2008.

ISBN: 978-3-540-68865-5

https://www.researchgate.net/profile/Robert-Schreiber-6/publication/220788980_Large-Scale_Parallel_Collaborative_Filtering_for_the_Netflix_Prize/links/00b7d5213e07f3dbcc000000/Large-Scale-Parallel-Collaborative-Filtering-for-the-Netflix-Prize.pdf

Zwillinger, Dan (Editor). (2018). *CRC Standard Mathematical Tables and Formulas, 33rd Edition*. CRC Press, ©2018.

ISBN: 978-1-4987-7780-3

[https://www.softouch.on.ca/kb/data/CRC%20Standard%20Mathematical%20Tables%20and%20Formulas%2033E%20\(2018\).pdf](https://www.softouch.on.ca/kb/data/CRC%20Standard%20Mathematical%20Tables%20and%20Formulas%2033E%20(2018).pdf)

8